



## **BIG DATA**

# **Map Reduce Programming model and Architecture**

---

**K V Subramaniam**

Computer Science and Engineering

# BIG DATA

## Map Reduce Programming model and Architecture

---



### What we have learnt so far..

- Large amounts of data to be processed.
- We have HDFS as a **distributed** store.
- We need to distribute the processing also.

# BIG DATA

## Map Reduce Programming model and Architecture

---

What is Map Reduce ?



# BIG DATA

## Map Reduce Programming model and Architecture



### Why Map-Reduce ?

- A new fundamental way to process extremely large data (?)
- We are going to
  - Study Map-Reduce paradigm
  - Study Hadoop architecture
    - Open Source implementation of Map-Reduce

# BIG DATA

## What is MapReduce ?

---



- Origin from Google, [OSDI'04]
- A simple programming model
- Functional model
- For large-scale data processing
  - Exploits large set of commodity computers
  - Executes process in distributed manner
  - Offers high availability

# Big Data: Map Reduce Motivating Example

# BIG DATA

## Map Reduce - Motivation

---

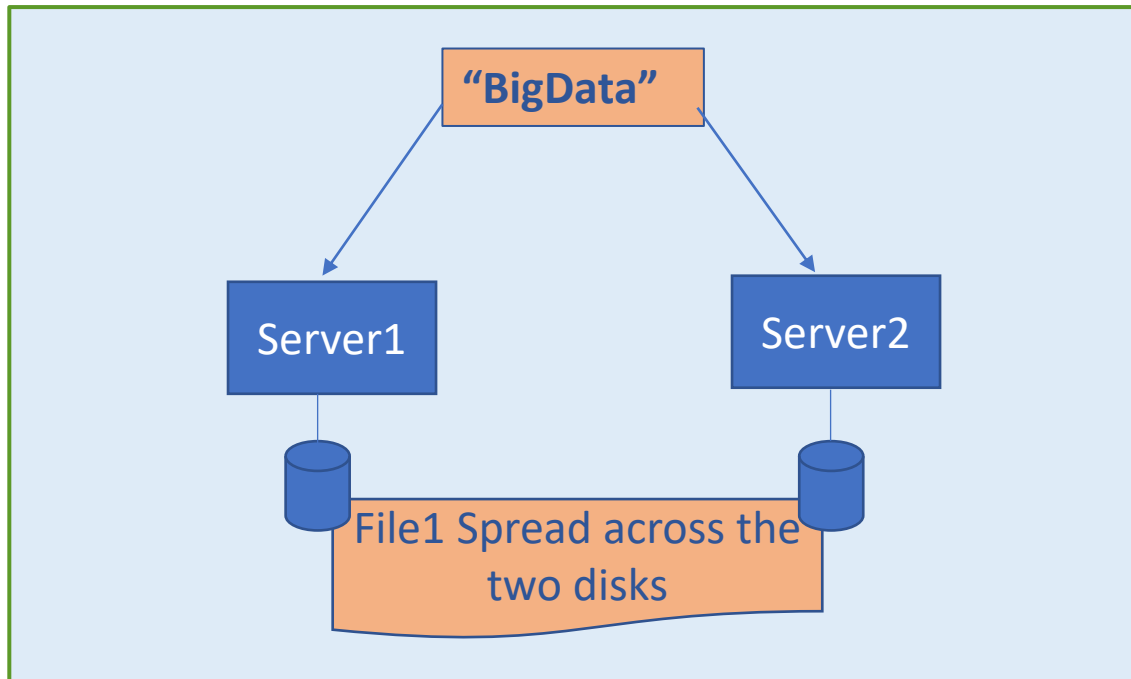


- Lots of demands for very large scale data processing
- Lots of machines needed (scaling)
- Two basic operations on the input
  - Map
  - Reduce
- To understand what Map/Reduce really are..

# BIG DATA

## A MapReduce Example

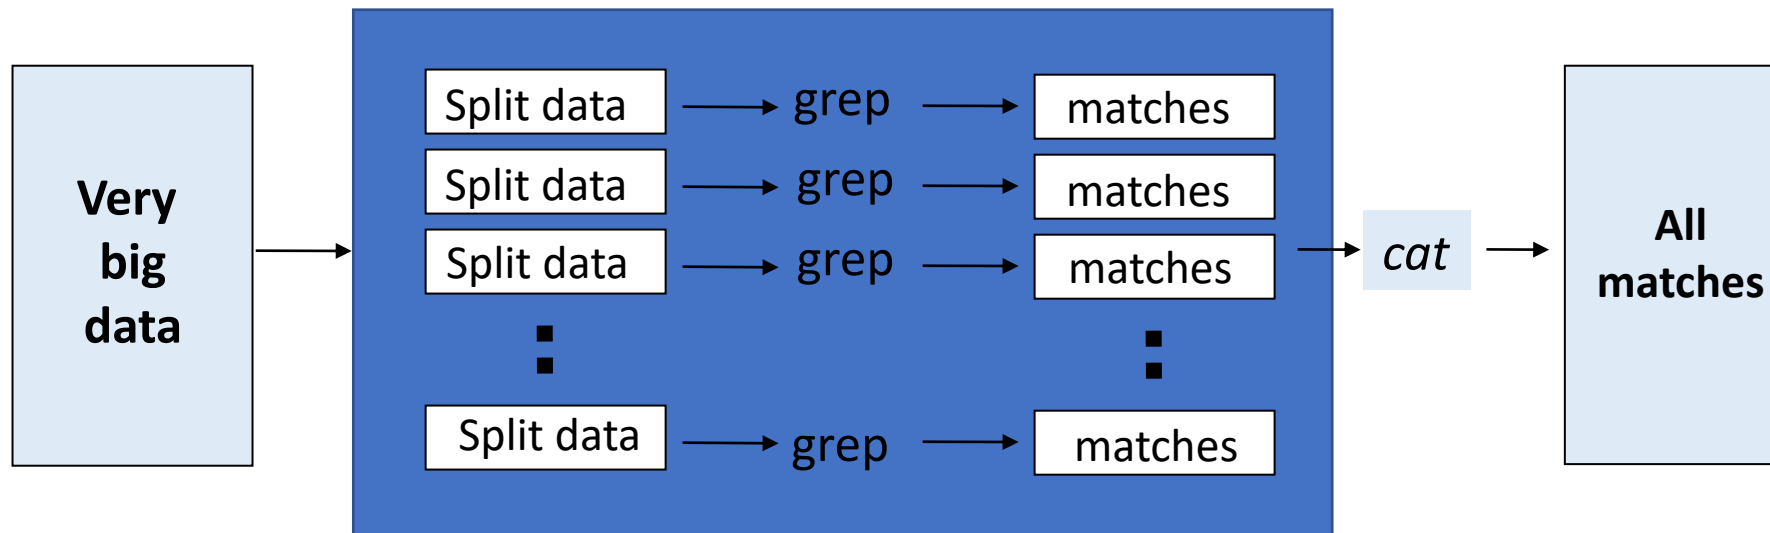
- Consider a very large text file and you want to determine if a word exists in this file?
  - The file is stored in HDFS across two machines.
  - How will you search to check if the word “BigData” is present in the file?





# BIG DATA

## Map Reduce: Distributed Grep-Solution



## Distributed Word Count



# BIG DATA

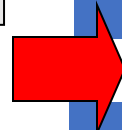
Example: find the number of restaurants offering each item?

Menu 1

Idli	Vada
Pizza	

Menu 2

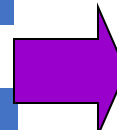
Dosa	Pizza
Burger	



Idli	1
Vada	1
Pizza	1

Dosa	1
Pizza	1
Burger	1



Merged  
results

Burger	1
Dosa	1
Idli	1
Pizza	2
Vada	1

- Suppose we need parallelism of the merge program.
- Would the earlier approach work?
- What do we need to do?

# BIG DATA


Example: find the number of restaurants offering each item?

Menu 1

Idli	Vada
Pizza	

Menu 2

Dosa	Pizza
Burger	



- Treat output of first program as a key value pair
- Partition the keys between the second program

# BIG DATA


## Distributed Word Count

Menu 1

Idli Vada  
Pizza

Menu 2

Dosa Pizza  
Burger



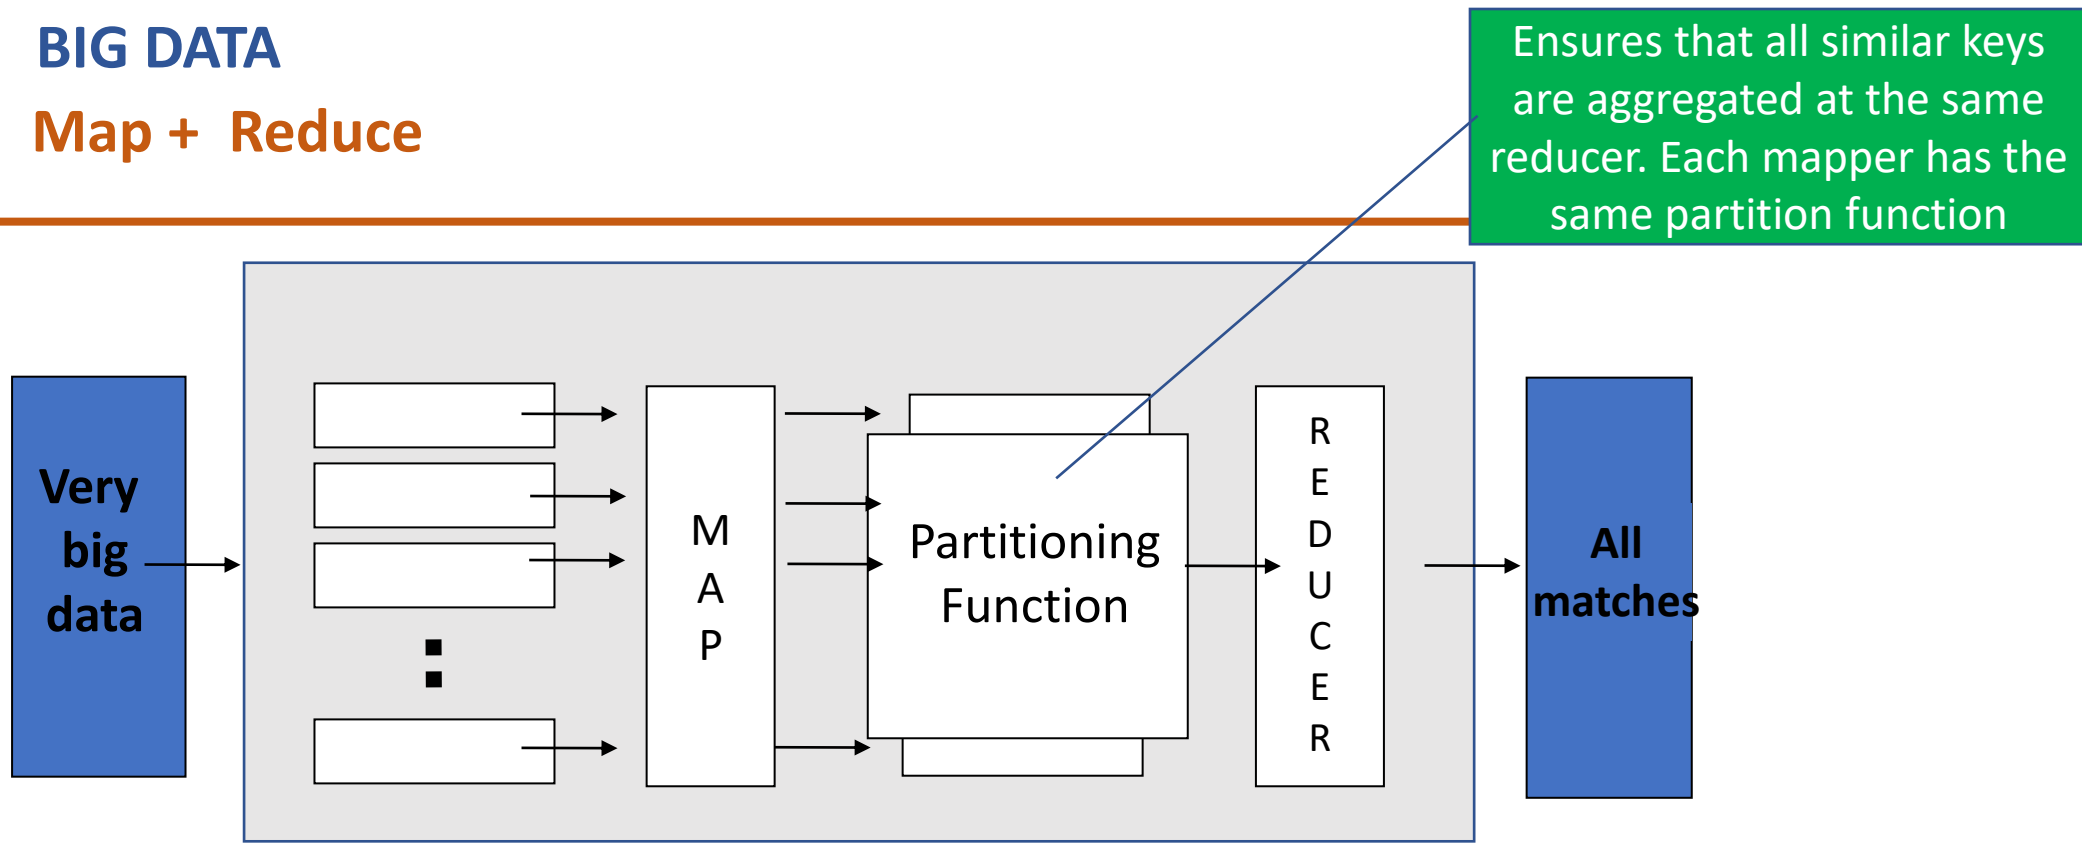
All keys starting  
from A-M

All keys starting  
from N-Z

How do we know that all  
the “pizza” keys must be  
aggregated in server 2?

# BIG DATA

## Map + Reduce



- Map:

- Accepts *input* key/value pair
- Emits *intermediate* key/value pair

- Reduce :

- Accepts *intermediate* key/value\* pair
- Emits *output* key/value pair

## Map Reduce: A look at the code

# BIG DATA

## Map Reduce Programming model

---



- **Data type:** key-value *records*
- **Map function:**

$$(K_{in}, V_{in}) \rightarrow \text{list}(K_{inter}, V_{inter})$$

- **Reduce function:**

$$(K_{inter}, \text{list}(V_{inter})) \rightarrow \text{list}(K_{out}, V_{out})$$



```
public static class TokenizerMapper
    extends Mapper<Object, Text, Text, IntWritable>{

    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    public void map(Object key, Text value, Context context
        ) throws IOException, InterruptedException {
        StringTokenizer itr = new StringTokenizer(value.toString());
        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            context.write(word, one);
        }
    }
}
```

(Key , value)

List (Key ,value)

$(K_{in}, V_{in}) \rightarrow list(K_{inter}, V_{inter})$

```
public static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values,
                      Context context
                      ) throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}
```

**Key ,List (value)**

**List (Key ,value)**

$(K_{inter}, list(V_{inter})) \rightarrow list(K_{out}, V_{out})$

```
public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    String[] otherArgs = new GenericOptionsParser(conf, args).
        getRemainingArgs();
    if (otherArgs.length < 2) {
        System.err.println("Usage: wordcount <in> [<in>...] <out>");
        System.exit(2);
    }
    Job job = new Job(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);

    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    for (int i = 0; i < otherArgs.length - 1; ++i) {
        FileInputFormat.addInputPath(job, new Path(otherArgs[i]));
    }
    FileOutputFormat.setOutputPath(job,
        new Path(otherArgs[otherArgs.length - 1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
```

**Set Mapper  
and  
Reducer class**

## Map Reduce: Sample Exercise

# BIG DATA

## Map Reduce , Search

---



**Input:** file(lineNumber, line) records and pattern

**Output:** lines matching a given pattern

**What will be the mapper and reducer? What will be the keys?**

**Map:**

**Reduce:**

# BIG DATA

## Map Reduce , Search - Solution

---



**Input:** file (lineNumber, line) records and pattern

**Output:** lines matching a given pattern

**Map:**

```
if(line matches pattern):  
    output(line)
```

**Reduce:** identity function

—Alternative: no reducer (map-only job)

# BIG DATA

## Map Reduce , Functions in the Model

---



- **Map**

- Process a key/value pair to generate intermediate key/value pairs
- Sorts all key/value pairs before sending to reducer

- **Reduce**

- Merge all intermediate values associated with the same key
- Runs after all Map tasks are finished (why?)

- **Partition**

- By default : **hash(key) mod R** (Well balanced)
- There are cases where this can be more complex

## Map Reduce: Revision exercise



# BIG DATA

## Map Reduce, Sort

---



**Input:** (key, value) records

**Output:** same records, sorted by key

**What will be the mapper and reducer?**

**What will be the partition function?**

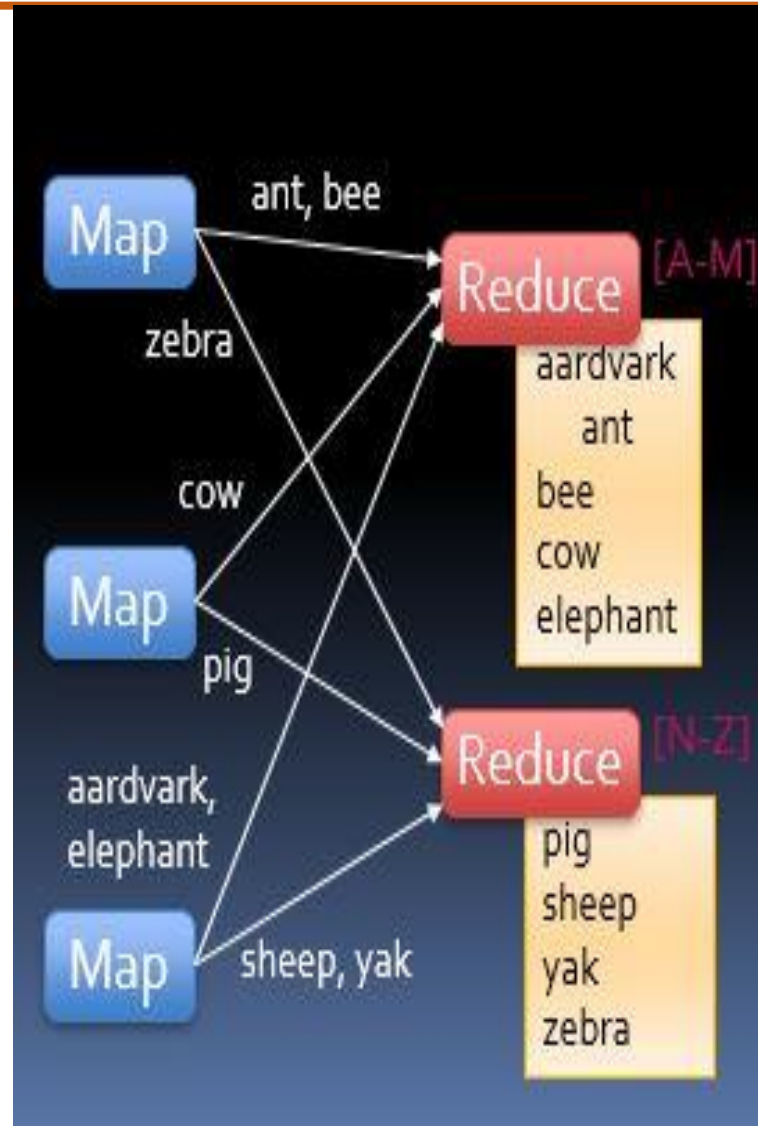
**Map:**

**Reduce:**

# BIG DATA

## Map Reduce, Sort - Solution

- **Input:** (key, value) records
- **Output:** same records, sorted by key
- **Map:** identity function
- **Reduce:** identity function
- **Trick:** Pick partitioning function  $p$  so that  $k_1 < k_2 \Rightarrow p(k_1) < p(k_2)$
- Works because map sorts output keys.




## Map Reduce: Job Submission Flow

# BIG DATA

## Map Reduce , Hadoop Flow.

- User submits job
  - input data, MapReduce program, and configuration information
- How is parallelization achieved?
  - Divide input into smaller chunks – called **input splits**
- Hadoop divides jobs into tasks
  - Map tasks, reduce tasks
  - One map task per split
  - Tasks run in parallel



# BIG DATA

## MapReduce Flow for A SINGLE REDUCE Task

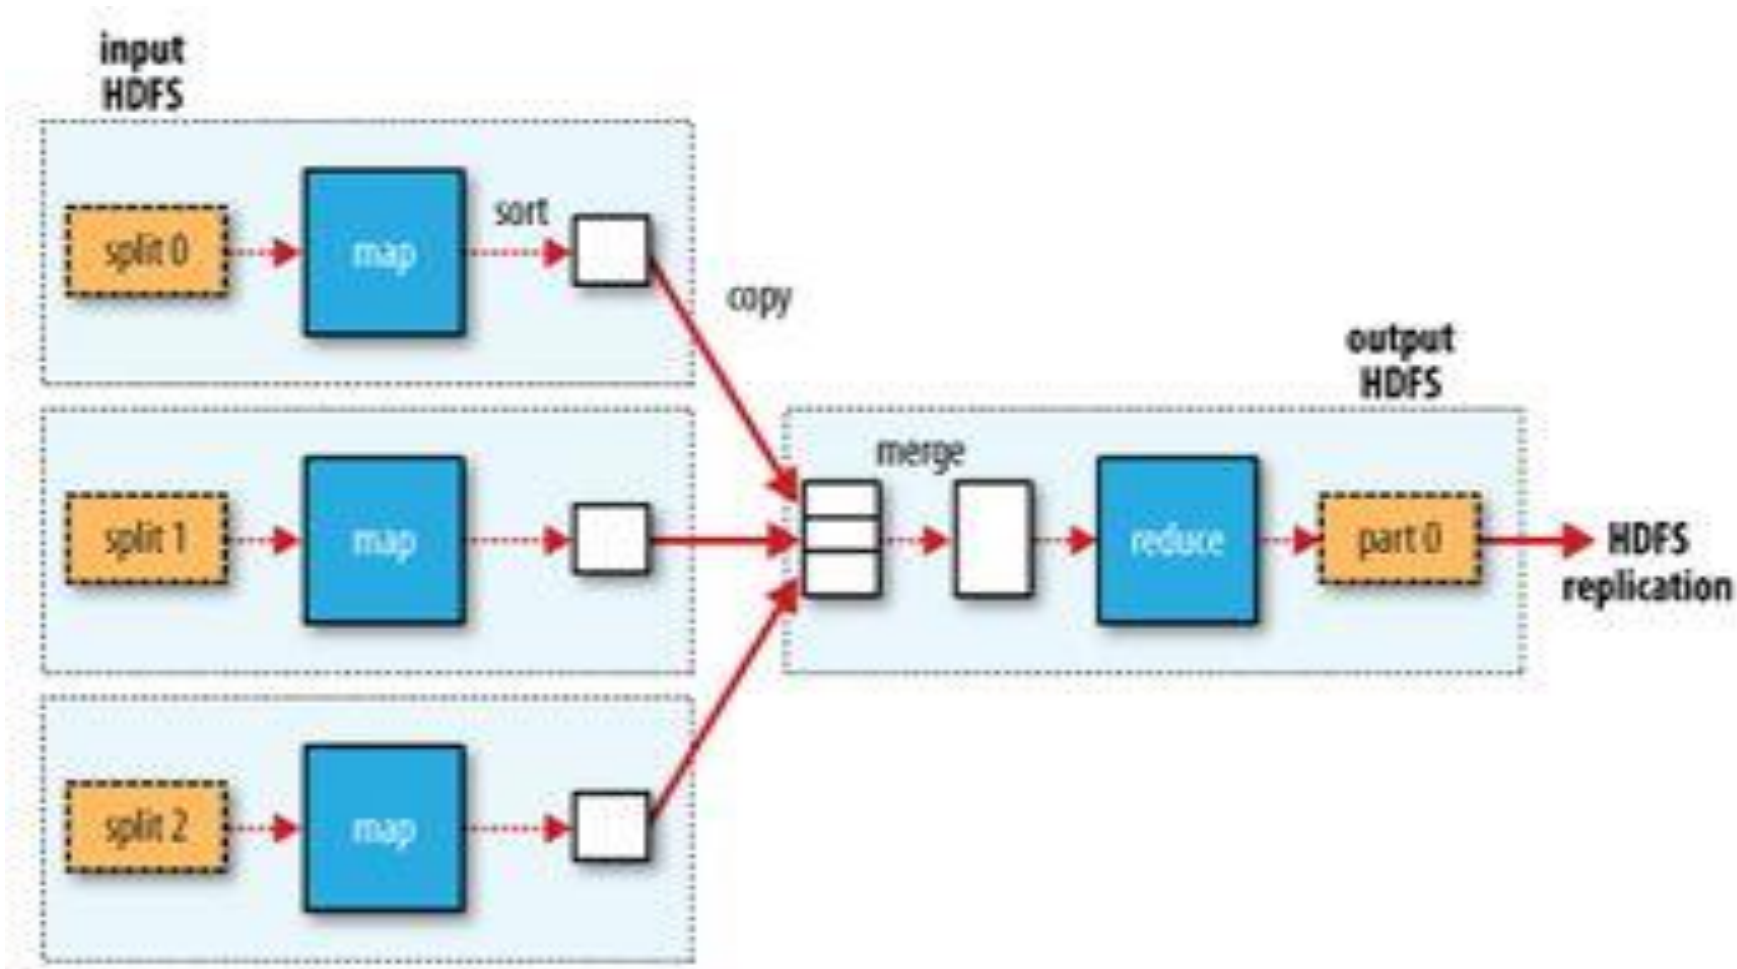


Figure: Map reduce data flow for single reducer task.

**Map Reduce:**

**Exercise : Job Submission Flow with two Reducers**

# BIG DATA

## MapReduce Flow

- How will it work when there are two reducers?
- Where will the outputs be?

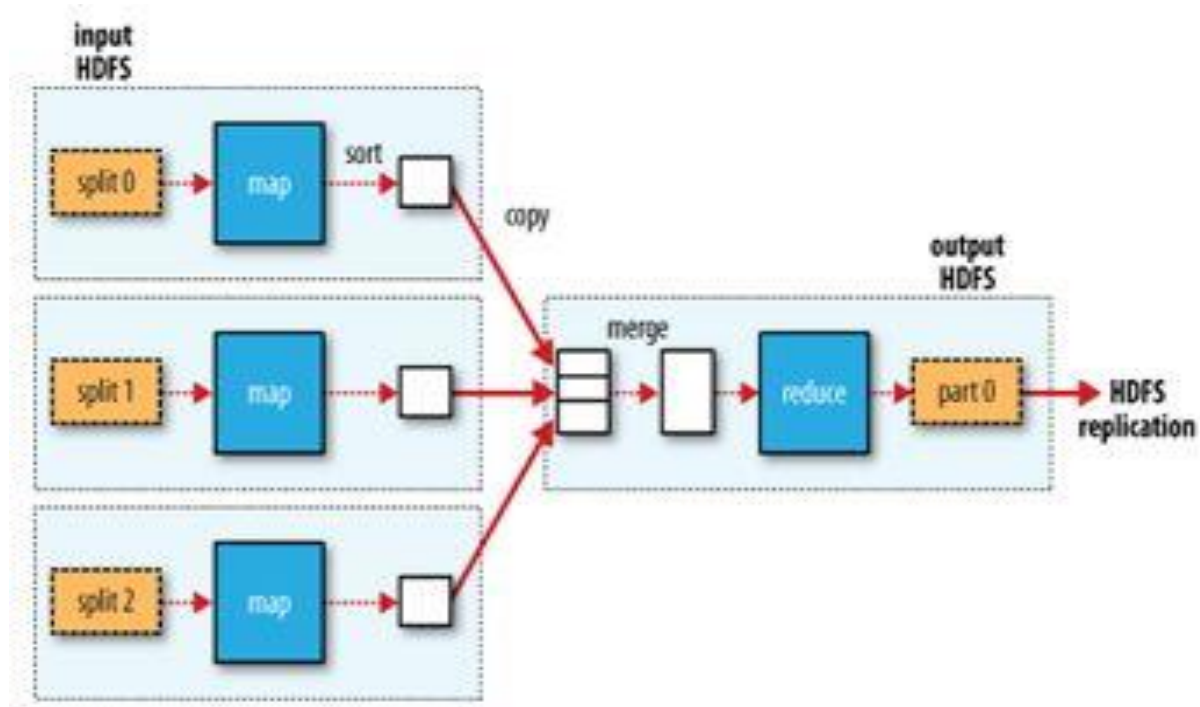



Figure: Map reduce data flow for single reducer task.

# BIG DATA

## MapReduce Flow for **MULTIPLE REDUCE** Tasks



- Outputs are created on two different nodes and have to be merged.
- But available through HDFS on any node



## Map Reduce: Splits

# BIG DATA

## Map Reduce Split size considerations

---



- Split size proportional to parallelism
- Small split size
  - Advantages
    - Large #splits
    - Increased parallelism
    - Increased load balancing
  - Disadvantages
    - the overhead of managing the splits and of map task creation
    - begins to dominate the total job execution time.

Optimal split size == size of HDFS block (128MB)  
(default) on Hadoop v2

### Split == Block size

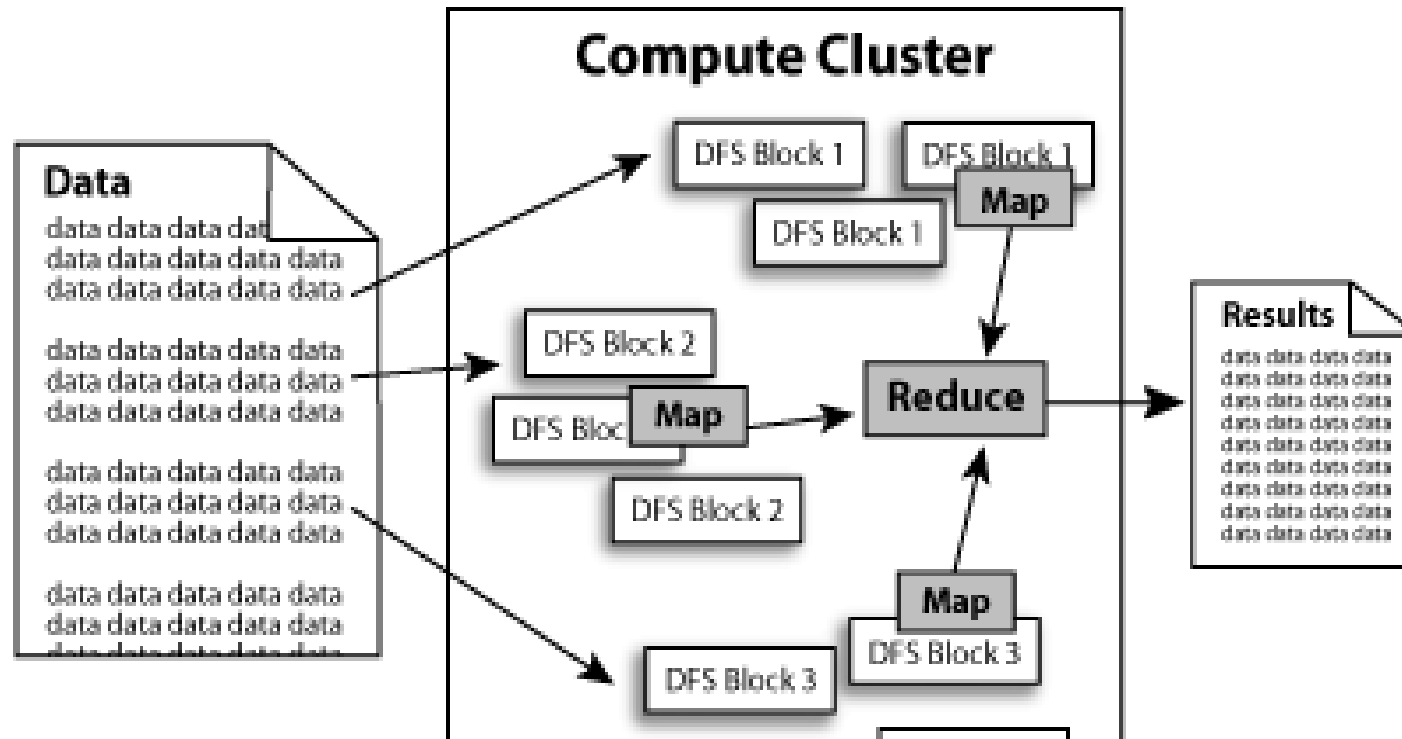
- All data required for Map
  - In the same node
  - No inter-node data transfer is required

### Split != block size

- Data transfer across multiple nodes
- Impacts performance

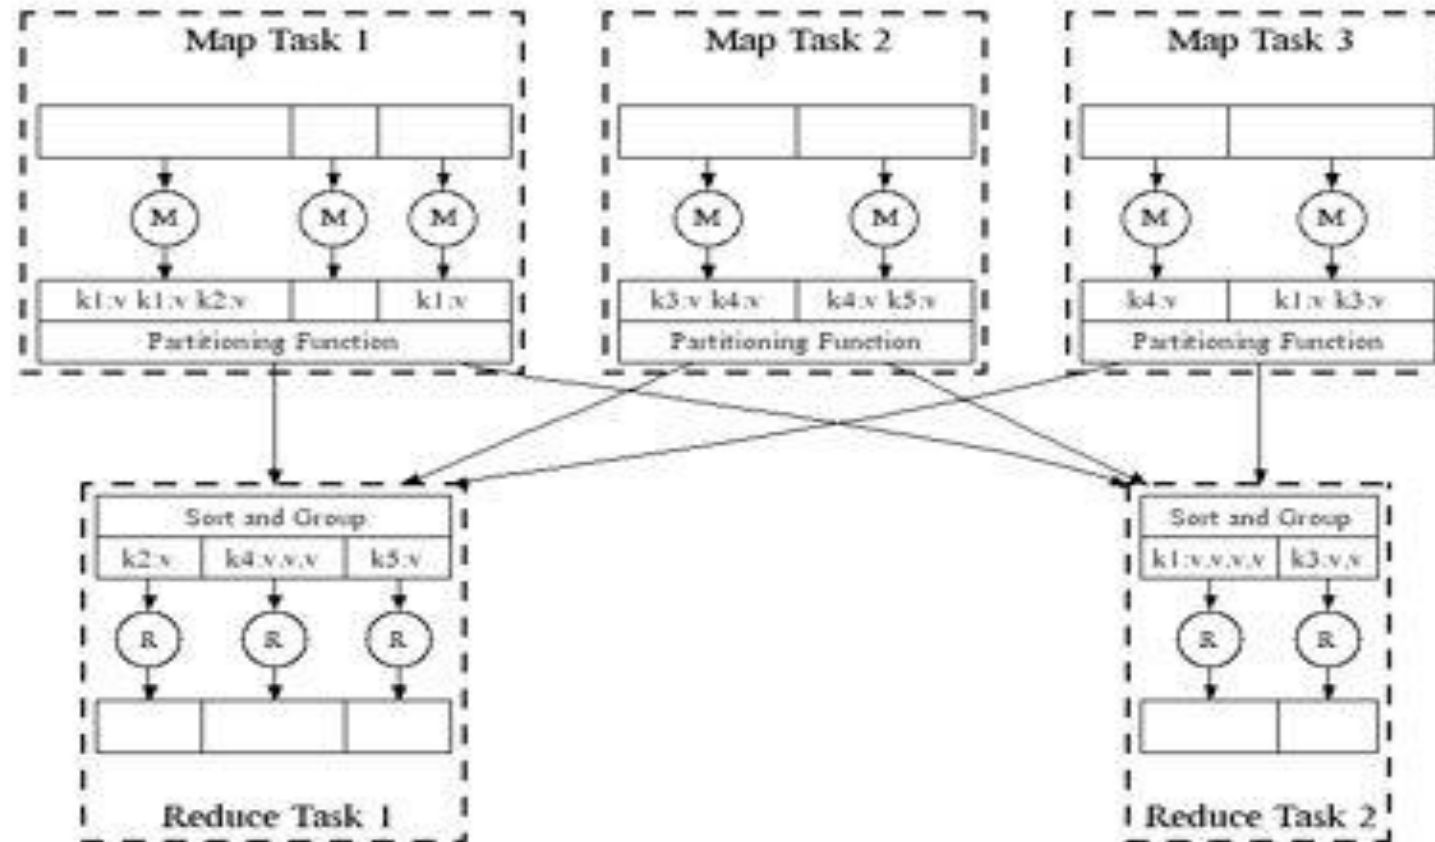
# BIG DATA

## Traditional: Move Data to Compute



# BIG DATA

## Big Data: Move Compute to Data



Parallel Execution

- **Where is Map output written to?**
  - Local disk and not HDFS
  - Why? Temporary output to be discarded after reduce.
- **Failure**
  - If the node running the map task fails
    - before the output has been consumed by reducer
  - Automatically rerun map task on another node

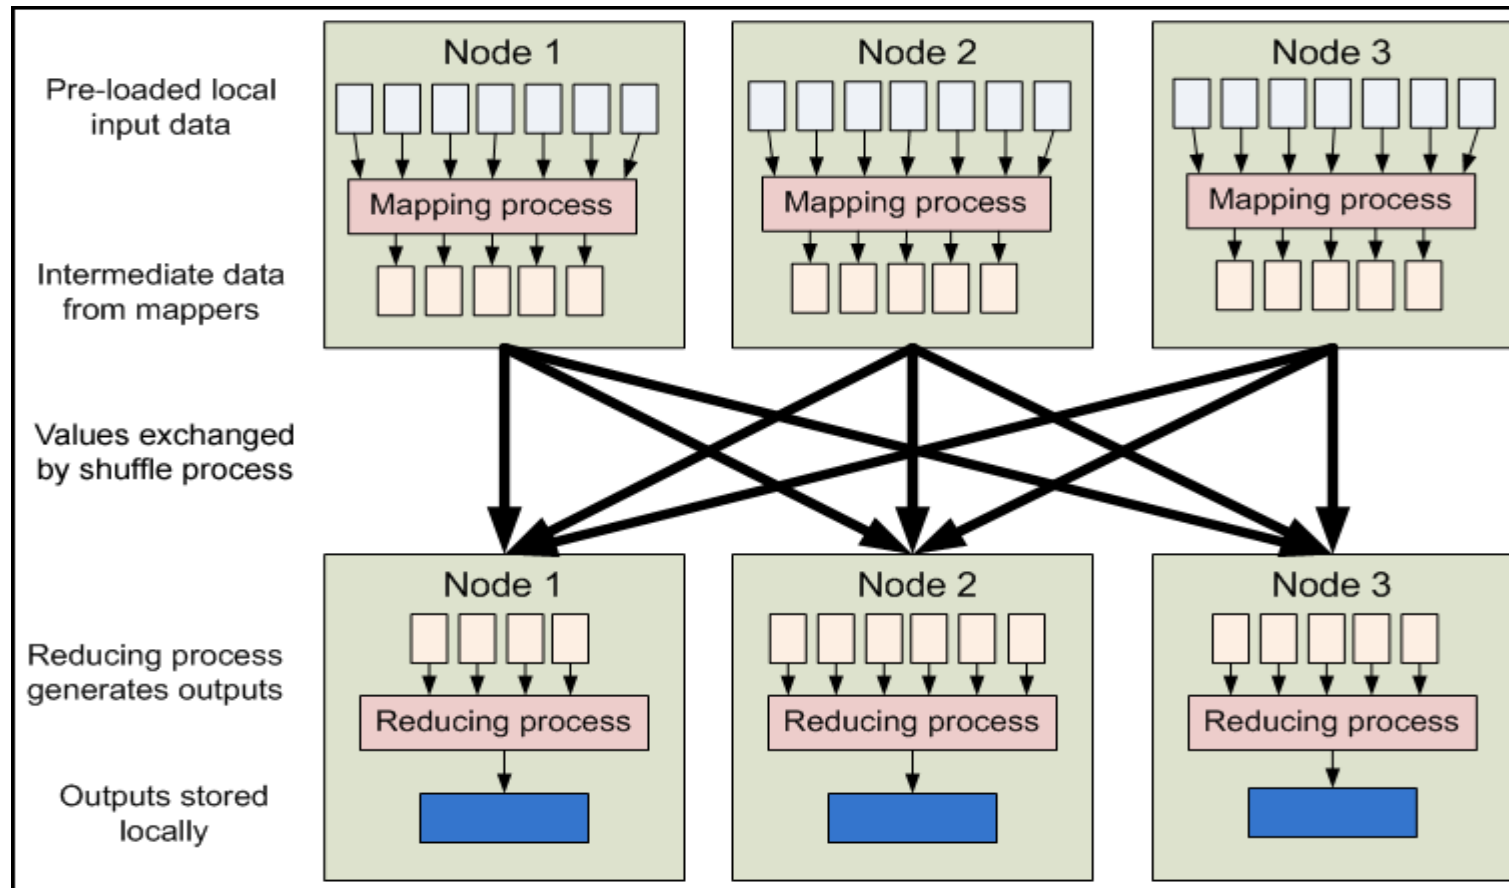
- Reduce tasks don't have the advantage of data locality
  - the input to a single reduce task is normally the output from all mappers.
- Sorted map outputs
  - have to be transferred across the network
  - Where to?
    - To the node where the reduce task is running
    - Merge data from different mappers
    - Then passed to the user-defined reduce function.
- The output of the reduce is normally stored in HDFS for reliability.
- Where is the reduce output stored
  - 1 on the local node where the reduce happens
  - Other replicas on off-rack nodes.
  - Consumes network bandwidth

## Map Reduce: Working



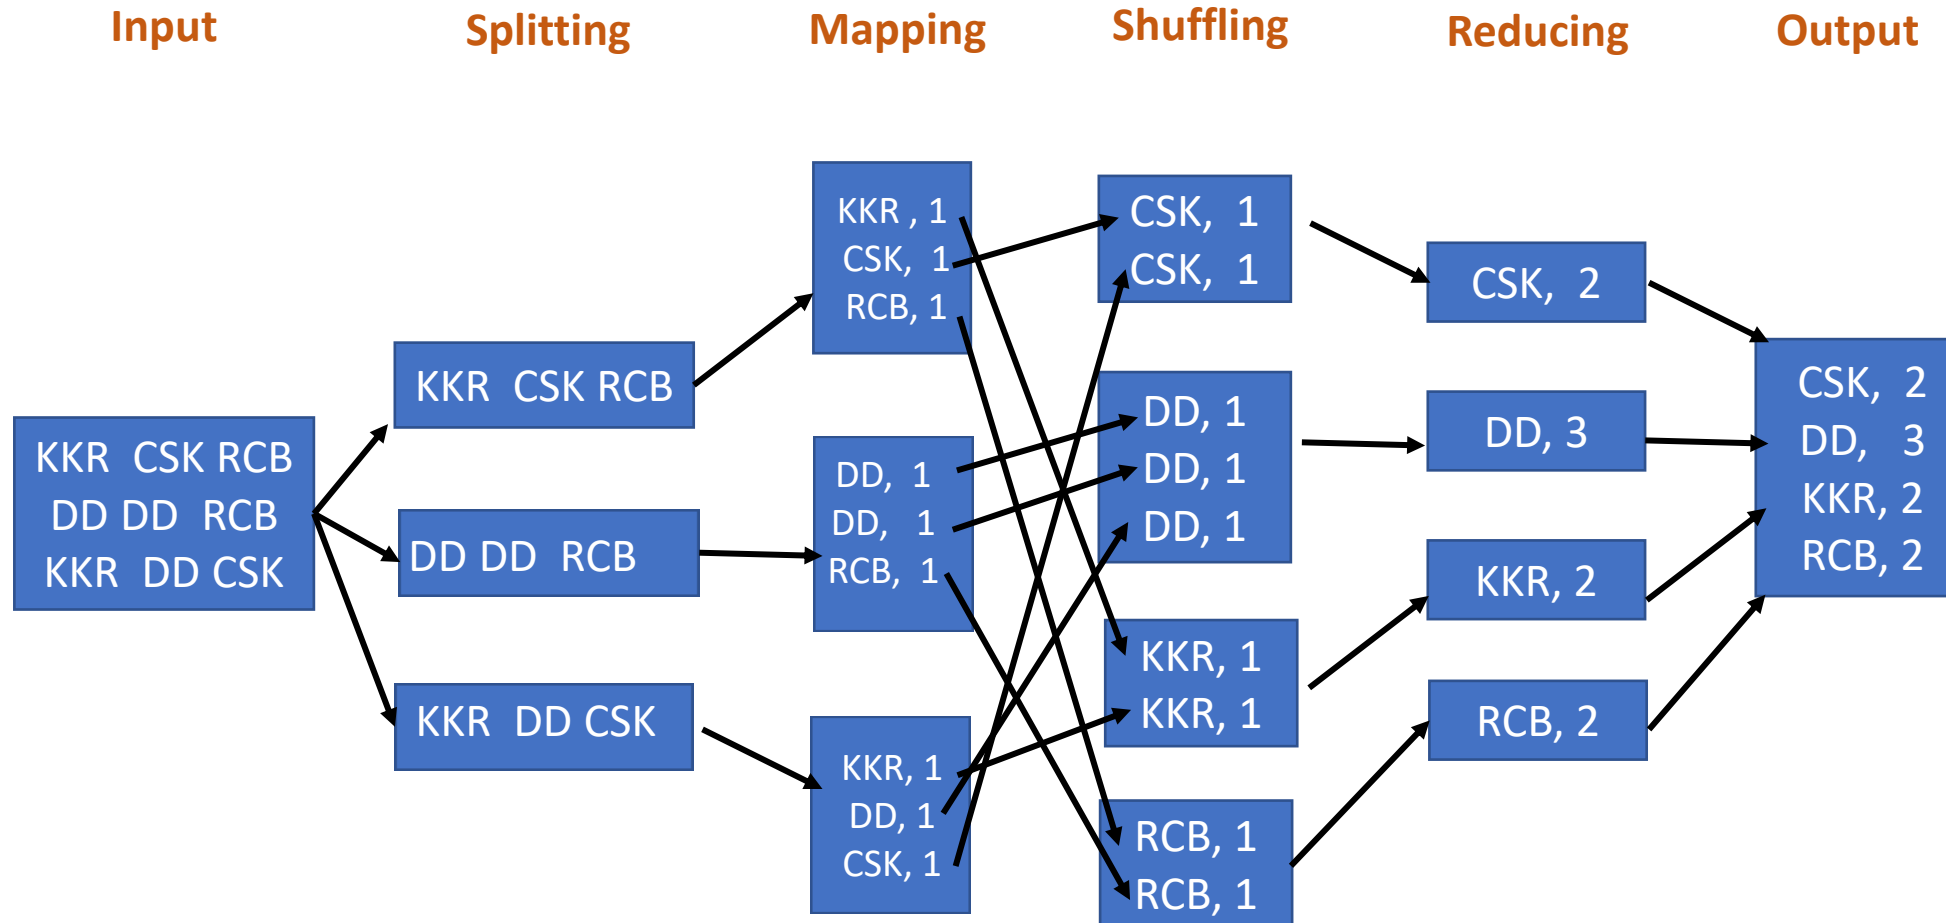
# BIG DATA

## High level view

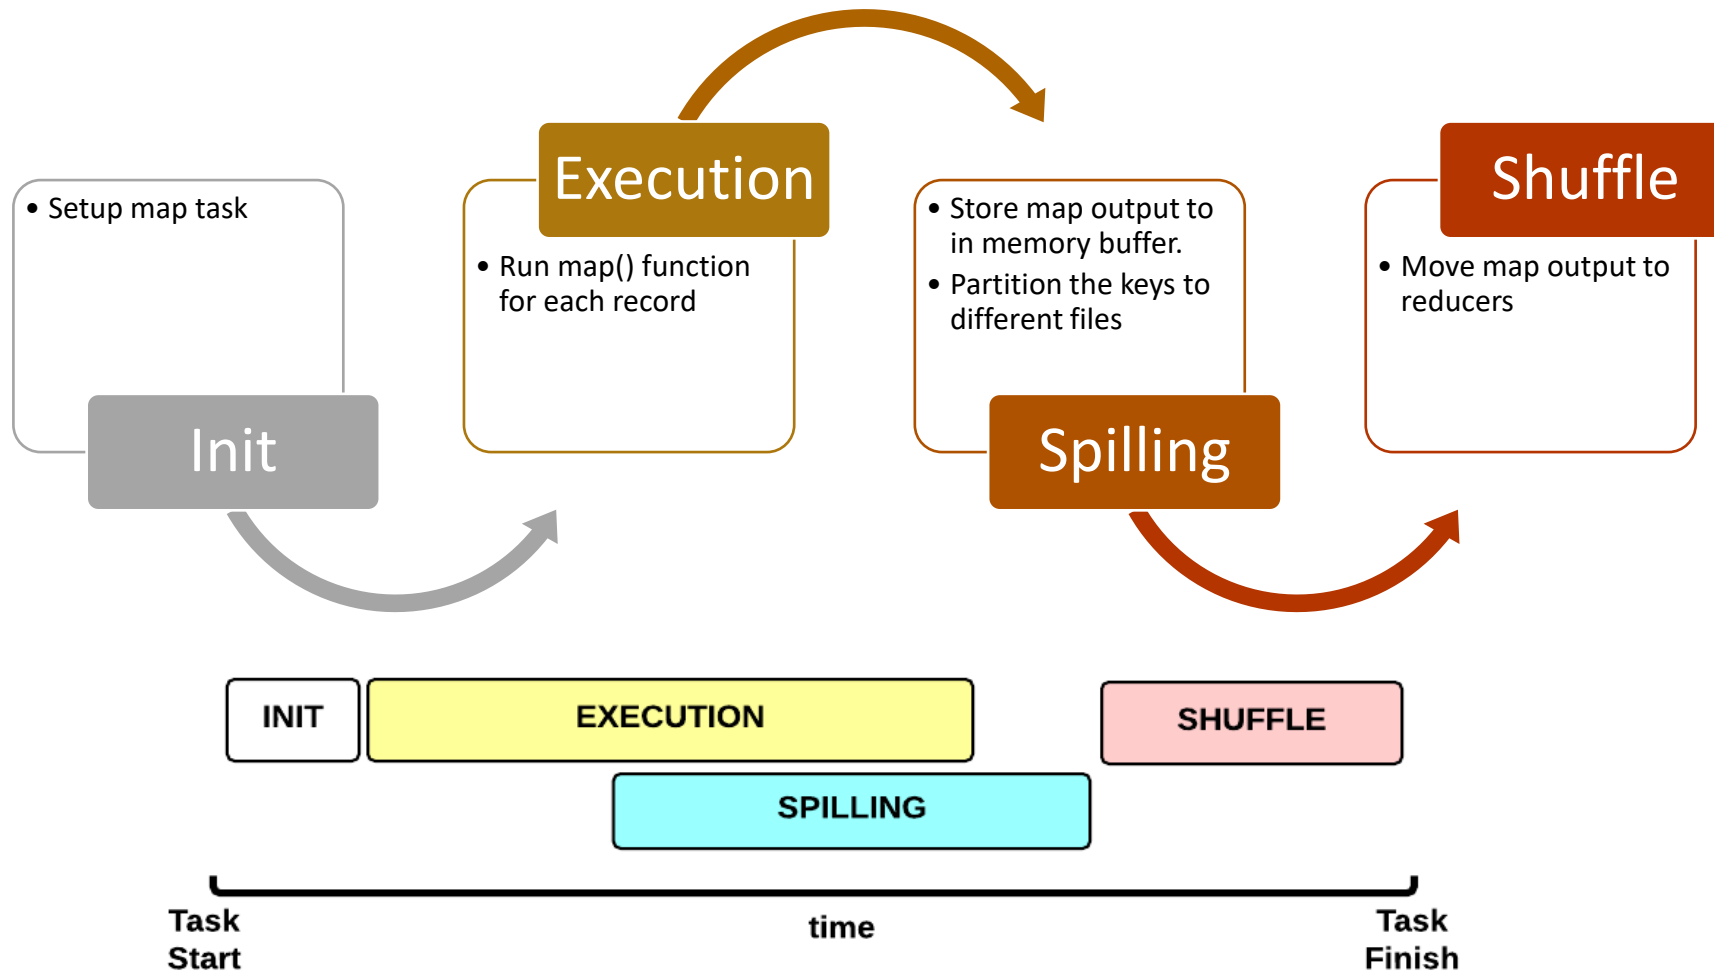


# BIG DATA

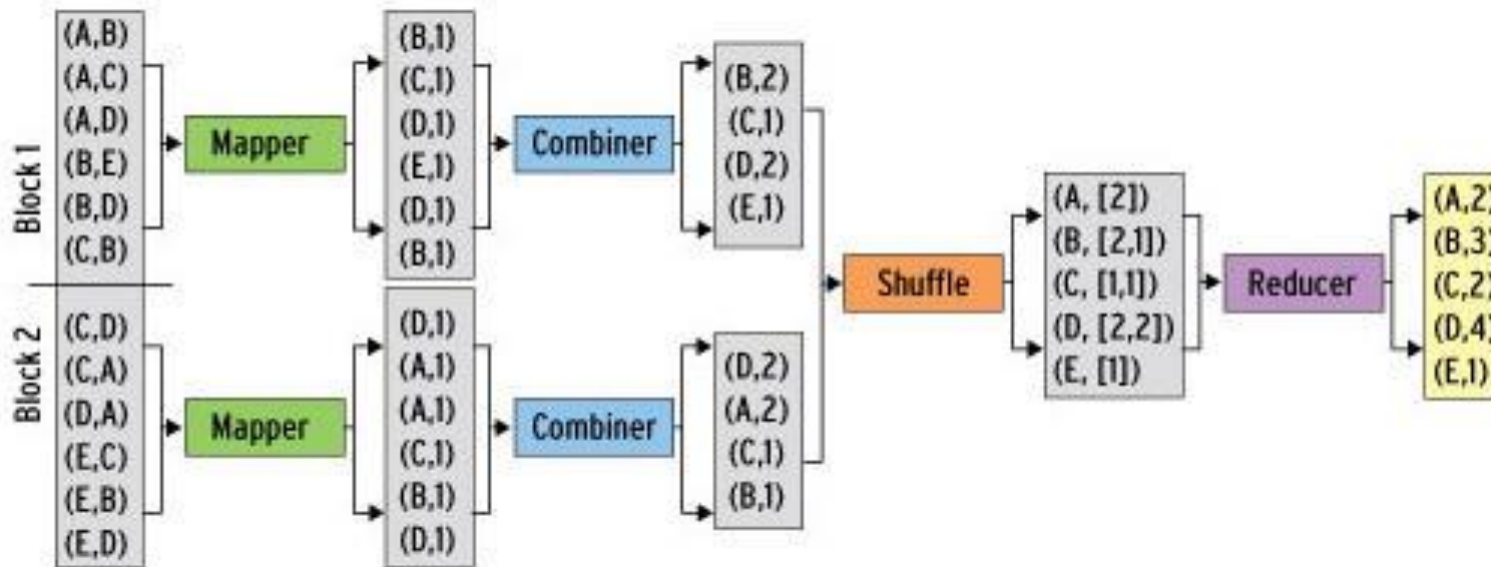
## Shuffling - example



The overall Map-Reduce word count Process



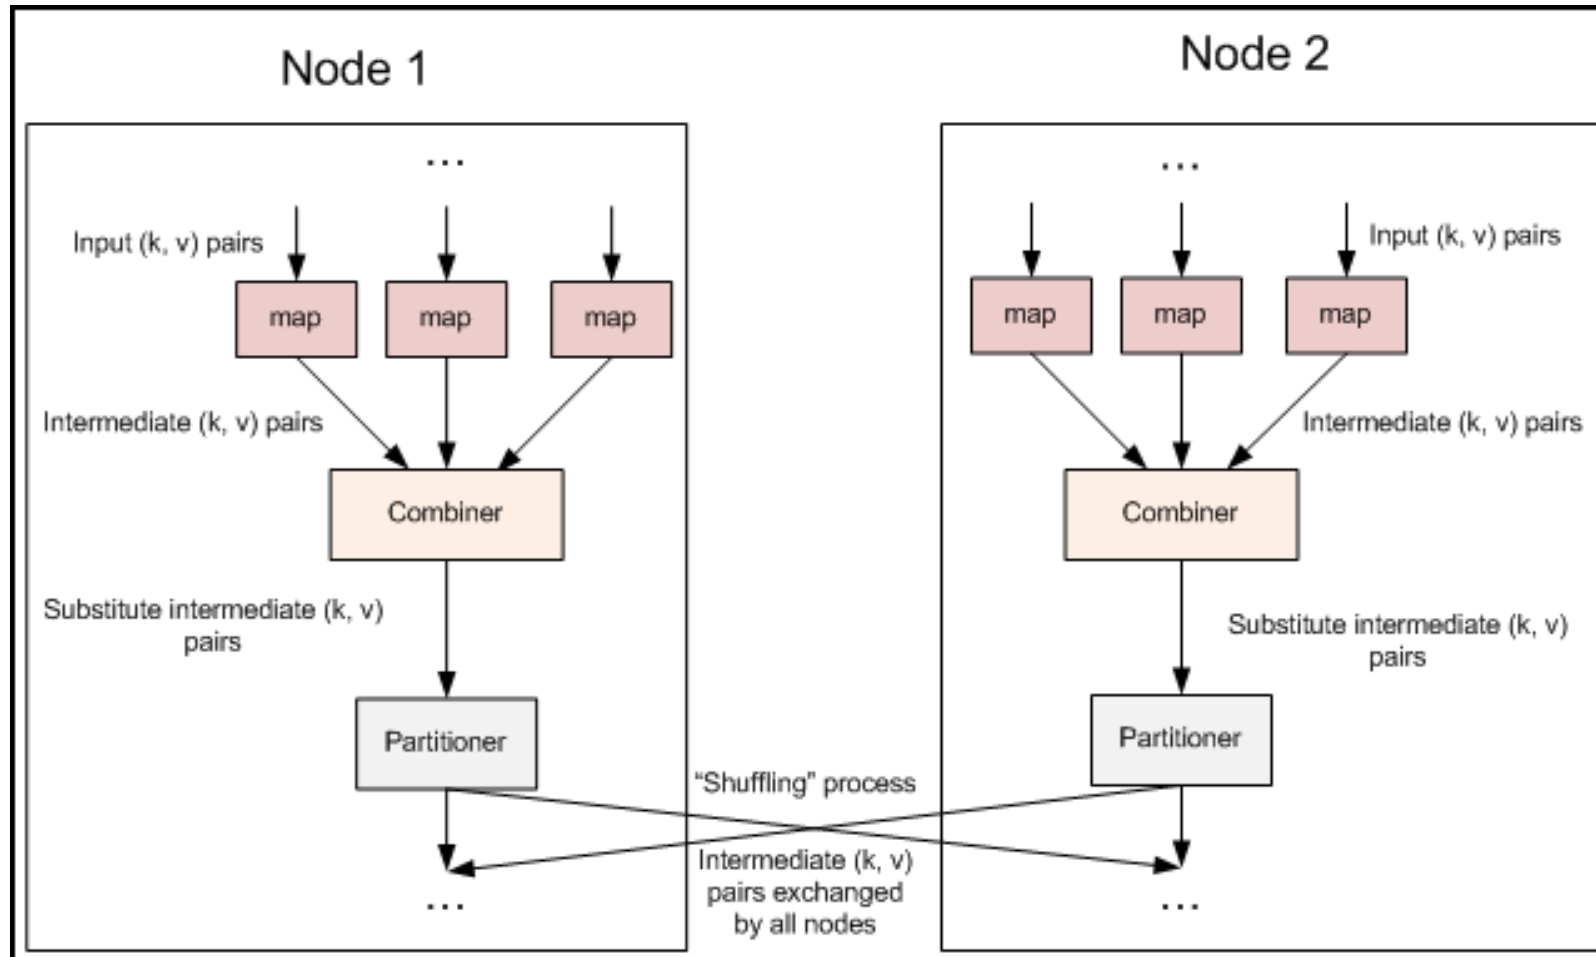
## Map Reduce: Combiners



- Combine multiple map outputs before doing a reduce
- Can write a combiner function in program
  - Combiner will be run before reduce
- Mini-reducer

# BIG DATA

## Combiner – when does it run?



```
public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    String[] otherArgs = new GenericOptionsParser(conf, args).
        getRemainingArgs();
    if (otherArgs.length < 2) {
        System.err.println("Usage: wordcount <in> [<in>...] <out>");
        System.exit(2);
    }
    Job job = new Job(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    for (int i = 0; i < otherArgs.length - 1; ++i) {
        FileInputFormat.addInputPath(job, new Path(otherArgs[i]));
    }
    FileOutputFormat.setOutputPath(job,
        new Path(otherArgs[otherArgs.length - 1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
```

**Combiner is  
set here**



## Review Questions



# BIG DATA

## Review Questions

---



- Questions from T1, LOR 2.4

- How many mappers and reducers will get started when trying to process a 230 MB file with Hadoop v2?
  - Ans: Block size = 128MB, so there will be two blocks. Assuming one block per split there will be 2 mappers
    - #reducers is configurable.
- Where is a combiner executed?
  - On the mapper.
- Write mappers and reducer pseudo code showing keys for counting #unique words in a file?
  - Similar to word count. Just that reducer does not have to write the count.

## **Additional Notes, References and Videos**

- Chapter 2.4 from T1 – Rajkamal
- Tom whites book is an excellent reference for the programming component.



# THANK YOU

---

**K V Subramaniam**

Dept. of Computer Science and Engineering

**[subramaniamkv@pes.edu](mailto:subramaniamkv@pes.edu)**