



BIG DATA

Columnar Databases for Analytics

K V Subramaniam

Computer Science and Engineering

Hadoop Available Storage Types

BIG DATA

HDFS and it's limitations



HDFS is good for batch processing (scans over big files)

Limitations

- Not good for record lookup
- Not good for incremental addition of small batches
- Not good for updates

BIG DATA

HIVE and it's limitations



- Doesn't store data
 - Uses HDFS or Hbase as actual store
- Provides an SQL Interface for querying data
- Data must be Structured: definite schema

Limitations

- Not good for record lookup
- Not good for incremental addition of small batches
- Not good for updates
- Not good for unstructured/semistructured data

BIG DATA

Hbase and Cassandra



- Built on the BigTable data model
- different in architecture

Advantages

- Fast record lookup
- Record level insertion
- Support for updates (Hbase creates new versions)
- Support for unstructured/semistructured data

Chang, Fay, et al. "Bigtable: A distributed storage system for structured data." *ACM Transactions on Computer Systems (TOCS)* 26.2 (2008): 4.

BIG DATA

Use case for different storage types



- HDFS
 - Unstructured data
 - Writes: no updates, only appends
 - Read entire file and analyze
- Hive
 - Structured data
 - Analytics via SQL
- HBase/Cassandra
 - Unstructured data
 - Arbitrary writes
 - Analytics

BIG DATA

Exercise

Which of these could be stored in HDFS, Hive or Hbase?

Parsed transaction logs of user activity in a website where relevant fields from the log have been extracted

Unparsed transaction logs of user activity

Database of users and friends at a social website, which is periodically analyzed for social networking analysis



BIG DATA

Solution

Which of these could be stored in HDFS, Hive or Hbase?

Parsed transaction logs of user activity in a website where relevant fields from the log have been extracted: HIVE

Unparsed transaction logs of user activity : HDFS

Database of users and friends at a social website, which is periodically analyzed for social networking analysis : HBASE



Columnar Storage

BIG DATA

Motivational Example – storage in DBMS

Row Key	Info:height	Info:age	School:House	School:sports
HarryPotter	4.5ft	11	Gryffindor	Quidditch
Voldemort	7ft	50	Slytherin	

- Row storage: DB is stored as a single file, one row per line
- Column storage: each column is a separate file, one value per line
- For using data, we need to perform an I/O to load data from disk
- Which method does less I/O for
 - Analyzing the relationship between age and earnings
 - Column storage
 - Adding a new row or read a row
 - Row storage

BIG DATA

History

- The first use of dbs were for transactions
 - Read a person's bank balance
 - Update bank balance
 - Row storage used since more efficient for transactions
- Column dbs
 - Became popular with Big Data systems
 - More efficient for analytics, particularly if db is large
 - To handle unstructured data



BIG DATA

Unstructured data

STRUCTURED VS. UNSTRUCTURED DATA

Structured Data

High Degree of organization, such as a relational database

Column	Value
Patient	Joe Brown
Date of Birth	02/13/1972
Date Admitted	02/05/2014

Unstructured Data

Information that is difficult to organize using traditional mechanisms

"The patient came in complaining of chest pain, shortness of breath, and lingering headaches...smokes 2 packs a day... family history of heart disease...has been experiencing similar symptoms for the past 12 hours...."

- How can the unstructured data above be stored in a structured relational db?
- How can it be stored in a unstructured db?

<https://www.slideshare.net/perficientinc/ibm-watson-content-analytics-discover-hidden-value-in-your-unstructured-data>

BIG DATA

Unstructured data

Customer id	Visit id	Date
		2-Oct 2017

Customer id	Visit id	Symptom id	Symptom
		1	Chest pain
		2	Headache

Customer id	Visit id	Info
		Date: {2-Oct 2017} Symptoms: {Chest pain, Headache}

Structured db

Unstructured db: simpler, more efficient

Hbase and Cassandra Data Model

BIG DATA

Hbase/Cassandra

- Hbase
 - Distributed column oriented database built on top of HDFS
 - Data is logically organized as rows/columns of a table
- Cassandra
 - Distributed database – peer to peer built by facebook
 - Inspired by Dynamo DB
 - Same data model as Hbase – inspired by BigTable

BIG DATA

Data Model – BigTable

- Key-Value pairs

Column Families: Each row divided into column families

Row Key

Name	Info:height	Info:age	School:House	School:sports
HarryPotter	4.5ft	11@2011	Gryffindor	Quidditch
Voldemort	7ft	50	Slytherin	

Timestamp

Each column family divided onto columns

HBase schema consists of several **Tables**

Each table consists of a set of **Column Families**

Columns are not part of the schema

HBase has **Dynamic Columns**

Because column names are encoded inside the cells

Different cells can have different columns

“School” column family has different columns in different cells



Name	Data
HarryPotter	Info:{height:"4.5ft", age: "11@2011"} School:{House:"Gryffindor", Sports:"Quidditch"}
Voldemort	Info:{height:"7ft", age: "50"} School:{House:"Syltherin", Role:"Prefect"}

BIG DATA

Data Model – BigTable

Row Key	Data
HarryPotter	Info: {height:"4.5ft", age: "11@2011"} School: {House:"Gryffindor", Sports:"Quidditch"}
Voldemort	Info: {height:"7ft", age: "50"} School: {House:"Syltherin", Role:"Prefect@1980, DarkLord@1995"}

Different types of
data into different
column families

Single column may
have different values
at different
timestamps

BIG DATA

Data Model – BigTable

Column family named “anchor”



Key

Byte array

Serves as the primary key for the table

Indexed for fast lookup

Column Family

Has a name (string)

Contains one or more related columns

Column

Belongs to one column family

Included inside the row

familyName:columnName

Column family named “Contents”

Row key	Time Stamp	Column “contents:”	Column “anchor:”	
“com.apache.www”	t12	“<html>...”		
	t11	“<html>...”		
	t10		“anchor:apache.com”	“APACHE”
“com.cnn.www”	t15		“anchor:cnn.com”	“CNN”
	t13		“anchor:my.look.ca”	“CNN.com”
	t6	“<html>...”		
	t5	“<html>...”		
	t3	“<html>...”		

BIG DATA

Data Model – BigTable

Version Number

Unique within each key

By default → System's timestamp

Data type is Long

Value (Cell)

Byte array (Hbase only)

Version number for each row

Row key	Time Stamp	Column “content s:”	Column “anchor:”	
“com.apac he.ww w”	t12	“<html> ...”		
	t11	“<html> ...”		
	t10		“anchor:apache .com”	“APACH E”
“com.cnn.w ww”	t15		“anchor:cnnsi.co m”	“CNN”
	t13		“anchor:my.look. ca”	“CNN.co m”
	t6	“<html> ...”		
	t5	“<html> ...”		
	t3	“<html> ...”		

value

Hbase Architecture

BIG DATA

Hbase Architecture – Master Slave

Region

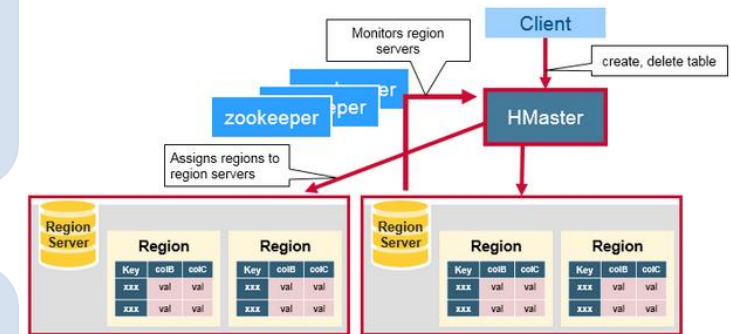
- A subset of a table's rows, like horizontal range partitioning
- Automatically done

RegionServer (many slaves)

- Manages data regions
- Serves data for reads and writes (*using a log*)
- Like datanode of HDFS

Master

- Responsible for coordinating the slaves
- Assigns regions, detects failures
- Admin functions
- Line namenode of HDFS



BIG DATA

Regions and region servers




PES
UNIVERSITY
ONLINE

Tables horizontally partitioned into regions

Regions stored on region servers

Start with a single region and then master monitors load and splits into multiple regions

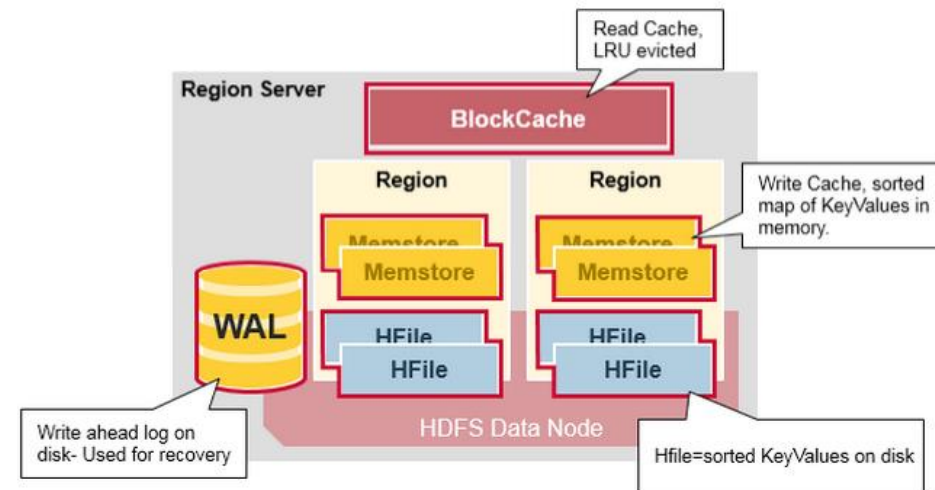
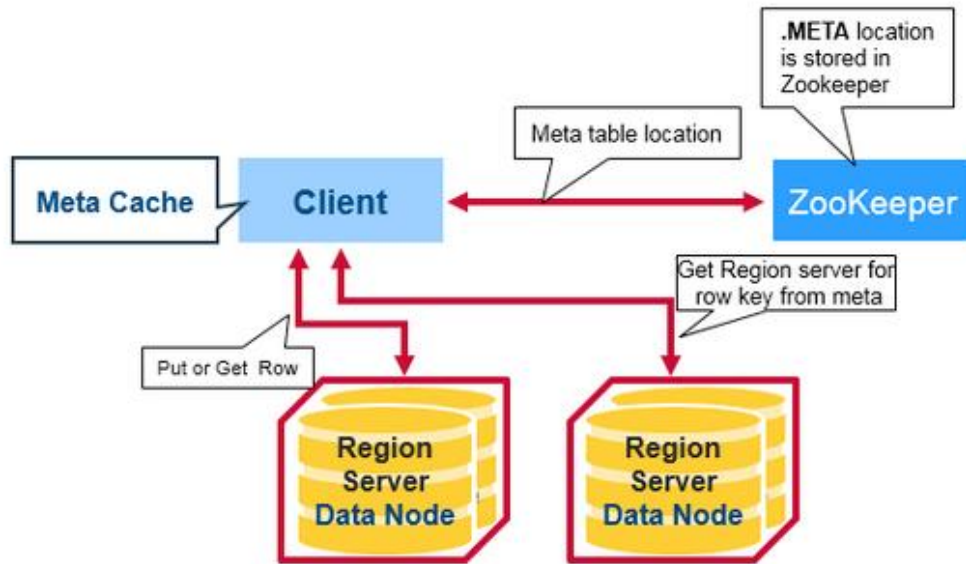
META Table:
keeps track of
regions



Refer <https://www.mapr.com/blog/in-depth-look-hbase-architecture>

BIG DATA

Read/Write Operation



Cassandra Architecture



BIG DATA

Cassandra Architecture – peer to peer architecture



- Differences from Hbase
 - Request coordination over a partitioned dataset – no Master
 - Ring membership and failure detection – no Master
 - Local persistence (storage) engine – does not rely on HDFS
- Cqlsh – for performing queries

Hbase usage



BIG DATA

Hbase: Creating a new table

```
hbase(main):001:0> create 'test', 'data'  
0 row(s) in 0.9810 seconds
```

Column
family name

Table name

BIG DATA

Hbase inserting data

```
hbase(main):003:0> put 'test', 'row1', 'data:1', 'value1'  
hbase(main):004:0> put 'test', 'row2', 'data:2', 'value2'  
hbase(main):005:0> put 'test', 'row3', 'data:3', 'value3'
```

Row key

Column
name

Value

BIG DATA

Hbase: retrieving data

Getting a
specific row

```
hbase(main):006:0> get 'test', 'row1'
COLUMN                                CELL
data:1                                timestamp=1414927084811, value=value1
1 row(s) in 0.0240 seconds
```

```
hbase(main):007:0> scan 'test'
ROW                                COLUMN+CELL
row1                                column=data:1, timestamp=1414927084811, value=value1
```

All rows



THANK YOU

K V Subramaniam

Dept. of Computer Science and Engineering

subramaniamkv@pes.edu