



BIG DATA

In memory analytics with Spark : Introduction

K V Subramaniam

Computer Science and Engineering

BIG DATA

Overview of lecture – Spark Introduction



- Why Spark – the motivation?
- Moving to in memory compute
- Distribute data in memory
- Handling fault tolerance
- Programming model – Operations in Spark
- Handling key-value based operations
- Putting it all together : Word count in Spark

Motivation for Spark

BIG DATA


Spark: Motivation

Most current cluster programming models are based on *acyclic data flow* from stable storage to stable storage

Hadoop → reads data from persistent storage in *Map* step


Writes data back to persistent store (HDFS) in reduce

Advantage – dynamically decide machines/handle failures



Consider the page rank exercise we did in the last unit?


What are the issues that we see with this?



BIG DATA

Challenges

- Acyclic data flow
 - consider operating on a *data working set*
 - **Working set:** same set of data reused
 - in page rank, we keep computing importance vector and reusing in next iteration.
 - Hadoop – inefficient in such cases.
- Example of use
 - Where we need to *iterate*
 - Graph processing
 - Machine Learning
 - Where we need to do *interactive analysis*
 - Python, R
- On every iteration, storing/reloading of data from persistent storage is time consuming.



In Memory Computation

BIG DATA

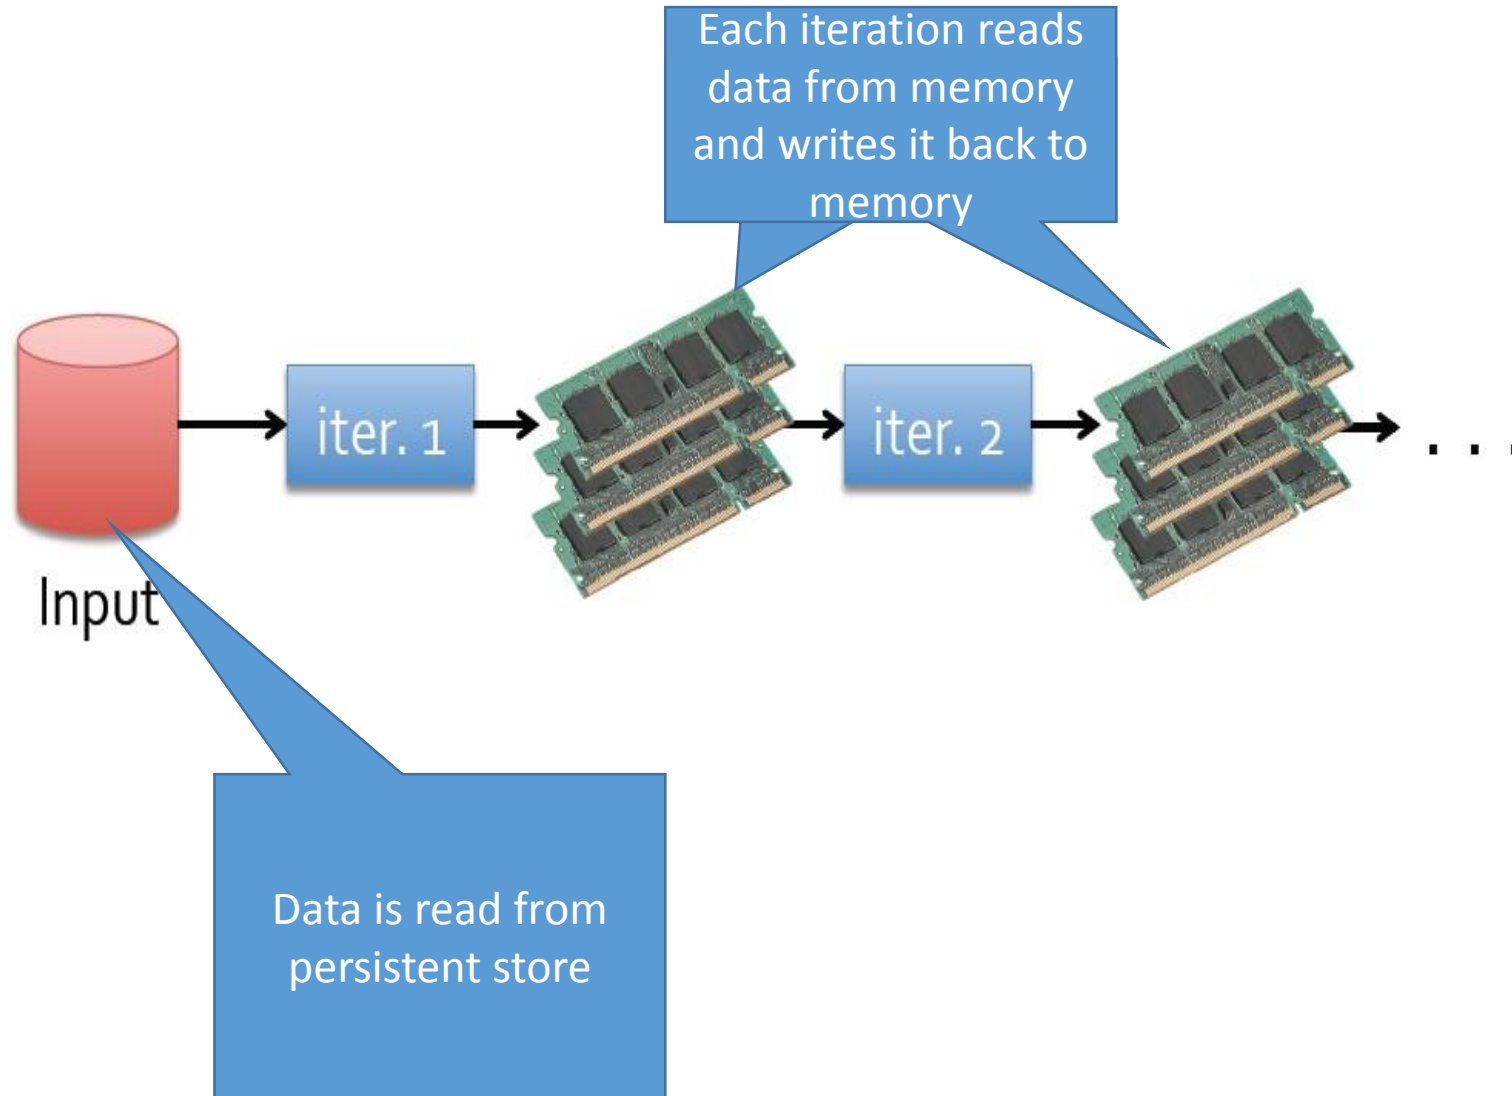
Recap: Word count in scala

```
val lines = scala.io.Source.fromFile("textfile.txt").getLines
val words = lines.flatMap(line => line.split(" ")).toIterable
val counts = words.groupBy(identity).map(words =>
    words._1 -> words._2.size)
val top10 = counts.toArray.sortBy(_._2).reverse.take(10)
println(top10.mkString("\n"))
```

Each operation creates
A data value that can be kept in
Memory and reused.


BIG DATA

Doing in memory processing – Iterative processing



BIG DATA

Doing in memory processing – interactive processing



BIG DATA

Challenges of in memory processing

- How do we distribute the data among the DRAM of the cluster?
- What happens if this memory is not sufficient?
- How do we handle failures because memory is volatile?

Distributed Dataset

BIG DATA

Example: Log Processing



Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(startsWithERROR())
messages = errors.map(split("\t"),2)
cachedMsgs = messages.cache()

cachedMsgs.filter(containsfoo()).count()
cachedMsgs.filter(containsbar()).count()
. . .
```

BIG DATA

Example: Log Processing

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(startsWithERROR())
messages = errors.map(split("\t"), 2)
cachedMsgs = messages.cache()

cachedMsgs.filter(containsfoo()).count()
cachedMsgs.filter(containsbar()).count()
. . .
```

Loads a file to an in memory struct called an RDD (think of it as a collection of strings)

Filter function to retain only those lines with an error.
Creates another RDD

Applies a function to each element(string) in RDD and produces a new RDD

Keep it in memory as it will be reused

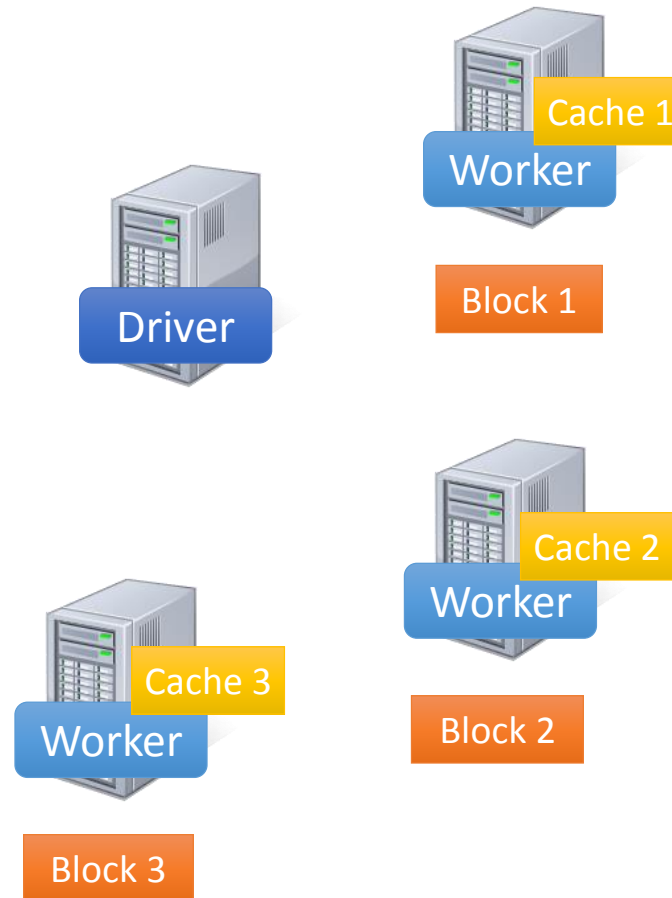
Counts #objects in the RDD

BIG DATA

Java and Scala: Spot the differences

```
lines = spark.textFile("hdfs://...")  
errors = lines.filter(startsWithERROR())  
messages = errors.map(split("\t"),2)  
cachedMsgs = messages.cache()
```

```
cachedMsgs.filter(containsfoo()).count()  
cachedMsgs.filter(containsbar()).count()
```



Adding fault tolerance – The RDD

Consider the following code:

```
Step1      Step2
messages = textFile(...).filter(startsWithERROR())
                                .map(split("\t")(2))
                                Step3
```



Step1: Read
in the file to
an in memory
RDD

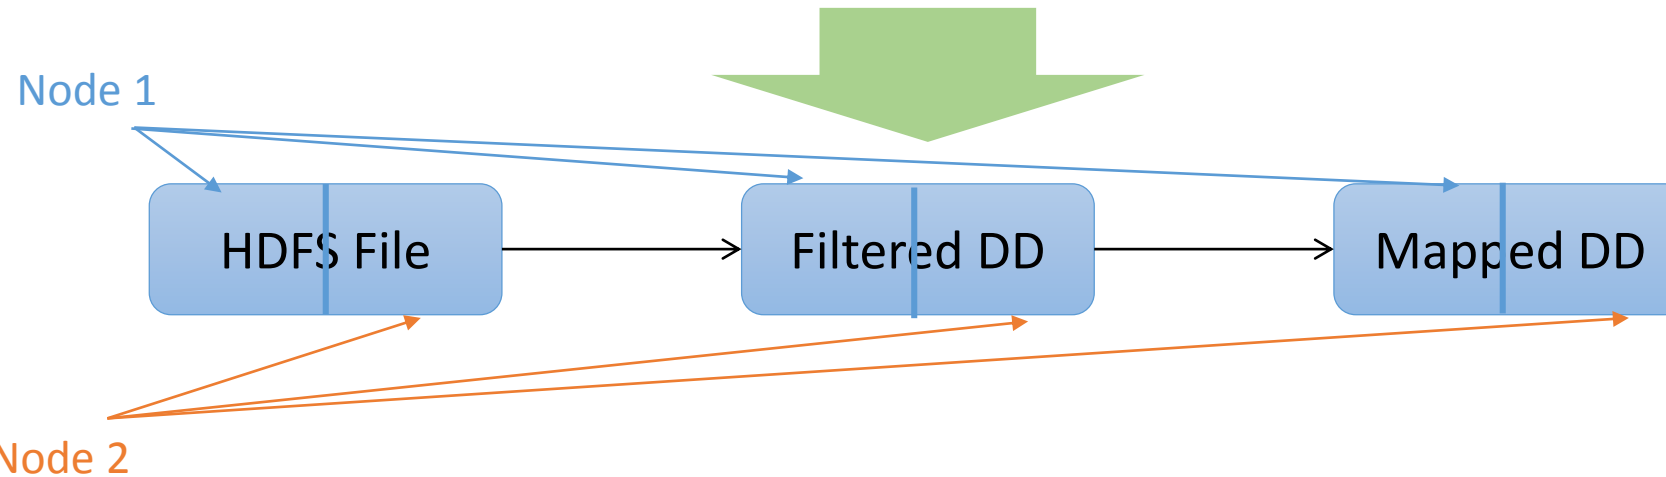
Step2: remove all
lines that don't
contain the term
ERROR

Step3: split the
line

map(func)	Return a new distributed dataset formed by passing each element of the source through a function <i>func</i> .
filter(func)	Return a new dataset formed by selecting those elements of the source on which <i>func</i> returns true.

Ex:

```
messages = textFile(...).filter(startsWithERROR())  
                        .map(split("\t")(2))
```



BIG DATA

What is an RDD?

- When we add lineage information to the concept of a Distributed Dataset
 - We add ability to recreate it in case of failure
 - So, this data is now *resilient* to failures.
 - Hence called an ***RDD: Resilient Distributed Dataset***

BIG DATA

How is lineage information stored

- Lineage information is stored by keeping track of
 - Operations that are performed on
 - An RDD
 - That results in another RDD
 - What types of operations are supported?



RDD Operations : Transformations and Actions

BIG DATA

Types of Operations

- Operations are of two types
 - *Transformations*
 - *Actions*

- Are operations that create a new dataset from an existing dataset
- For example:
 - *map()* is a transformation
 - Each line on input RDD is passed through the *map()* function
 - result of *map()* function applied on each value is stored in the output RDD.
- Note it is similar to the Map of map-reduce, but is more generic.

BIG DATA

Transformations

Table 3-2. Basic RDD transformations on an RDD containing {1, 2, 3, 3}

Function name	Purpose	Example	Result
<code>map()</code>	Apply a function to each element in the RDD and return an RDD of the result.	<code>rdd.map(x => x + 1)</code>	{2, 3, 4, 4}
<code>flatMap()</code>	Apply a function to each element in the RDD and return an RDD of the contents of the iterators returned. Often used to extract words.	<code>rdd.flatMap(x => x.to(3))</code>	{1, 2, 3, 2, 3, 3, 3}
<code>filter()</code>	Return an RDD consisting of only elements that pass the condition passed to <code>filter()</code> .	<code>rdd.filter(x => x != 1)</code>	{2, 3, 3}
<code>distinct()</code>	Remove duplicates.	<code>rdd.distinct()</code>	{1, 2, 3}
<code>sample(withReplacement, fraction, [seed])</code>	Sample an RDD, with or without replacement.	<code>rdd.sample(false, 0.5)</code>	Nondeterministic

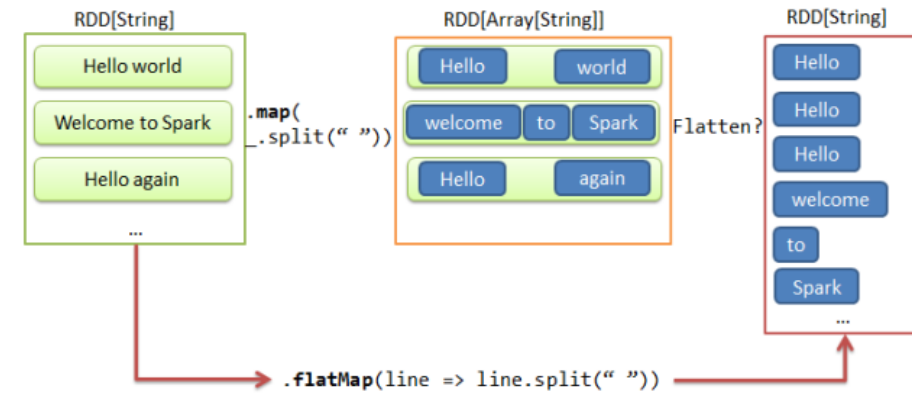


Table 3-3. Two-RDD transformations on RDDs containing {1, 2, 3} and {3, 4, 5}

Function name	Purpose	Example	Result
<code>union()</code>	Produce an RDD containing elements from both RDDs.	<code>rdd.union(other)</code>	{1, 2, 3, 3, 4, 5}
<code>intersection()</code>	RDD containing only elements found in both RDDs.	<code>rdd.intersection(other)</code>	{3}
<code>subtract()</code>	Remove the contents of one RDD (e.g., remove training data).	<code>rdd.subtract(other)</code>	{1, 2}
<code>cartesian()</code>	Cartesian product with the other RDD.	<code>rdd.cartesian(other)</code>	{(1, 3), (1, 4), ... (3,5)}

- Are operations that return a value
- For example:
 - Reduce() is an action
 - Aggregates all elements of a RDD to produce a result.

BIG DATA

Actions

Table 3-4. Basic actions on an RDD containing {1, 2, 3, 3}

Function name	Purpose	Example	Result
collect()	Return all elements from the RDD.	rdd.collect()	{1, 2, 3, 3}
count()	Number of elements in the RDD.	rdd.count()	4
countByValue()	Number of times each element occurs in the RDD.	rdd.countByValue()	{(1, 1), (2, 1), (3, 2)}
take(num)	Return num elements from the RDD.	rdd.take(2)	{1, 2}

top(num)	Return the top num elements the RDD.	rdd.top(2)	{3, 3}
takeOrdered(num)(ordering)	Return num elements based on provided ordering.	rdd.takeOrdered(2)(myOrdering)	{3, 3}
takeSample(withReplacement, num, [seed])	Return num elements at random.	rdd.takeSample(false, 1)	Nondeterministic
reduce(func)	Combine the elements of the RDD together in parallel (e.g., sum).	rdd.reduce((x, y) => x + y)	9
fold(zero)(func)	Same as reduce() but with the provided zero value.	rdd.fold(0)((x, y) => x + y)	9

RDD Operations : Working with key-value pairs

- Consider our earlier operation of map/reduce using Spark
- Worked on datasets with only single values
- Let's consider how to represent *<key, value>* pairs
- Spark provides
 - Separate RDDs called pair RDDs for this
 - Separate operations to function on Pair RDDs


Table 4-1. Transformations on one pair RDD (example: $\{(1, 2), (3, 4), (3, 6)\}$)

Function name	Purpose	Example	Result
<code>reduceByKey(func)</code>	Combine values with the same key.	<code>rdd.reduceByKey((x, y) => x + y)</code>	$\{(1, 2), (3, 10)\}$
<code>groupByKey()</code>	Group values with the same key.	<code>rdd.groupByKey()</code>	$\{(1, [2]), (3, [4, 6])\}$
<code>combineByKey(createCombiner, mergeValue, mergeCombiners, partitioner)</code>	Combine values with the same key using a different result type.	See Examples 4-12 through 4-14.	

<code>mapValues(func)</code>	Apply a function to each value of a pair RDD without changing the key.	<code>rdd.mapValues(x => x+1)</code>	<code>{(1, 3), (3, 5), (3, 7)}</code>
<code>flatMapValues(func)</code>	Apply a function that returns an iterator to each value of a pair RDD, and for each element returned, produce a key/value entry with the old key. Often used for tokenization.	<code>rdd.flatMapValues(x => (x to 5))</code>	<code>{(1, 2), (1, 3), (1, 4), (1, 5), (3, 4), (3, 5)}</code>
<code>keys()</code>	Return an RDD of just the keys.	<code>rdd.keys()</code>	<code>{1, 3}</code>
<code>values()</code>	Return an RDD of just the values.	<code>rdd.values()</code>	<code>{2, 4, 6}</code>

`countByKey(k, V)` → returns a HashMap of (k, Int) key value pairs
with count of each key

Word Count in Spark



Create a spark context: tell
Spark to create a new job

```
val sc = new SparkContext(new SparkConf().setAppName("Spark Count"))
```

Read in text file
Split it into words

```
val tokenized = sc.textFile(args(0)).flatMap(_.split(" "))
```

```
val wordCounts = tokenized.map((_, 1)).reduceByKey(_ + _)
```

Each of these
is an RDD

Reduce by key. Can also
use countByKey

Map each word to 1

- What is Apache Spark? Matei Zaharia
 - <https://www.youtube.com/watch?v=p8FGC49N-zM>
- [RDD, DataFrames and Datasets](#)
 - <https://www.youtube.com/watch?v=pZQsDloGB4w>



THANK YOU

K V Subramaniam, Usha Devi

Dept. of Computer Science and Engineering

subramaniamkv@pes.edu

ushadevibg@pes.edu