



BIG DATA

Kafka Architecture and Usage

K V Subramaniam

Computer Science and Engineering

BIG DATA

Kafka Overview



- Need for data pipelines
- Kafka architecture
- Kafka components – messaging model
- Communication and Routing
- Messages
- Scaling
- Fault Tolerance


Kafka : Need for data pipelines

BIG DATA

The need for processing events



- Stream processing requires requires processing of events
- Think of an event as something that happens at a time
- Events are processed on the server
- For example
 - “Student with ID 23489 entered building”




Data pipelines starts like this.

BIG DATA

What are Data Pipelines




- Data sources are varied and
 - Stream data at varied rates
 - Are very lightweight and not capable of compute
- Analysis is done on servers
- Data has to be transferred from sources to the compute servers
- Point to point connection is a simple way to connect



Data pipelines starts like this.

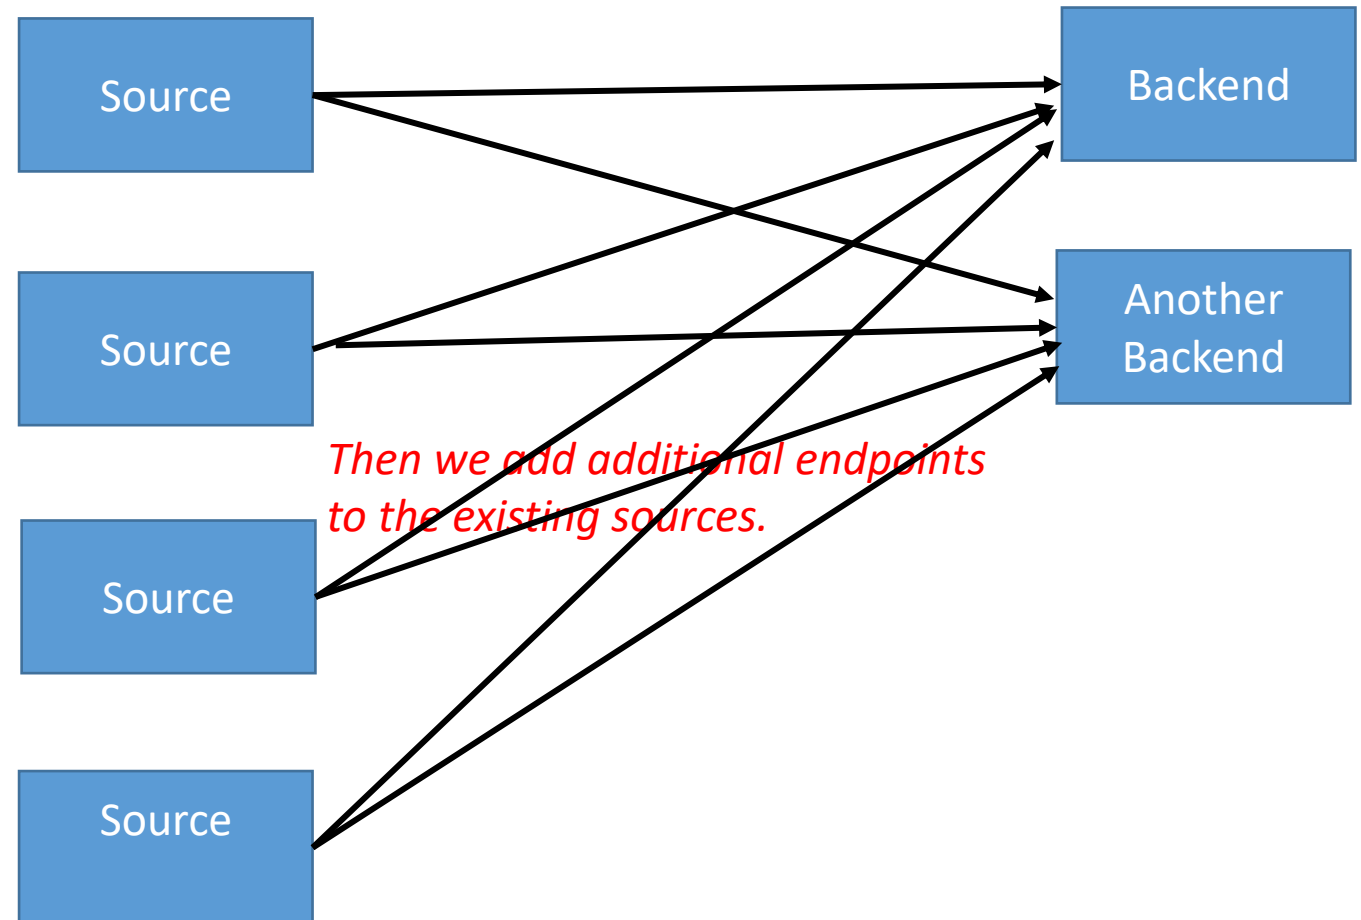
- But sources are many
- We often have to read from multiple data sources
 - E.g: Multiple cameras connected to a central processing system
- We can connect multiple clients each over a pool of connections



BIG DATA

Need for multiple backends


- But we might have multiple servers
- On which to process the same data
- Use case: stock trading – may have different analysts wanting to analyze the behavior



BIG DATA


The maze of connections

- Then it starts to look like this.



BIG DATA

The maze of connections




With may be some of these.

- As Distributed systems and services increasingly become part of modern architecture, this makes for a fragile system.

BIG DATA

Exercise 1 (5 minutes)


- Give an example of how datapipelines could be used
 - What are some examples of backends?



BIG DATA

Exercise 1 (5 minutes)

- Give an example of how datapipelines could be used
 - What are some examples of backends?

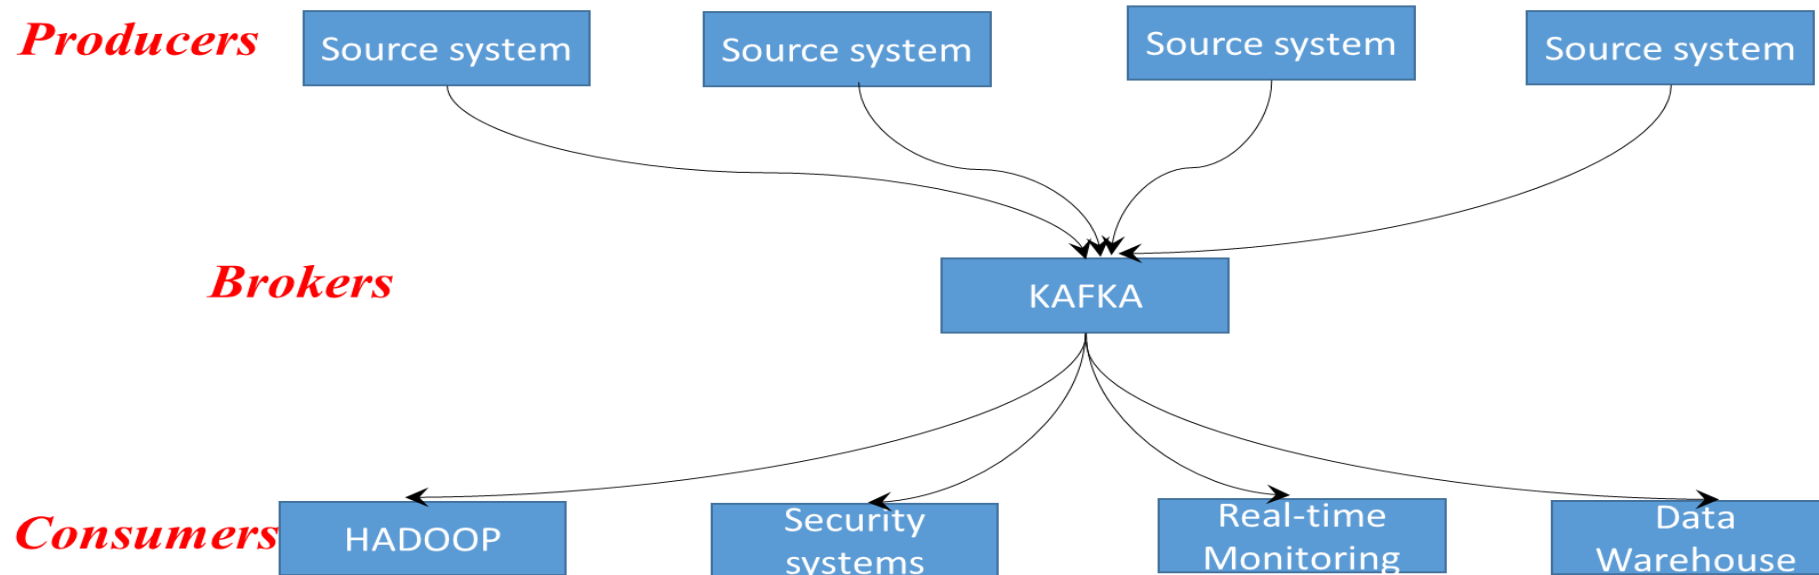


Kafka architecture

BIG DATA

Why Kafka?

- Can we have an intermediary that connects
 - Different sources
 - Multiple backends
- Decouple the pipeline so that producers and consumers do not need to know about each other



Kafka Decouples Data Pipelines

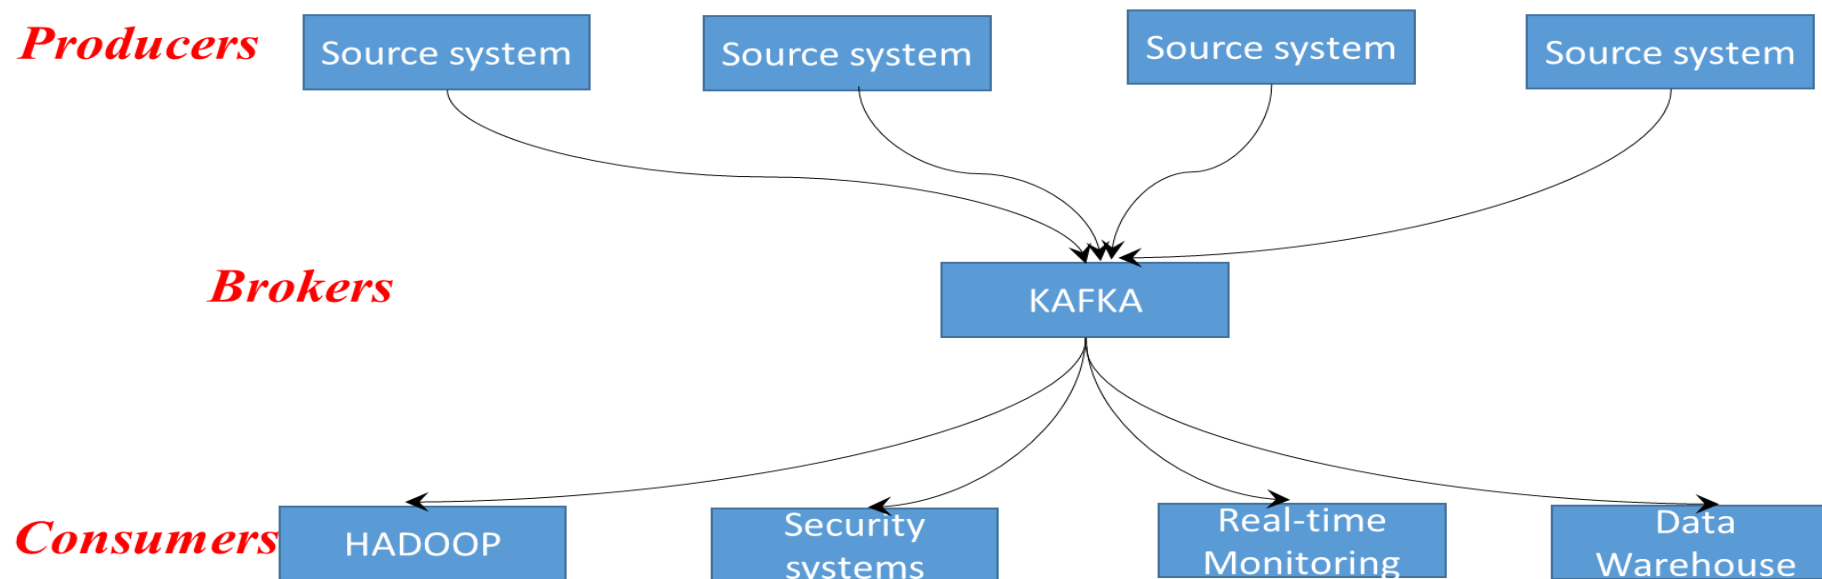
Kafka Components – Publish Subscribe Model

BIG DATA

KAFKA components

The Three essential Elements /Components:

1. Publisher
2. Subscriber
3. Communication Infrastructure



Kafka Decouples Data Pipelines

Publishers and Subscribers

So,

Okay . So

Q) What does a Publisher do ..?

**A. It publishes messages to the
Communication
Infrastructure.**

Q) What does a Subscriber do ..?

**A. It subscribes to a category of
messages.**

The Routing Mystery



*What If I tell you
I am the Publisher
and
I have no idea about
the Subscriber.*

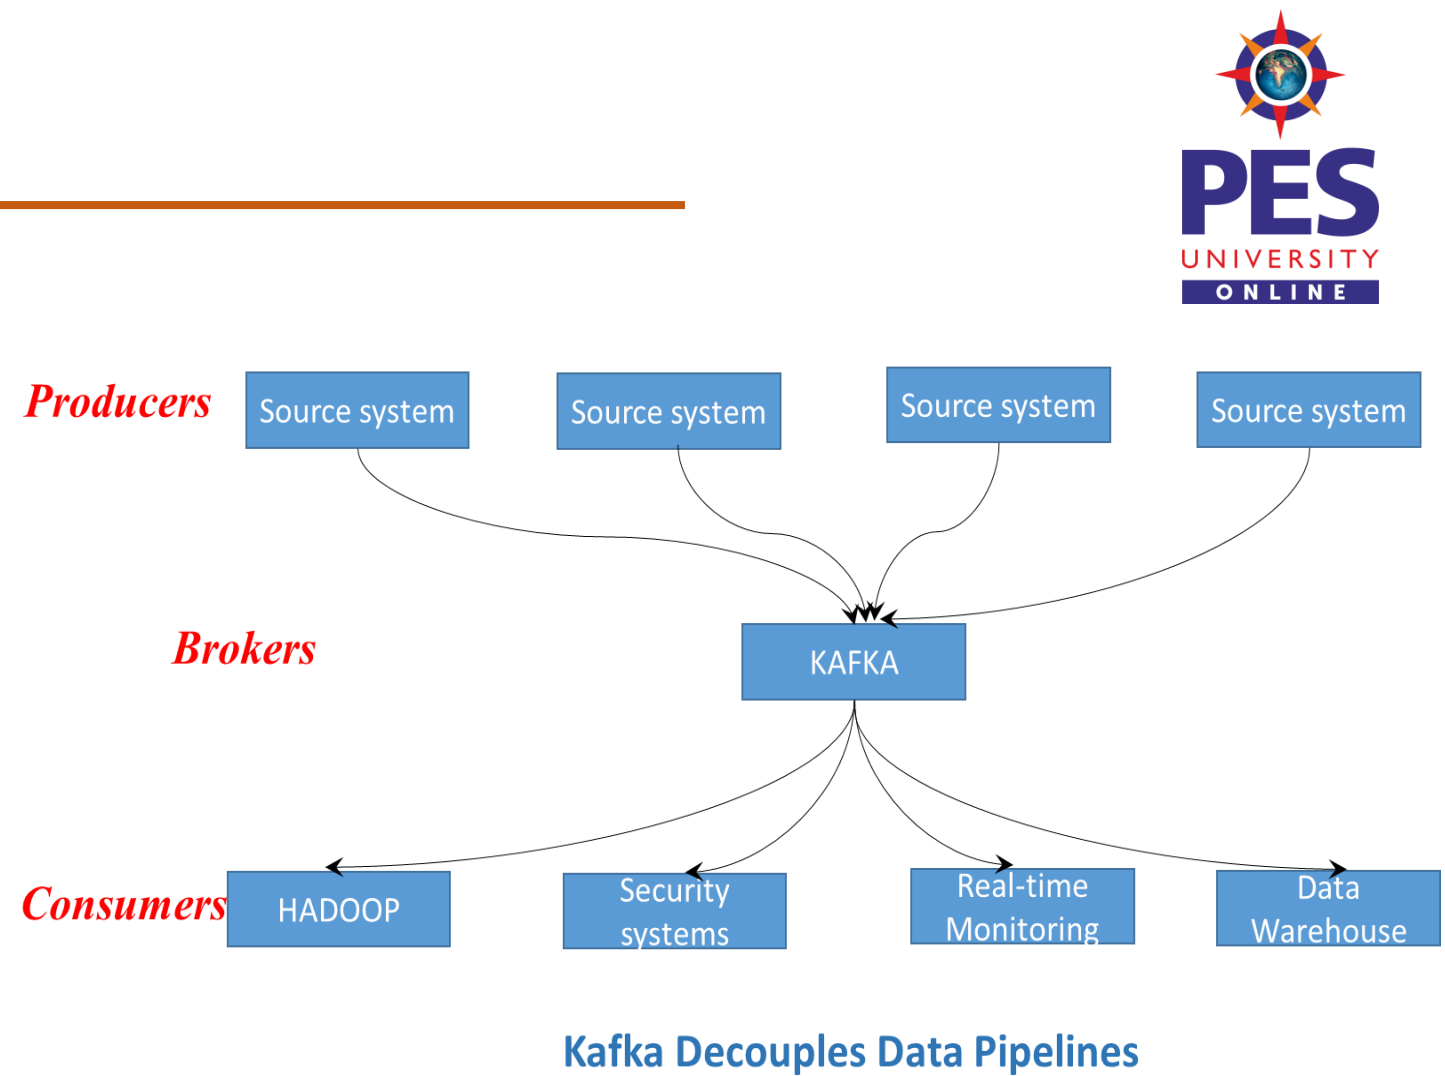


*What If I tell you
I am the Subscriber
and
I have no idea about
the Publisher.*

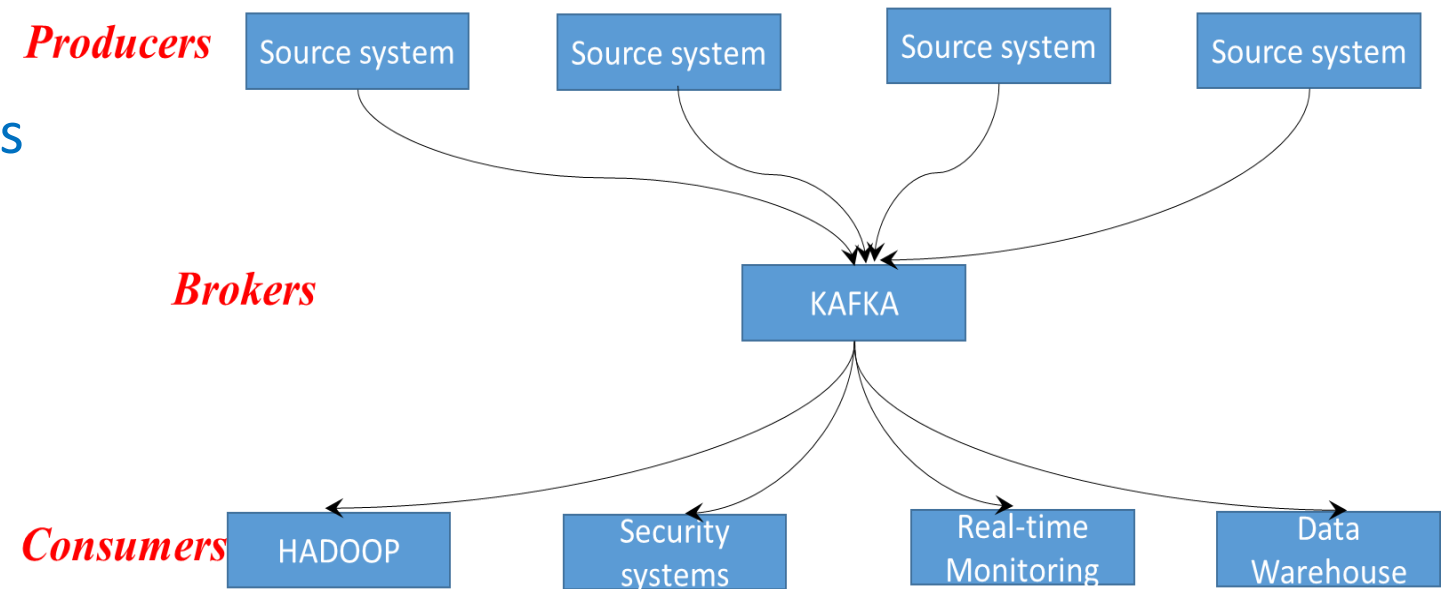
BIG DATA

Role of a producer

- Defines what data it wants to send
- Packages that go into a message and..
- It publishes messages on the communication infrastructure
- Simplest part → called the *publisher*



- Tells the communication infrastructure
 - What type of messages it wants to receive
- Does not tell “who” to receive message from.
- Messages are delivered to the consumer by the communication infrastructure.
- Called the subscriber



Kafka Decouples Data Pipelines

Kafka Communication and Routing



Which Messages get Delivered to Each Subscriber?

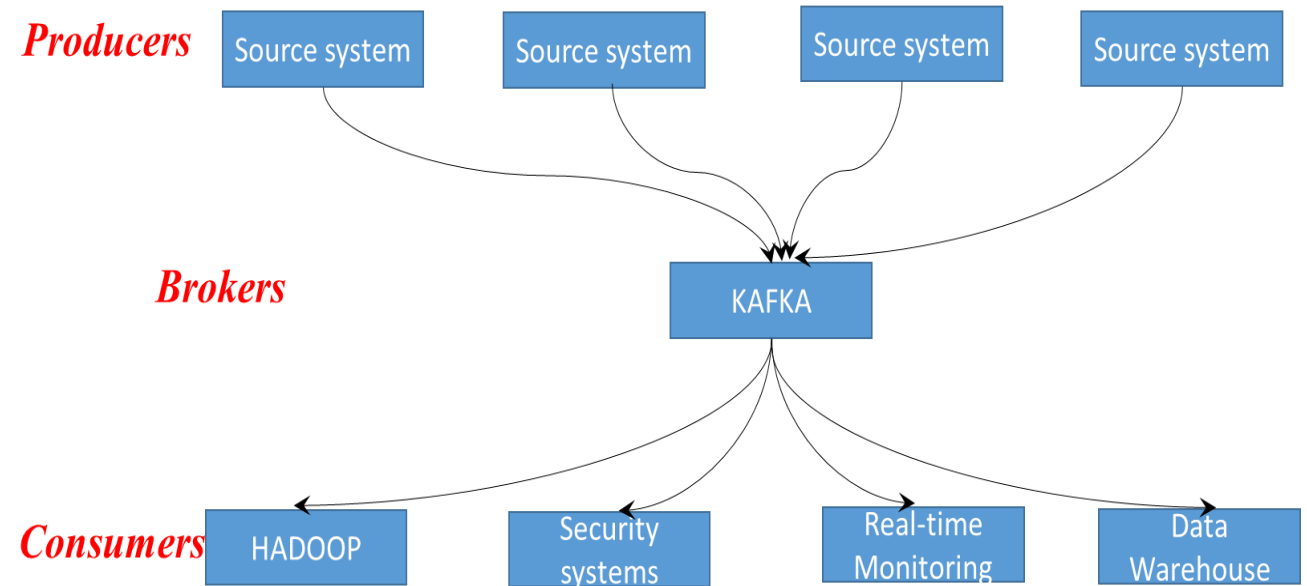
- Topic based system
- Content-based system



BIG DATA

The Communication infrastructure

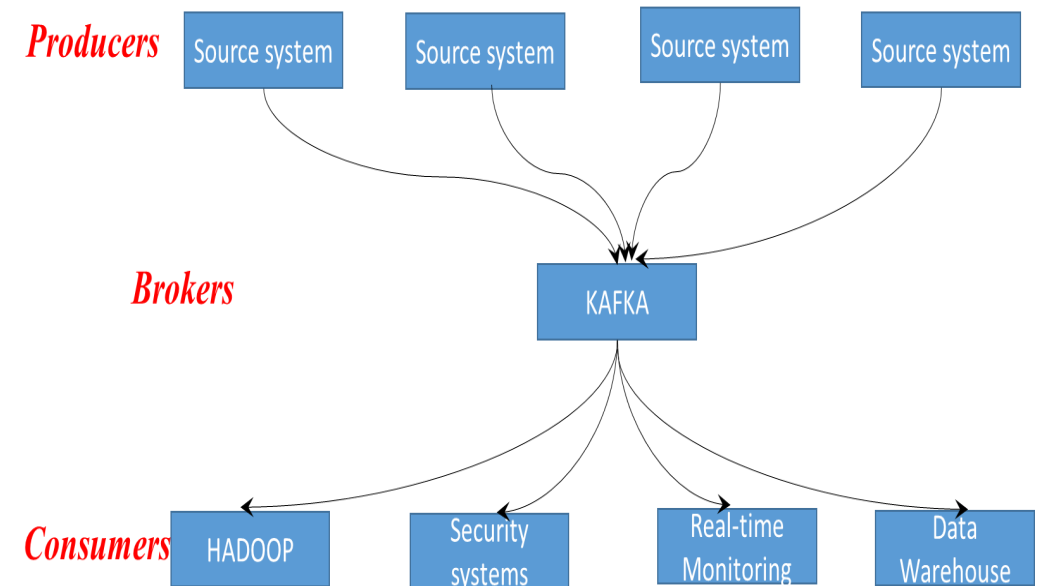
- The most complex part
- Which messages are delivered to each subscriber
- Two models of Routing
 - Topic based systems
 - Content based systems



Kafka Decouples Data Pipelines

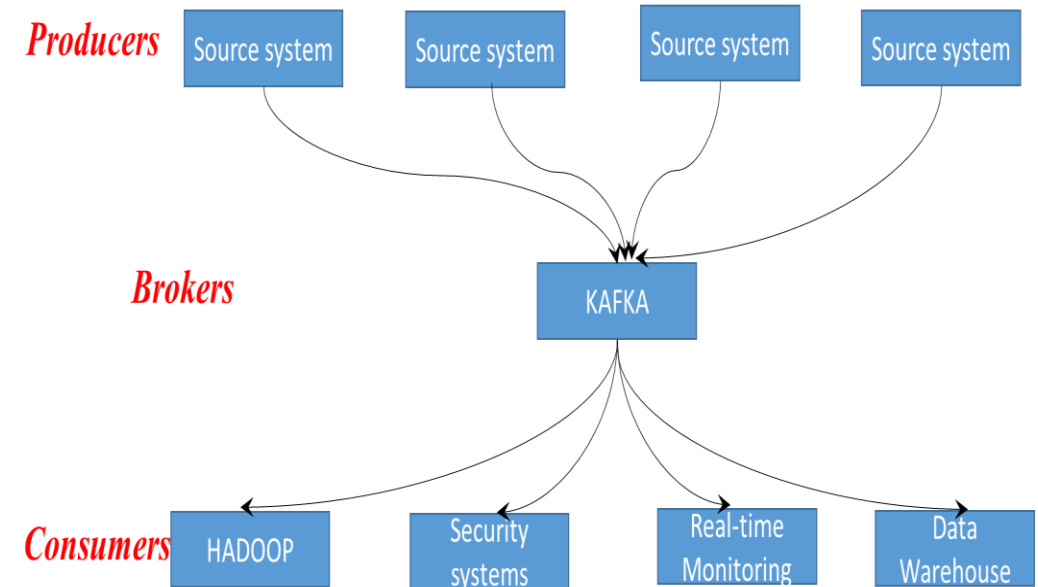


- Publishers send messages with topic labels
- Subscribers subscribe to topics
 - And will receive all messages on that topic
- Example
 - Subscribe to all fire sensors in block



Kafka Decouples Data Pipelines

- Subscribers define matching criteria
- And will receive all messages that match the criteria
- Example
 - Subscribe to advertisements that feature ViratKohli
 - This is not supported by Kafka
- Pattern based
 - Supports a simpler version called *pattern* based where we can give a wildcard expression for a topic
 - Get all topics that have ipl*



Kafka Decouples Data Pipelines

BIG DATA

Communication advantages/disadvantages

Pros

- No hard-wired connections between publishers and subscribers
- Flexible: Easy to add and remove publishers or subscribers


Cons

- Design and maintenance of topics
- Performance overhead due to communication infrastructure

BIG DATA

Exercise 2 (5 minutes)


- Consider a bookstore portal with various activities such as
 - Login
 - List books
 - Get book details
 - Buy book
 - Check status of order
 - Return book
 - Logout
- Assume we have 3 backend modules
 - Security
 - Order processing
 - Book information
- **Would you use a topic-based or content-based system? What would be the topics / content..?**



BIG DATA

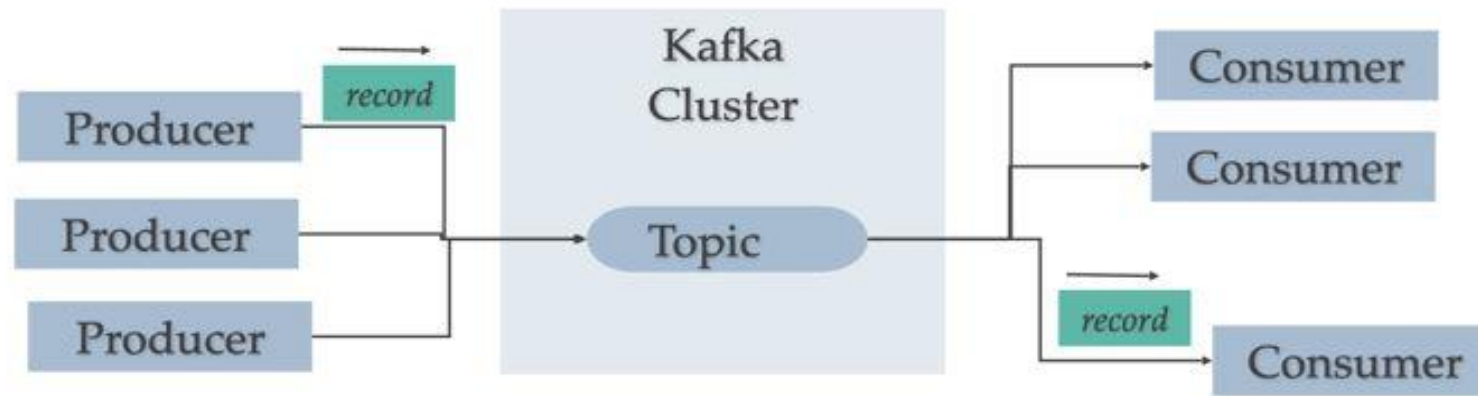
Exercise 2 (Solution)

- Would you use a topic-based or content-based system? What would be the topics / content..?
 - **Probably topic-based, since each message type can be a topic**



Kafka Messages

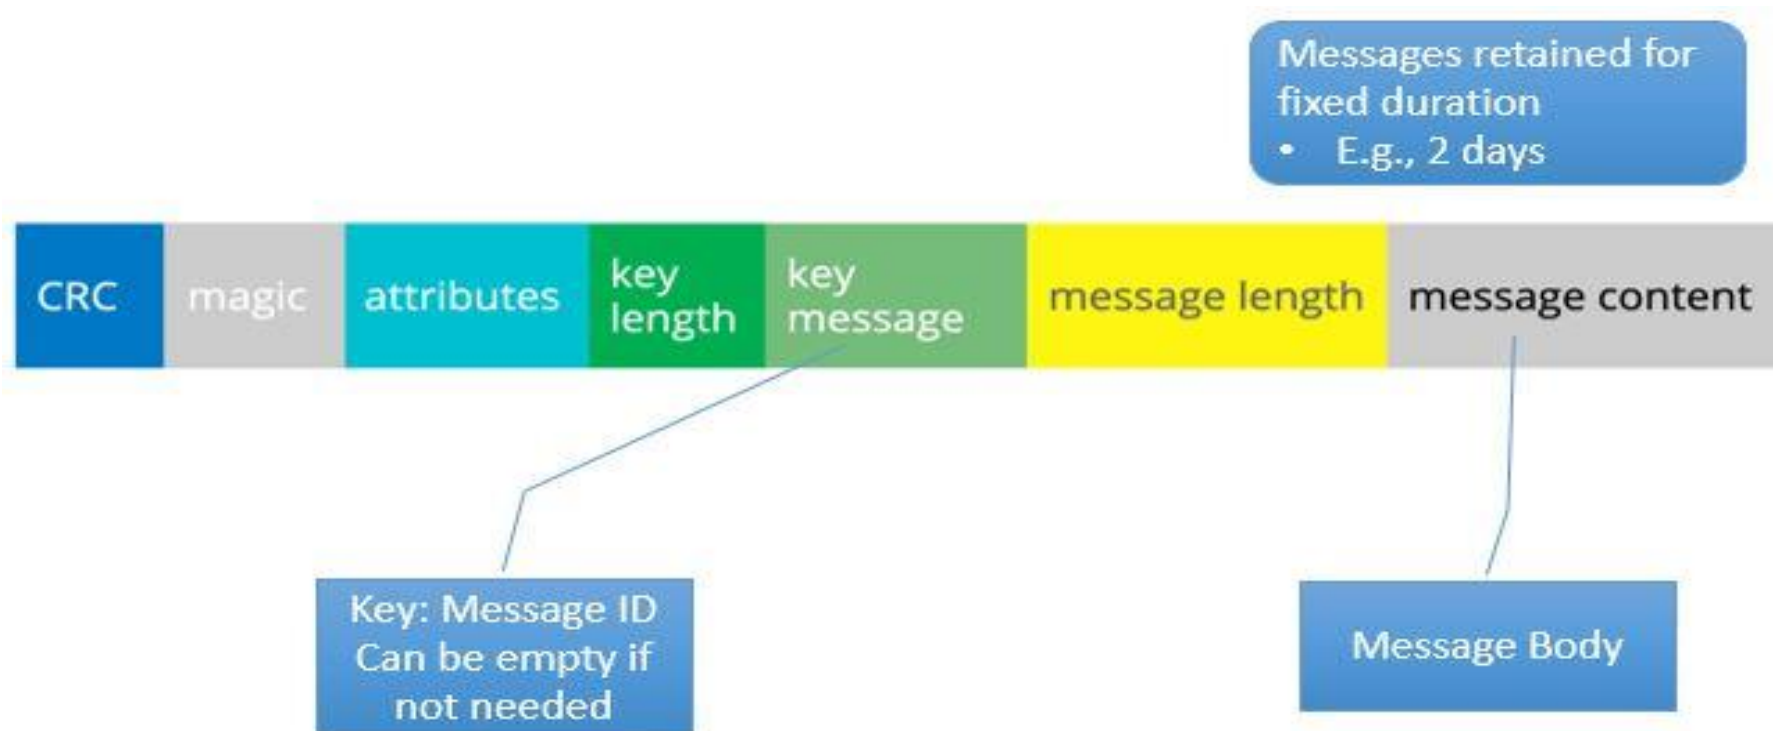
Topics, Producers and Consumers.



Topic name: “/shopping-cart-done“, “/user-signups”

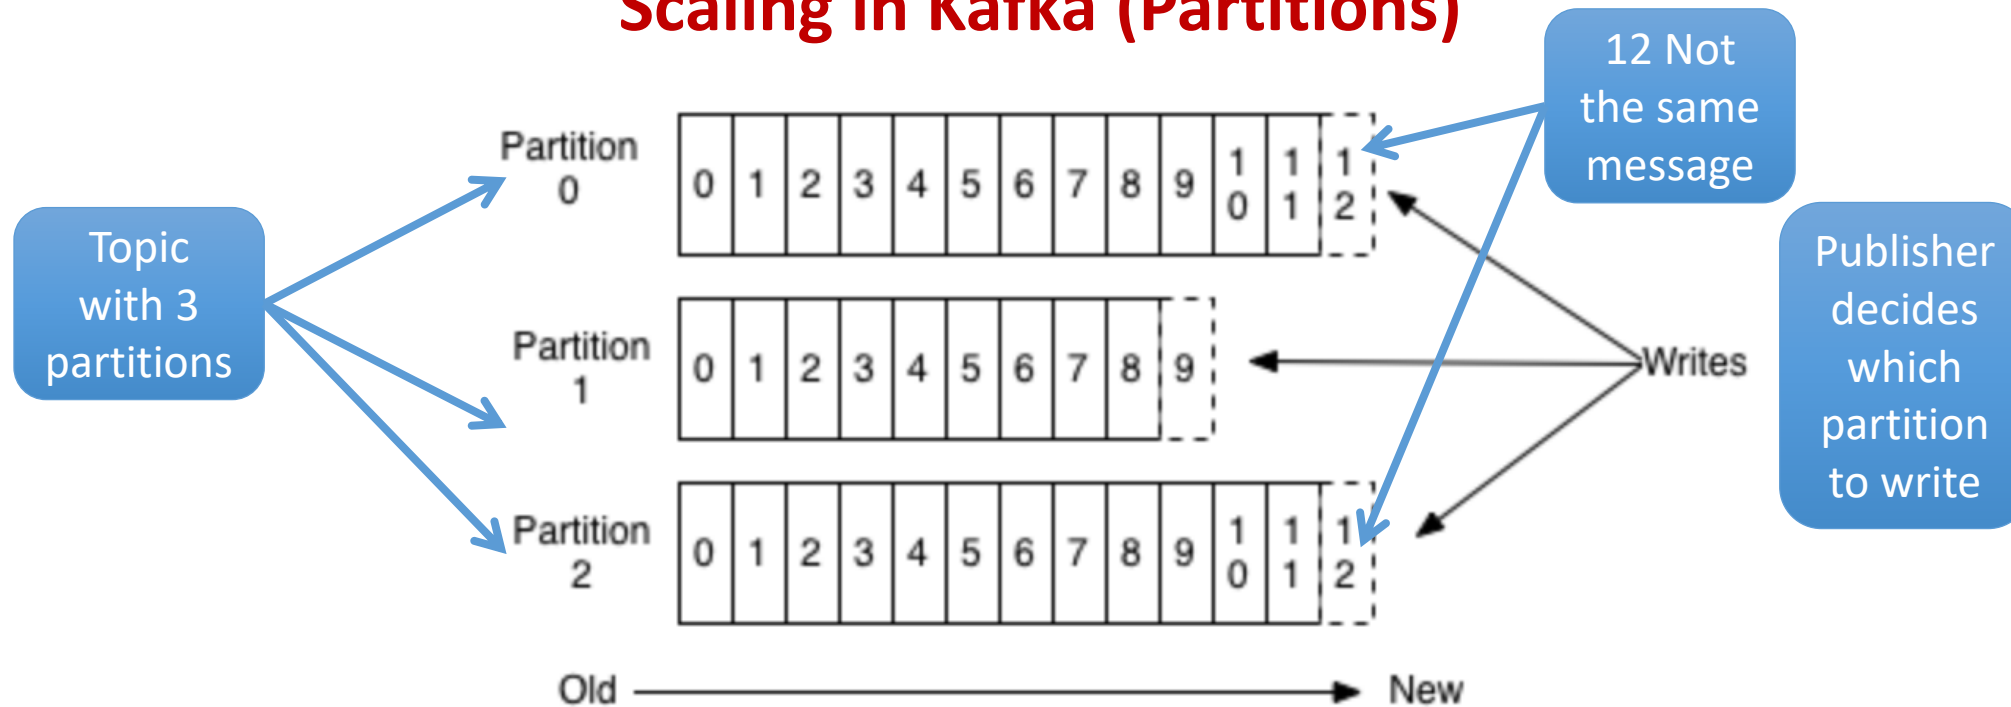
Topic: log data structure on disk

A Kafka Message




Kafka Scaling

Scaling in Kafka (Partitions)



- Partitions allow
 - Log greater than disk size
 - Throughput > single server
- Distributed over servers
- Publisher can load balance
 - Round-robin
 - Based on key
- Messages have an offset

- How can reliability be guaranteed in Kafka?
 - Hint: How does HDFS guarantee reliability?



Fault Tolerance in Kafka

BIG DATA

Fault-tolerance in Kafka (Replication)

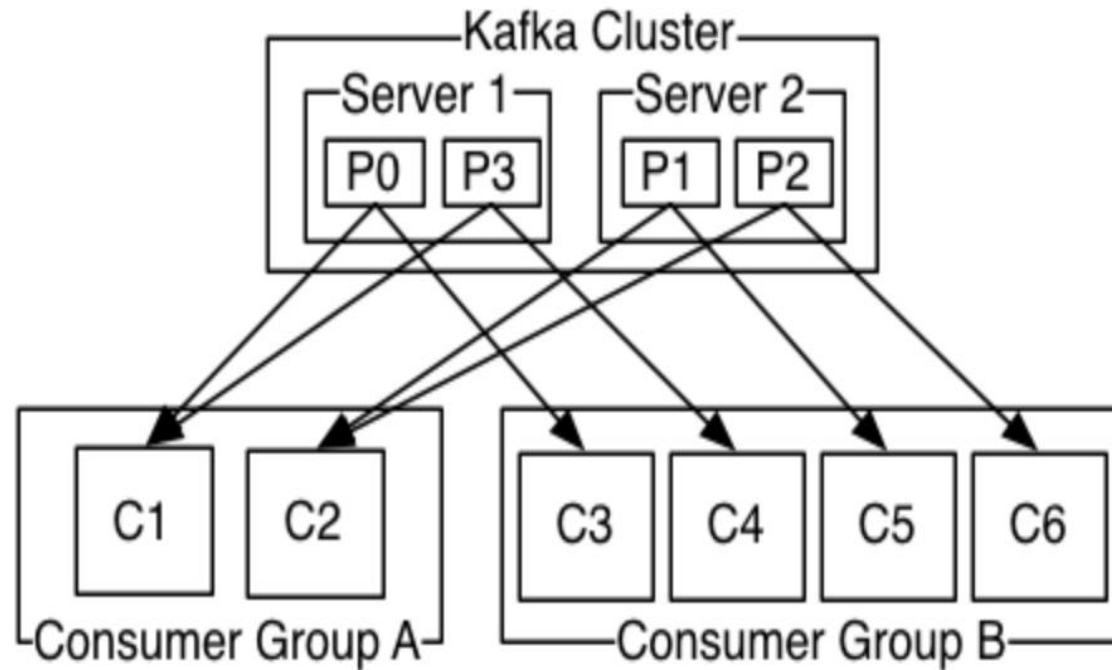
- Partitions can be replicated
 - Leader: all reads, writes
 - Follower replicates
 - New leader if leader fails
- Durability levels
 - Sync: after quorum writes
 - Async
 - 0 = leader only
 - -1 = no write
- Recommended
 - Replicas = 3
 - Quorum = 2



BIG DATA

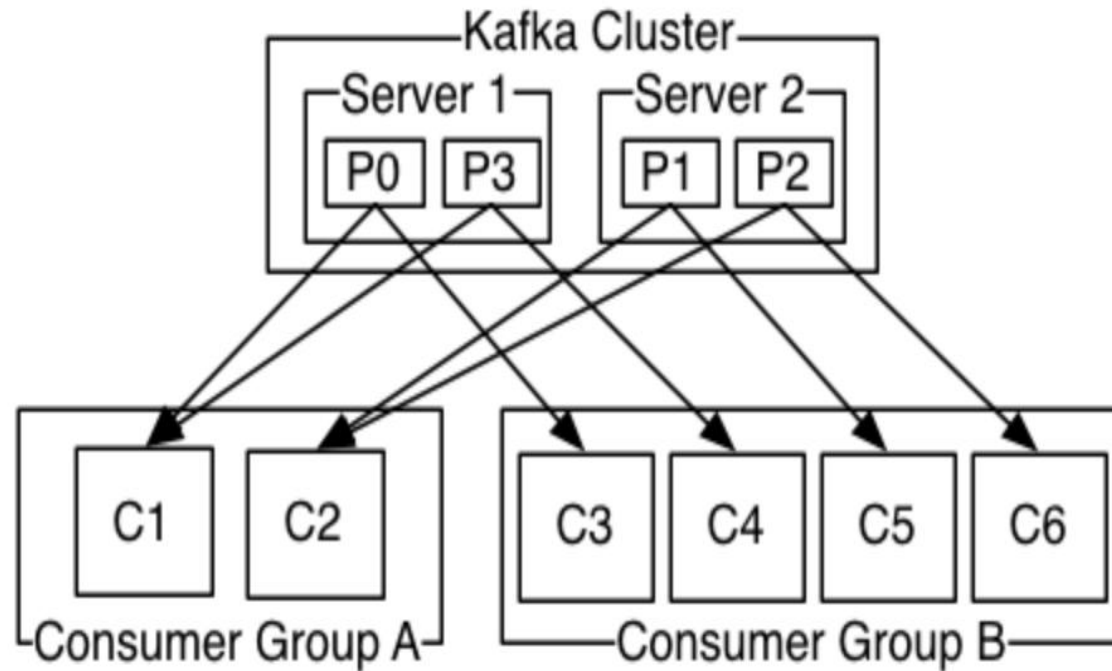
Message delivery to consumers

Topic
with 4
partitions



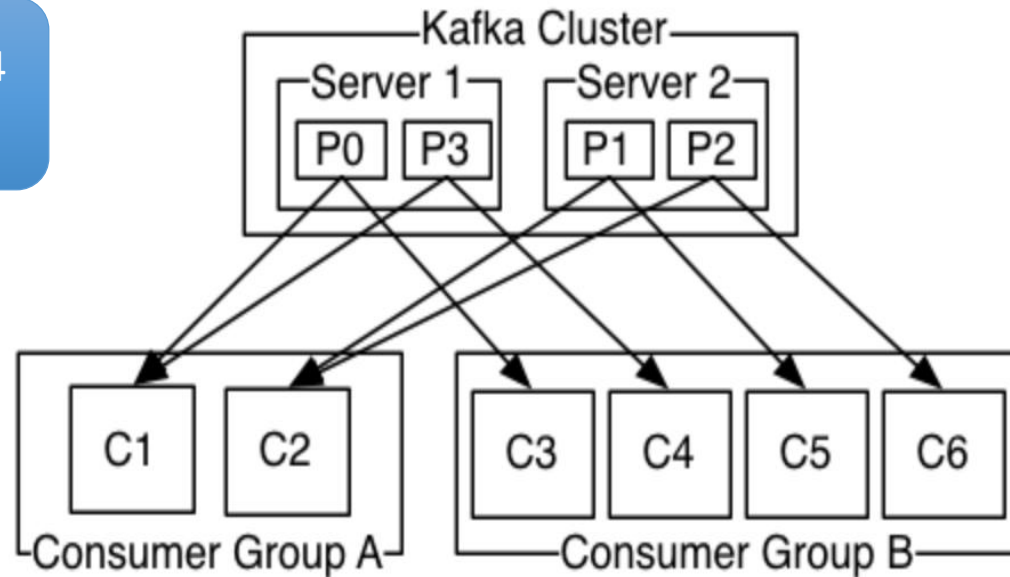
- Consumer group
 - Typically multiple instances of an application
- Partition delivers message to ONE of the group members
 - Load balancing

Topic
with 4
partitions



- In the above configuration, how is the load balanced over all the instances?

Topic with 4 partitions



- Group A
 - C1 is assigned P0, P3
 - C2 is assigned P1, P2

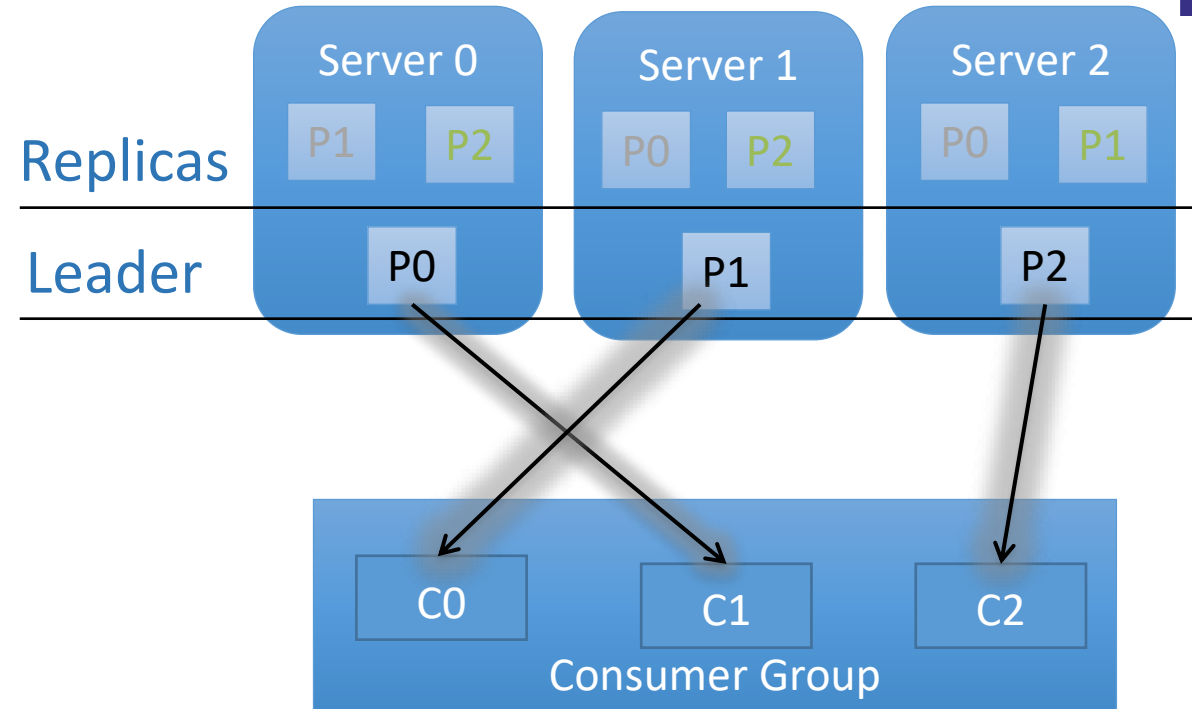
- Group B
 - C3 is assigned P0
 - C4 is assigned P3
 - ...

- Each instance has the same number of partitions


- Suppose we have a Kafka system
 - 1 topic
 - 3 servers
 - 3 partitions
 - 3 replicas per partition
 - Consumer group with 3 instances
- Draw a diagram showing
 - Servers
 - Partitions
 - Consumer instances
 - Partition assignments



- Suppose we have a Kafka system
 - 1 topic
 - 3 servers
 - 3 partitions
 - 3 replicas per partition
 - Consumer group with 3 instances
- Draw a diagram showing
 - Servers
 - Partitions
 - Consumer instances
 - Partition assignments



Kafka Performance



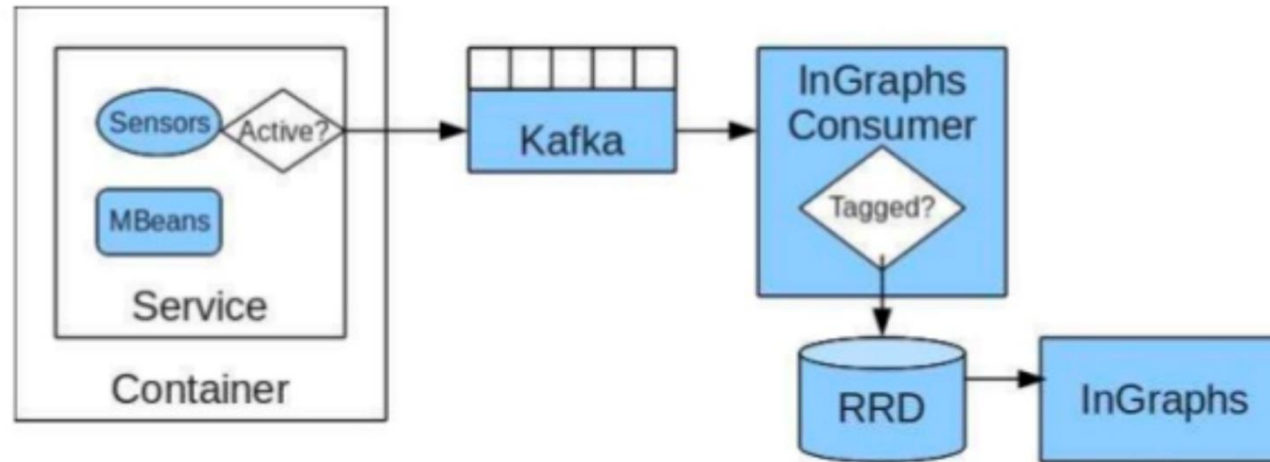
- I/O
 - Sequential reads by consumers
 - Sequential writes by producers
- Zero Copy I/O
 - DMA
 - No copy from kernel to user
- Write batching

Who uses Kafka ?

- ***LinkedIn***: Activity data and Operational metrics.
- ***Twitter***: Uses it as part of Storm – stream Processing infrastructure.
- ***Square***: Kafka as bus to move all system events to various Square data centers(logs, custom events, metrics, and so on). Outputs to Splunk, Gtaphite,Esper-like alerting systems.
- ***Spotify, Uber, Tumbler, Goldman Sachs, PayPal, Box, Cisco, Cloud Fatr, DataDog, LucidWorks, MailChimp, Netflix, etc.***

BIG DATA

Kafka at LinkedIn



- Multiple data centers
- System monitoring
 - 320,000,000 metrics/minute
 - 530 TB of disk space
 - > 210,000 metrics / service
- Analysis – Hadoop, ...

BIG DATA

Additional Reading

- Introduction to Kafka
 - <https://kafka.apache.org/intro>
 - Watch the video on the introduction.
- Kafka in a nutshell
 - <https://sookocheff.com/post/kafka/kafka-in-a-nutshell>





THANK YOU

K V Subramaniam, Usha Devi

Dept. of Computer Science and Engineering

subramaniamkv@pes.edu

ushadevibg@pes.edu