



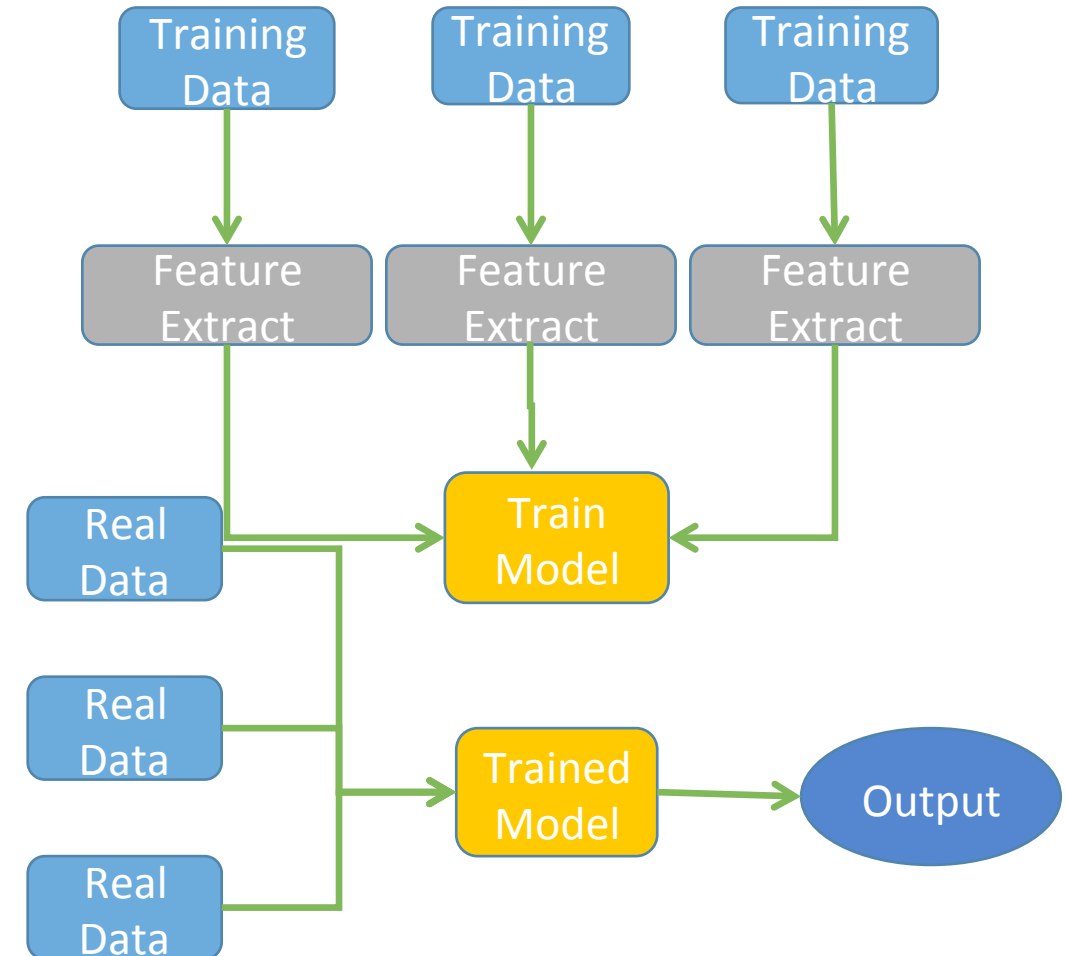
BIG DATA

Machine Learning Algorithms At Scale - Clustering

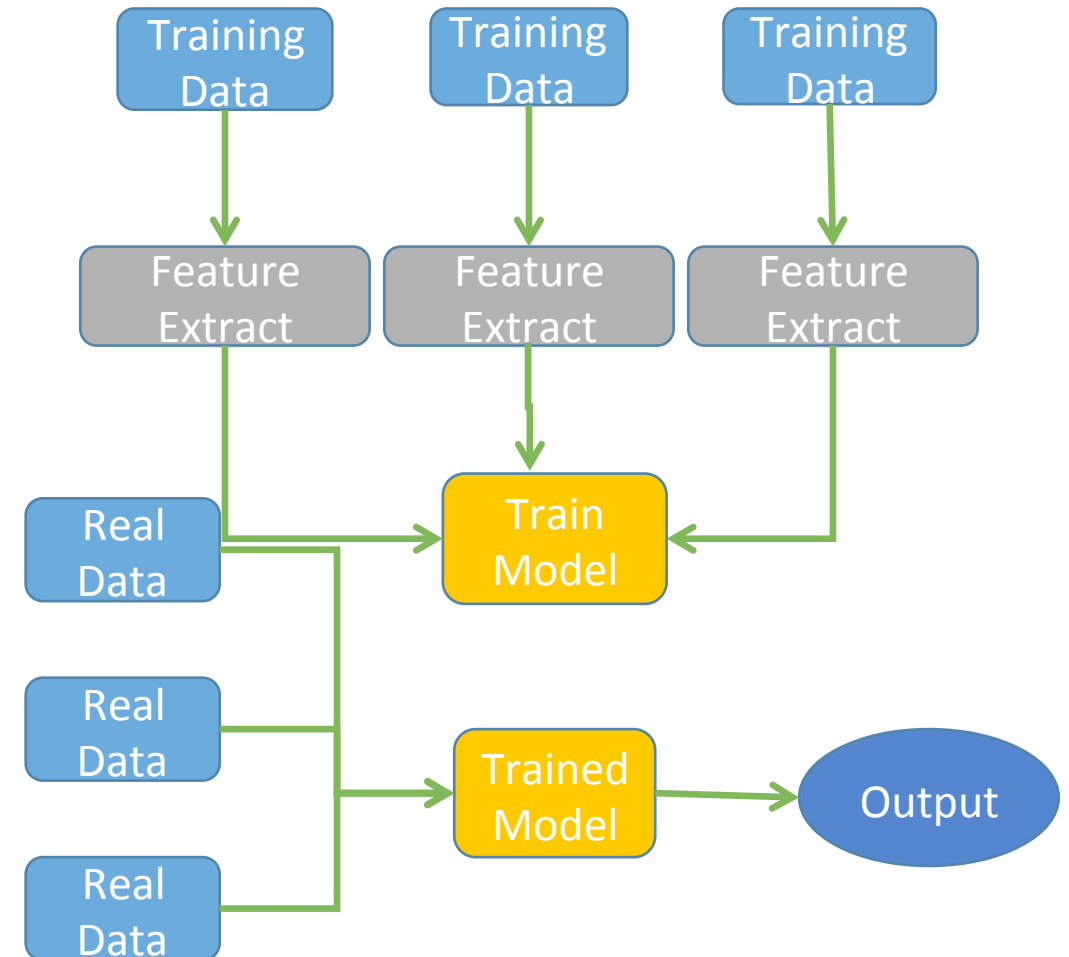
K V Subramaniam

Computer Science and Engineering


- Sometimes, problems are complex
 - We don't want to write explicit programs
 - E.g., recognizing syllables in speech recognition
 - Theoretically, each syllable is a mixture of frequencies
 - Simpler to give the computer examples of syllables and ask the computer to “learn”
- Machine learning is an area of artificial intelligence



- The examples we give are called *training data*
- The program may process the training data to calculate certain quantities called *features*
 - E.g., in speech recognition, the frequencies of each syllable or akshara
- It then uses the features to build a *model*
 - E.g., in speech recognition, which frequencies are associated with each syllable
 - Face recognition: for each person, e.g., how big are eyes, nose, mouth
- This process is called *training*



- To see how good the model is, we can input test data to the program
 - E.g., input syllables to the model
- We can calculate the accuracy
 - How many syllables are correct
- If the accuracy is good, we can use the program in a product
- Input real data, get the output



- You want to write a program to separate out the rocks into different categories.
- What will your approach be?



- There are two types of learning
 - *Supervised* learning
 - *Unsupervised* learning

BIG DATA

Supervised Learning

- We define the concepts we want the computer to learn
- Consider the photograph of pebbles on the right
 - We can input examples of each kind of pebble
 - Pebble 1 Large
 - Pebble 2 Medium
 - Pebble 3 Small
- The program will learn to classify pebbles



- The program
 - Computes various features of the pebbles
 - Groups similar pebbles together
 - It's own classification of pebbles
- These may be different from the way a human being would classify them
- The same program can come up with different classifications if we change some parameters
 - E.g., if we ask the program to classify the pebbles into 3 groups or 4 groups
- Finds structure that's already there in the data



Different Example required for this slide

- Why is unsupervised learning useful?
 - Recall the IPL class project
 - We asked you to group the batsmen into groups
 - We can manually define the groups; i.e., groups like “opener” “attacking” and so on
 - Supervised learning
 - Simpler to feed data about the batsmen and let the program group similar batsmen
 - Unsupervised learning



Supervised

- Input data is labeled
- Input training set consists of a pair
 - Data point or example
 - Classification

Unsupervised

- Input only training data points
- No labels
- Algorithm groups similar data points together

- What is the basic method of machine learning and big data?
- Supervised vs Unsupervised learning
- What is the basic method of ML and big data?
 - Feature extraction
 - Model, train
 - Predict
- Supervised vs unsupervised learning
 - Predefined vs no predefined concepts

- Consider the list of machine learning applications on the right. Which use *supervised learning* and which use *unsupervised learning*?
- In Google News, grouping together similar articles.
- Determining if a particular credit card transaction is fraudulent
- Analyzing an image to determine if a lump is cancerous
- Recommending a product based on what the user buys
- Market segmentation: dividing customers into various groups

- Consider the list of machine learning applications on the right. Which use *supervised learning* and which use *unsupervised learning*?

- In Google News, grouping together similar articles. unsupervised
- Determining if a particular credit card transaction is fraudulent. supervised
- Analyzing an image to determine if a lump is cancerous. supervised
-
- Recommending a product based on what the user buys. unsupervised
- Market segmentation: dividing customers into various groups. either depending on whether we already have pre-defined groups or not

Supervised

- Logistic regression
- Support Vector Machines
- Decision trees
- K-nearest neighbors

Unsupervised

- Principal Component Analysis
- Mixture models
- Hidden Markov models
- K-means




- This class, focus on scalable or large-scale machine learning
 - Google search index (finding page rank over millions of pages)
 - Amazon recommendation (recommendations for millions of users over thousands of products)
- Challenges (as usual)
 - Data size is huge
 - Huge amount of computation
 - Failure is likely (huge amount of hardware)
- Solution
 - Use the right infrastructure (Hadoop, Spark,...)
 - Scalable algorithms
- In this class, we talk about *K-means* and *Alternating Least Squares* algorithm on MapReduce

Scalable machine learning algorithms


- K-means introduction

- Clustering
 - Partition a number of data points into related groups called clusters.
 - K-means clustering partitions a dataset into a specified number k of clusters
 - The points in a cluster should be similar to each other
- E.g., the IPL modeling, batsmen can be characterized by many parameters
 - E.g., strike rate, highest score, position (opener, ...)
- To divide batsmen into groups, we need to be able to measure how similar batsmen are to each other

- We can consider each batsman to be a point in an n -dimensional space
 - n is the number of parameters we are measuring
- A *distance metric* measures the similarity (distance) between the two points
- For simplicity, consider a 2D space
 - Euclidean: Geometric distance
 - Manhattan: city blocks, used in traffic
 - Cosine: measures angle between points – used if points can have very different distances from origin




Manhattan



Euclidean


BIG DATA

Example for clustering:




Note the outliers in the clusters.


Demonstration of the standard algorithm




1) k initial "means" (in this case $k=3$) are randomly selected from the data set (shown in color).



2) k clusters are created by associating every observation with the nearest mean. The partitions here represent the **Voronoi diagram** generated by the means.



3) The **centroid** of each of the k clusters becomes the new means.



4) Steps 2 and 3 are repeated until convergence has been reached.

Iterative algorithm until convergence

- **Initialize:** Select K points at random (Centers)
- **Step 1:** For each data point, assign it to the closest center
 - Now we formed K clusters
- **Step 2:** For each cluster, re-compute the centers
 - E.g., in the case of 2D points →
 - X: average over all x-axis points in the cluster
 - Y: average over all y-axis points in the cluster
- **Loop check:** If the new centers are different from the old centers (previous iteration) → Go to Step 1

- How can the k -means algorithm be modified to run with MapReduce?
- What is the output of Map and Reduce stages?

Hints:

- Iterative algorithm like page rank
- Which steps can be done in Map and which in Reduce?

- **Initialize:** Select K points at random (Centers)
- **Step 1:** For each data point, assign it to the closest center
 - Now we formed K clusters
- **Step 2:** For each cluster, re-compute the centers
 - E.g., in the case of 2D points →
 - X: average over all x-axis points in the cluster
 - Y: average over all y-axis points in the cluster
- **Loop check:** If the new centers are different from the old centers (previous iteration) → Go to Step 1

Scalable machine learning algorithms

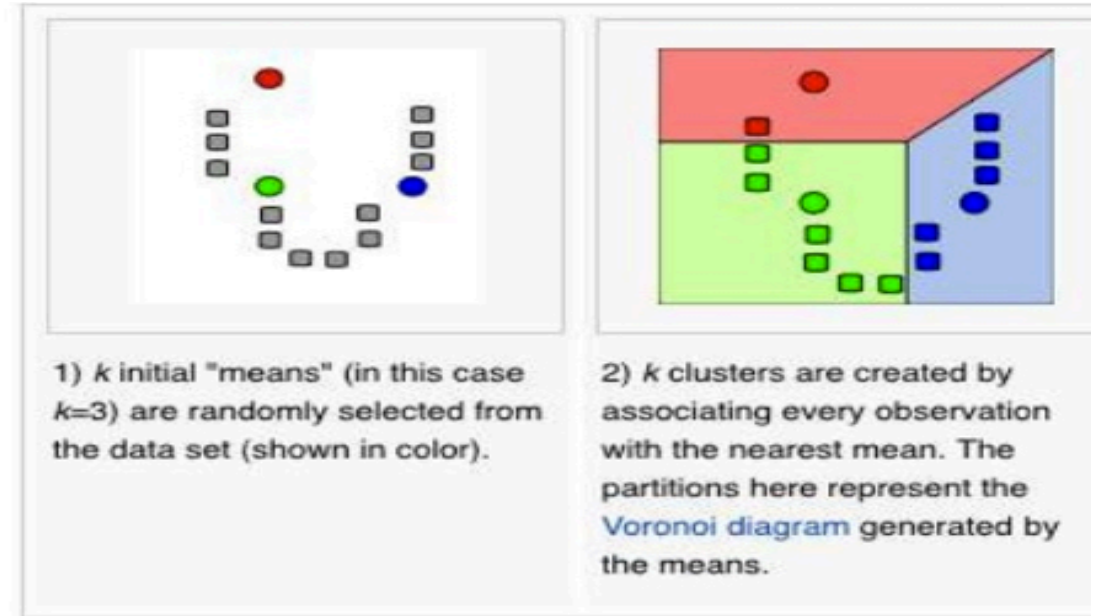
- K-means with Map-Reduce

- **Input**


- Dataset (set of points in 2D) --Large
- Initial centroids (K points) --Small

- **Map (reads 2 files as input)**

- Each map reads the K-centroids + one block from dataset
- Assign each point to the closest centroid
- Output <centroid, point> - centroid is the key



- **Reduce**
 - Gets all points for a given centroid
 - Re-compute a new centroid for this cluster
 - Output: <new centroid>
- **Loop check**
 - Compare the old and new set of K-centroids
 - If similar → Stop
 - Else
 - If max iterations has reached → Stop
 - Else → Start another Map-Reduce Iteration



- Given the following points
 - 20, 30, 99, 102,
 - 53, 9, 11, 54
- Partition them into two clusters using k-means assuming initial centroids are 20, 30.
- Assume that each row of numbers is on a different machine
- Show what the keys and values are for one iteration of k-means

- Mapper1 output
 - 20, 20,
 - 30, 30,
 - 30, 99,
 - 30, 102
- Mapper 2 output
 - 30, 53,
 - 20, 9,
 - 20, 11,
 - 30, 54
- Reducer input
 - 20 , <20, 9, 11>
 - 30, <30, 99, 102, 53, 54>
- Reducer output
 - 13.33 and 67.6

Scalable machine learning algorithms

- K-means optimizations










- **Use of Combiners**
 - Similar to the reducer
 - Computes for each centroid the local sums (and counts) of the assigned points
 - Sends to the reducer <centroid, <partial sums>>
- **Use of Single Reducer**
 - Amount of data to reducers is very small
 - Single reducer can tell whether any of the centers has changed or not
 - Creates a single output file

Scalable machine learning algorithms - Alternating least squares

4 star rating

Unknown rating

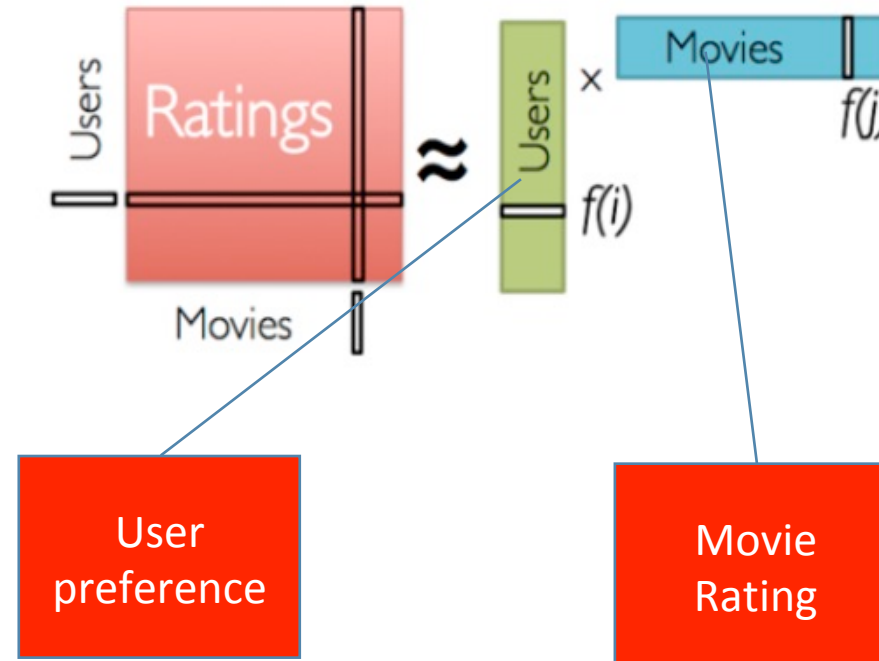
User-Item Rating matrix is generally very sparse i.e., most entries are unknown

			
	★	★★★★	?
	★	★★★	★★
	★★★★	?	★
	★	?	★★
	?	★★★	★★
	★★★★	★★	?

- Recover a rating matrix from a subset of its entries.

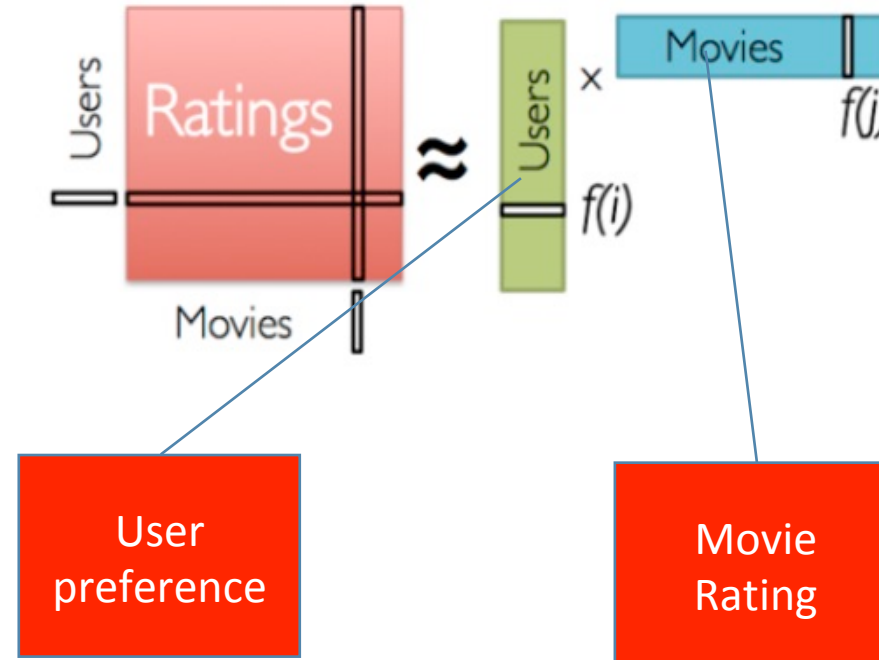


- Express User-Item Rating matrix (R) as a product of
 - User vector (A) dimension n =no of users
 - Item vector (B) dimension m =no of movies
 - Calculate A, B such that $R \approx AB$
 - A is a $n \times 1$ vector, B is a $1 \times m$ vector
 - R will be an $n \times m$ vector
- Suppose we need to find r_{ij} which is unknown
 - This is the rating of user i for item j
 - Calculate $R' = AB$
 - Use the ij^{th} element of R'



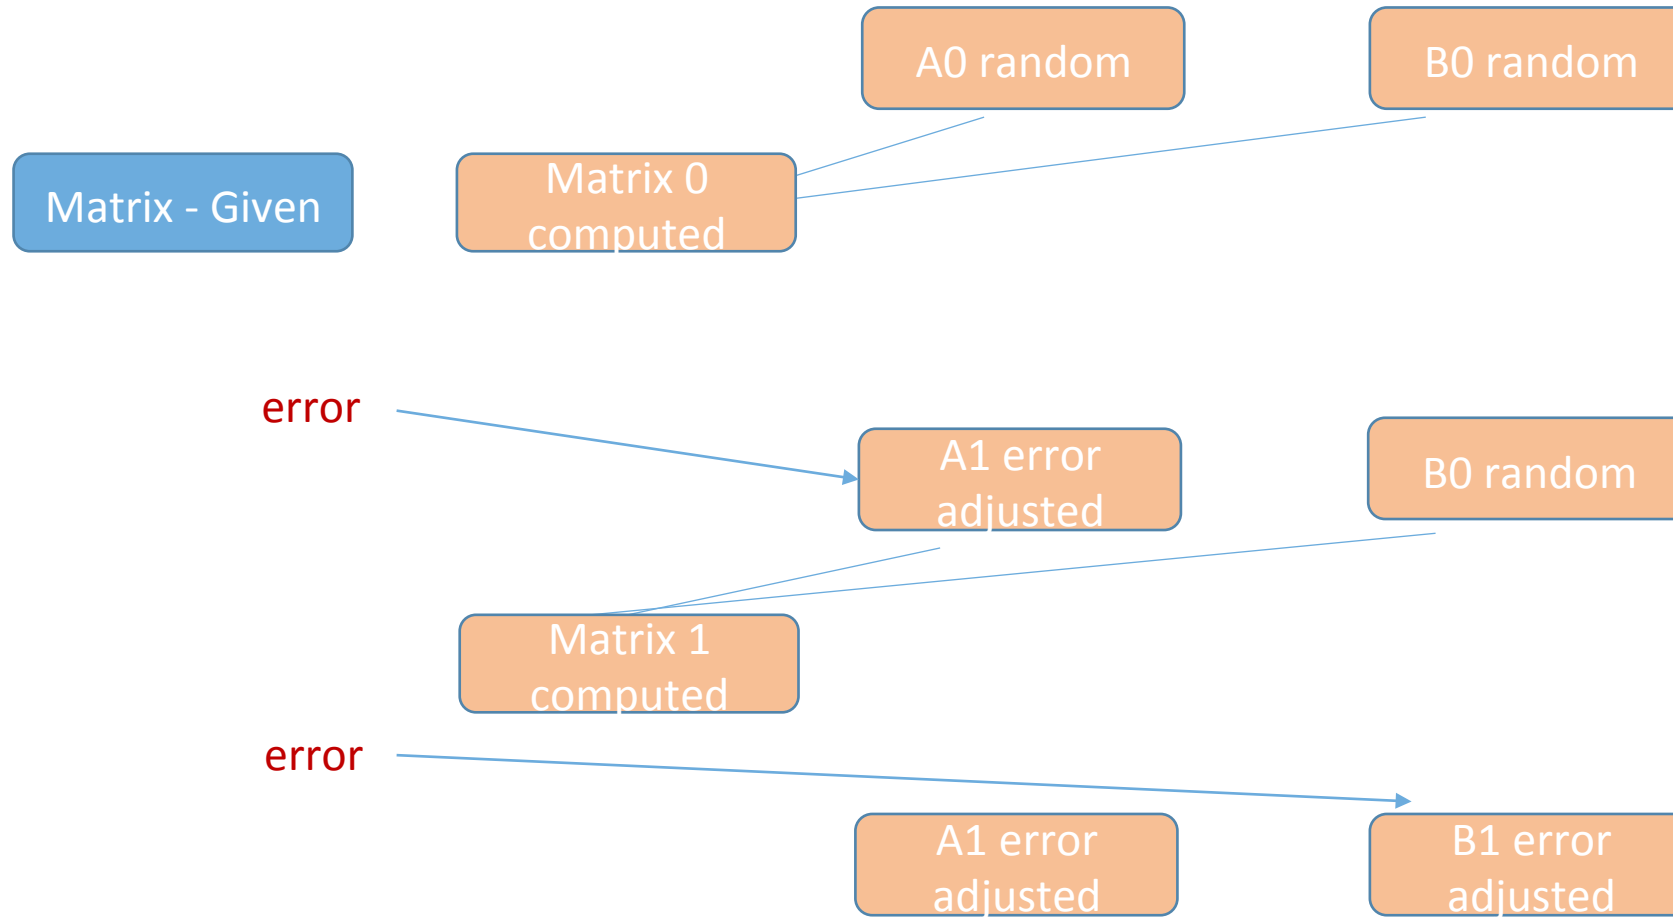
Xiangrui Meng, *MLlib: scalable machine learning on Spark*, Spark Workshop April 2014, <http://stanford.edu/~rezab/sparkworkshop/>

- Both A and B are not known
- This is like an optimization problem and can use Gradient Descent
 - But GD is too slow.
- Alternative – Factorize the matrix R into A and B
- We have to factorize R to get
 - A and B



- Start with Random A and B
- The algorithm will loop until the correct value is calculated
 - Each iteration, the program will calculate new values for A and B .
 - Let A_i and B_i be the values of A and B on the i^{th} iteration of the loop.
- On the i^{th} iteration of the loop
 - We have calculated A_{i-1} and B_{i-1} on the previous iteration
 - Step 1: assume B_{i-1} is correct. Calculate best value for A_i
 - Step 2: assume A_i is correct. Calculate best value for B_i
 - Loop until converged

- On the i^{th} iteration of the loop
 - We have calculated A_{i-1} and B_{i-1} on the previous iteration
 - Step 1: assume B_{i-1} is correct. Calculate best value for A_i How???
 - Consider $R - A_i B_{i-1}^T$
 - B_{i-1} and R are fixed. For any value of A_i , we can find $R - A_i B_{i-1}^T$
 - For any value of A_i , $R - A_i B_{i-1}^T$ is like an error term
 - The difference between R (the correct rating) and $A_i B_{i-1}^T$
 - The smaller the value of $R - A_i B_{i-1}^T$, the better
 - Since $R - A_i B_{i-1}^T$ can be -ve, we take $||R - A_i B_{i-1}^T||$ (determinant) and find A_i that will minimize
 - It can be shown that the solution is $A_i = (B_{i-1}^T B_{i-1})^{-1} B_{i-1}^T R^T$
 - Similarly for B_i
 - For the mathematics lovers, this is a least squares regression estimate



- How can the ALS algorithm be modified to run with MapReduce?

- Start with Random A and B
- The algorithm will loop until the correct value is calculated
 - Each iteration, the program will calculate new values for A and B .
 - Let A_i and B_i be the values of A and B on the i^{th} iteration of the loop.
- On the i^{th} iteration of the loop
 - We have calculated A_{i-1} and B_{i-1} on the previous iteration
 - Step 1: assume B_{i-1} is correct. Calculate best value for $A_i = (B_{i-1}^T B_{i-1})^{-1} B_{i-1}^T R^T$
 - Step 2: assume A_i is correct. Similarly calculate best value for B_i
 - Loop until converged

Scalable machine learning algorithms - Alternating least squares with MR

How can the ALS algorithm be modified to run with MapReduce?

Solution

- In Step 1, A_i is calculated by doing a number of matrix multiplications and inversions
- We have studied how to do matrix multiplication using MapReduce
- There are similar algorithms for doing matrix inverse using MapReduce

- Start with Random A and B
- The algorithm will loop until the correct value is calculated
 - Each iteration, the program will calculate new values for A and B .
 - Let A_i and B_i be the values of A and B on the i^{th} iteration of the loop.
- On the i^{th} iteration of the loop
 - We have calculated A_{i-1} and B_{i-1} on the previous iteration
 - Step 1: assume B_{i-1} is correct. Calculate best value for $A_i = (B_{i-1}^T B_{i-1})^{-1} B_{i-1}^T R^T$
 - Step 2: assume A_i is correct. Similarly calculate best value for B_i
 - Loop until converged



THANK YOU

K V Subramaniam, Usha Devi

Dept. of Computer Science and Engineering

subramaniamkv@pes.edu

ushadevibg@pes.edu