



**BIG DATA**

**Machine learning Case Study**

**Spark MLlib**

---

**K V Subramaniam**

Computer Science and Engineering

**Goal : Given a text Document , Predict its topic.**

Dataset: "20 Newsgroups"  
From UCI KDD Archive

### Features

Subject: Re: Lexan Polish?  
Suggest McQuires #1 plastic  
polish. It will help somewhat  
but nothing will remove deep  
scratches without making it  
worse than it already is.  
McQuires will do something...

\  
text, image, vector, ...




### Label

1: about science  
0: not about science

\  
CTR, inches of rainfall, ...

Goal : Find best line separating two sets of points.

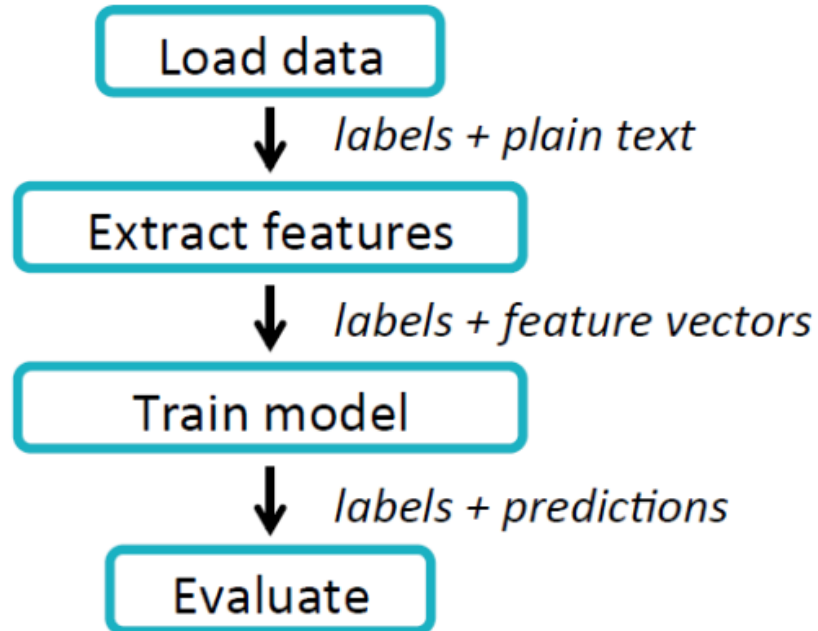


Classify into science and non-science. Each point represents a document.

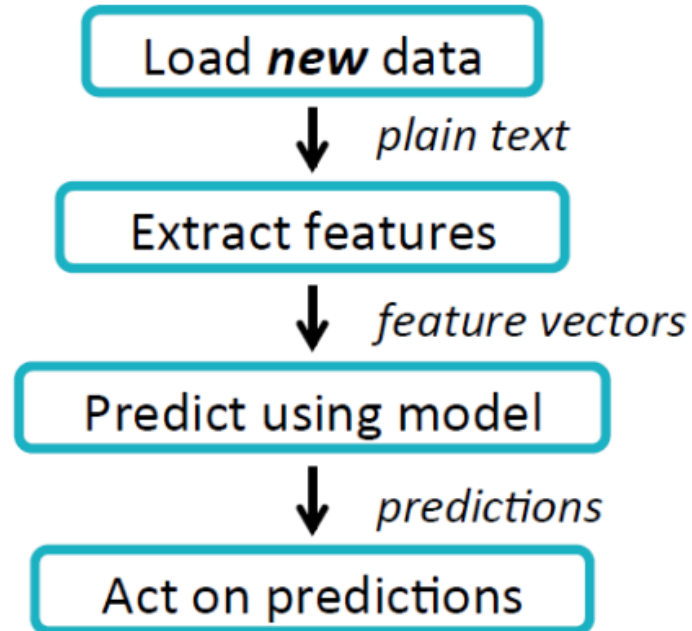
Further details on logistic regression can be found at - [https://en.wikipedia.org/wiki/Logistic\\_regression](https://en.wikipedia.org/wiki/Logistic_regression)

# Machine Learning Workflow and Challenges

### TRAINING



### TESTING/PRODUCTION



*Almost  
identical  
workflow*

## What are the pain points?

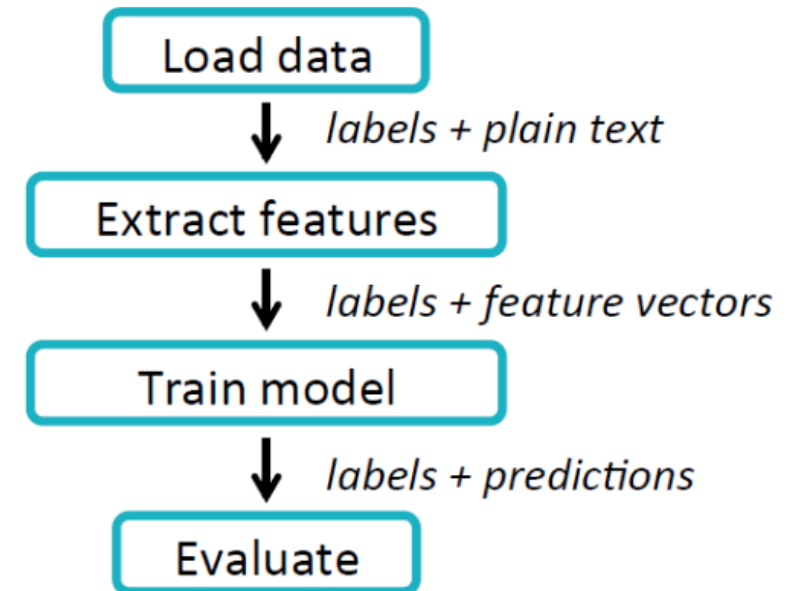
---

- Create and Handle Many RDDs and data types
  - Labels, features, predictions..
- Write as a script
  - Whole pipeline needs to be coded as a script
  - Not modular
- Tune parameters
  - Key part of ML
  - Training many models
    - For different splits of data
    - Different sets of parameters

## Solving the Machine learning challenges

- Make RDDs easier to read
  - Have to explicitly break up the fields in RDD
  - E.g., break line into blank separated tokens
- As developers we would like to just
  - Program to extract features
  - Specify the model to be used
- However, ML needs additional work
  - Write a script to do all the steps
  - Train the model
  - Evaluate the error of the model by testing

### TRAINING





- Reading RDDs: DataFrame

Solves the RDD creation pain-point

- ML Pipeline

- Transformers
- Estimators
- Evaluators

Solves the Scripting..

- Parameters

- API
- Tuning

Solves the parameter tuning pain point

- Recall
- Announced Feb 2015
- Inspired by data frames in R and Pandas in Python
- Works in:



<http://training.databricks.com/intro.pdf>

## What is a Dataframe?

- a distributed collection of data organized into named columns
- Like a table in a relational database

### Features

- Scales from KBs to PBs
- Supports wide array of data formats and storage systems (Hive, existing RDDs, etc)
- State-of-the-art optimization and code generation via Spark SQL Catalyst optimizer
- APIs in Python, Java

## Dataframe : RDD + Schema + DSL

Named columns with types


```
label: Double  
text: String  
words: Seq[String]  
features: Vector  
prediction: Double
```

label	text	words	features
0	This is ...	["This", "is", ...]	[0.5, 1.2, ...]
0	When we ...	["When", ...]	[1.9, -0.8, ...]
1	Knuth was ...	["Knuth", ...]	[0.0, 8.7, ...]
0	Or you ...	["Or", "you", ...]	[0.1, -0.6, ...]

Domain-Specific Language

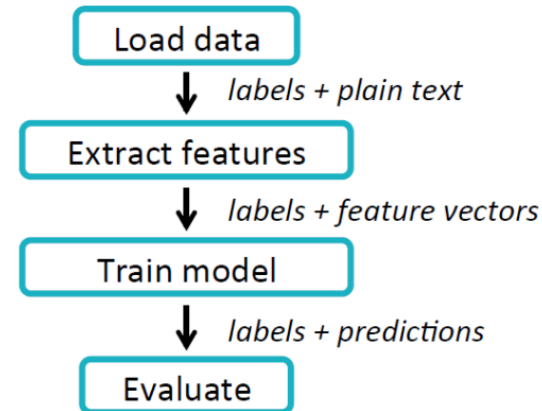
```
# Select science articles  
sciDocs =  
  data.filter("label" == 1)  
  
# Scale labels  
data("label") * 0.5
```

## ML Pipelines

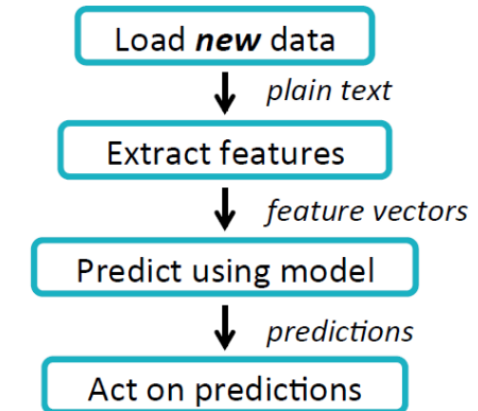


- Introduced in Spark 1.2 and 1.3
- Allows developers to just
  - Program to extract features
  - Specify the model to be used
- Automates the process of
  - Write a script to do all the steps
  - Train the model
  - Evaluate the error of the model by testing
  - Or deploy in production

### TRAINING



### TESTING/PRODUCTION



### Transformers

- Extract features from DataFrame
- Features are stored in a new DataFrame

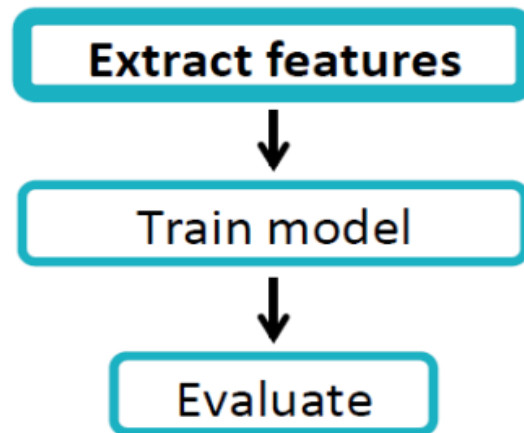
### Estimators

- ML Algorithms
- MLLib has standard defined ML algorithms (e.g., Logistic Regression)
- User can add his own


### Evaluators

- Compute predictions and estimate metrics such as error
- Tune algorithm parameters
- Evaluator depends upon estimator
  - Evaluator that trains Logistic Regression cannot be used for Decision Trees

### TRAINING



```
def transform(DataFrame) : DataFrame
```




Label	Text
0	
1	


Label	Text	Features
0		
1		



## TRAINING



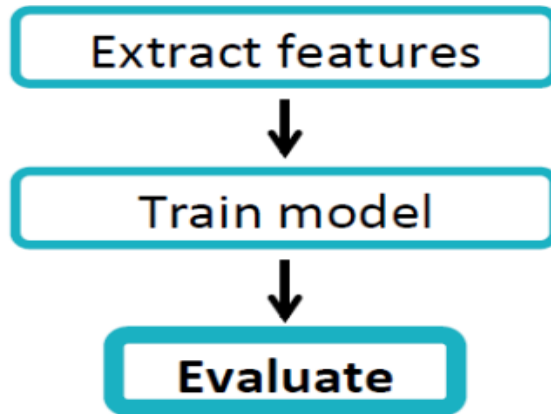
```
def fit(DataFrame): Model
```




LogisticRegression  
Model

Label	Text	Features
0		
1		

### TRAINING



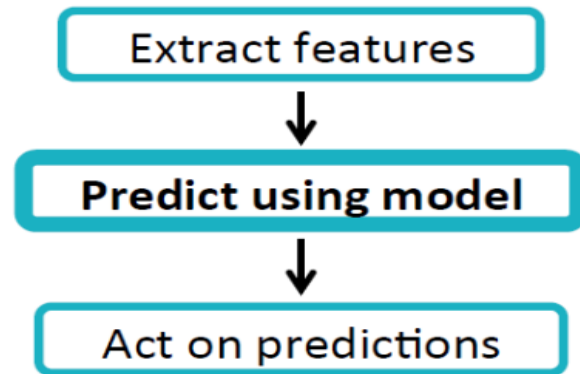
```
def evaluate(DataFrame) : Double
```



Metric:  
accuracy  
AUC  
MSE  
...


Label	Text	Features	Prediction
0			
1			

### TESTING/PRODUCTION

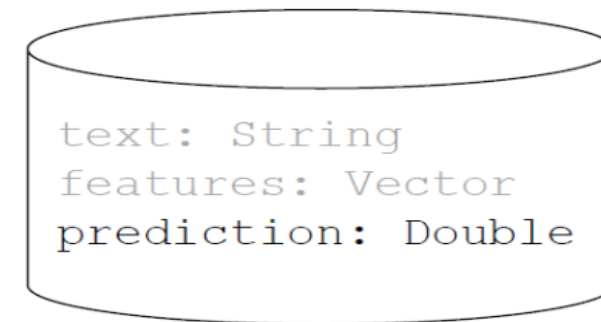


**Model is a type of Transformer**

```
def transform(DataFrame): DataFrame
```




Text	Features




Text	Features	Prediction

### TRAINING



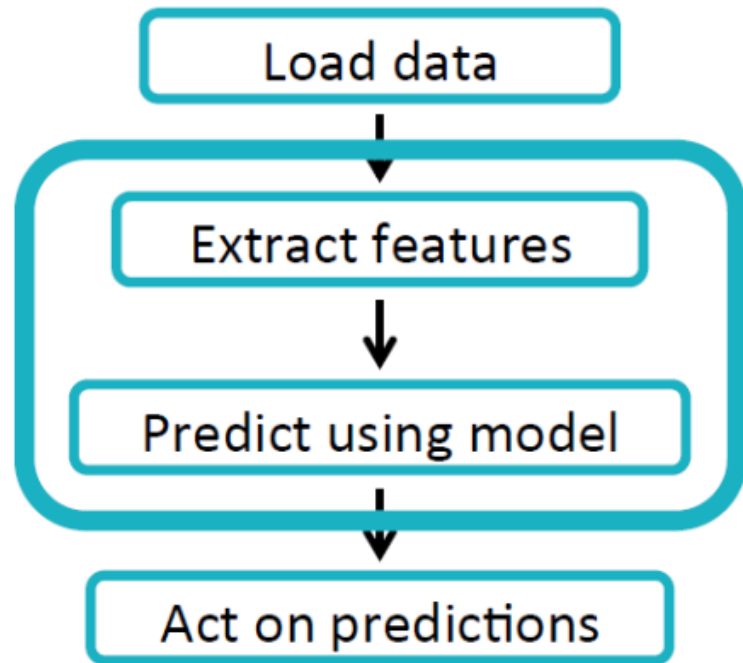
**Pipeline is a type of Estimator**

```
def fit(DataFrame): Model
```




PipelineModel

### TESTING/PRODUCTION

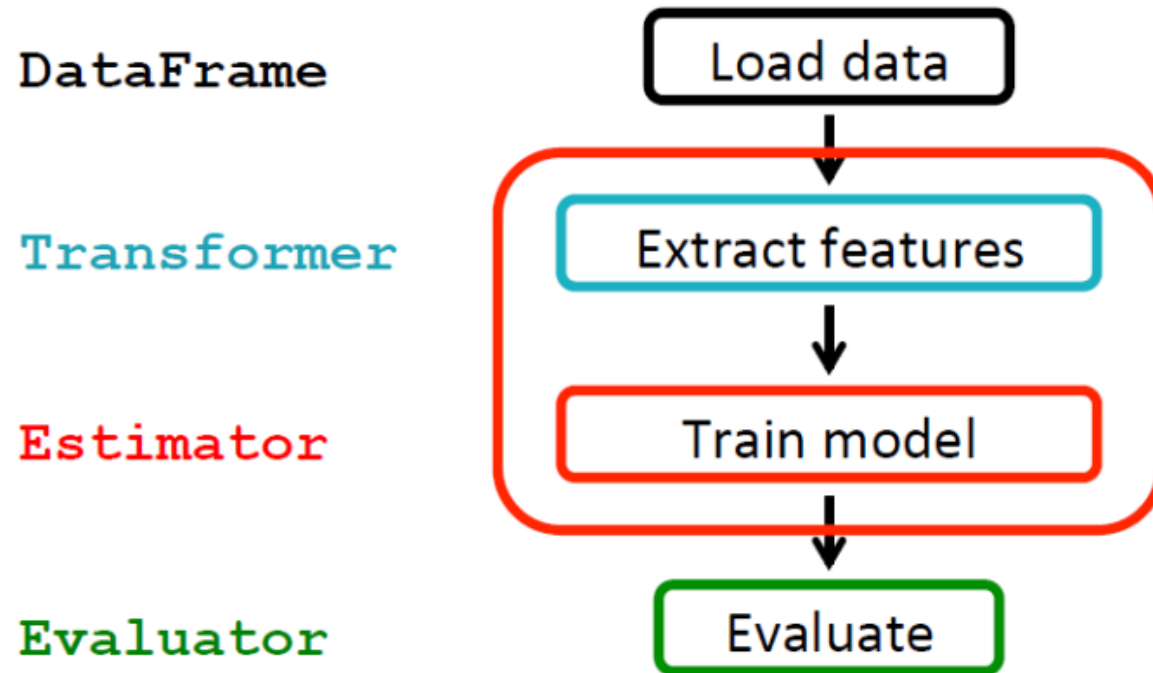


### PipelineModel is a type of Transformer

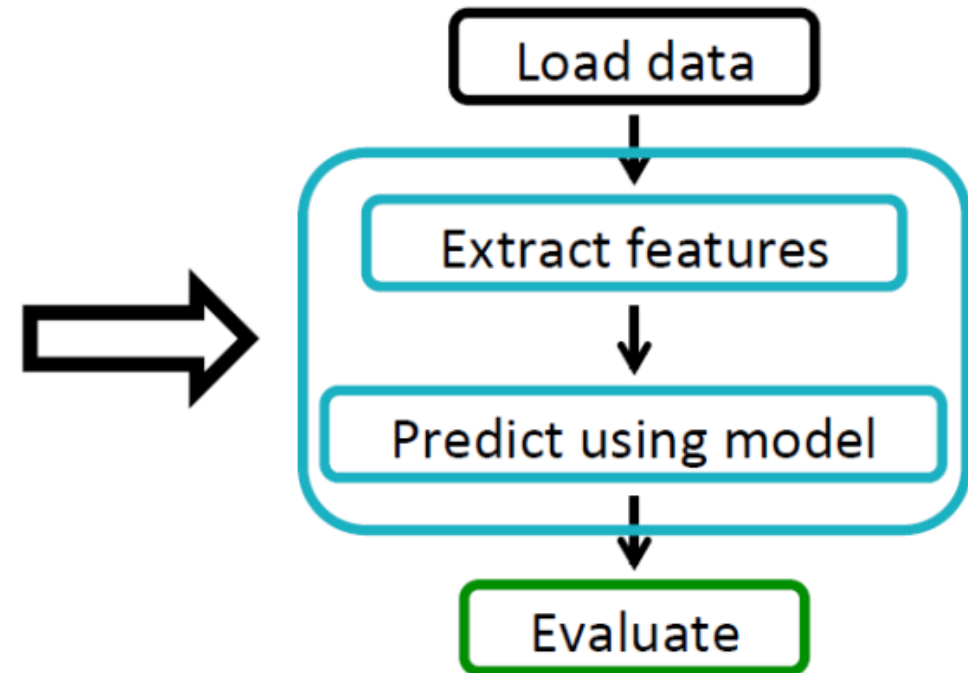
```
def transform(DataFrame): DataFrame
```



### TRAINING



### TESTING

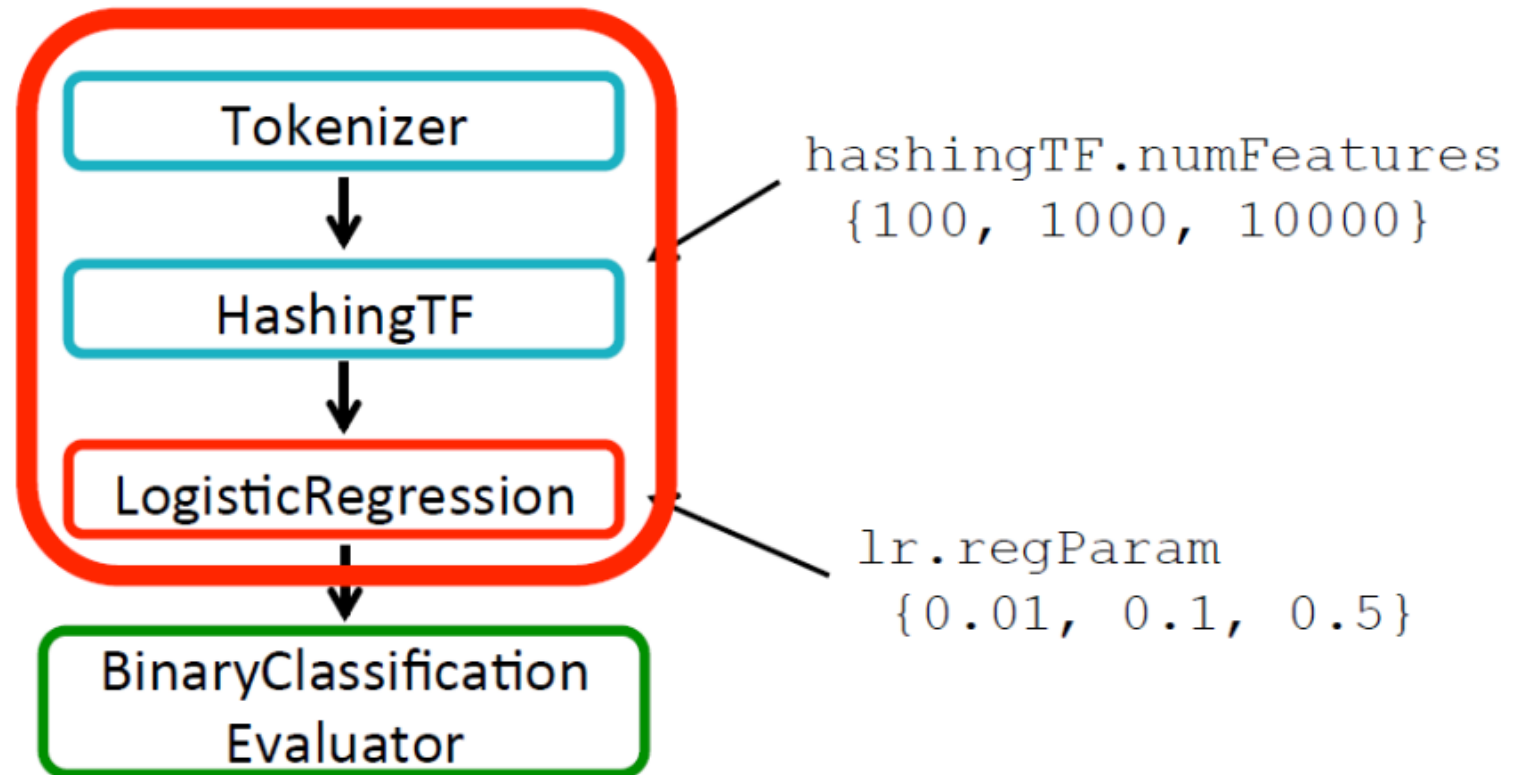


Given:

- Estimator
- Parameter Grid
- Evaluator

Find the Best parameters

**CrossValidator**



- Suppose we have a dataset in which each line has a recording of a noise, and its classification
- E.g., <bell.wav>, bell



- What would be the input DataFrame be?
- Suppose we want to recognize sounds by
  - Extracting the frequencies from the wav file
  - Gaussian model
    - Find the average frequency of each sound
    - For a new sound, calculate average frequency
    - Find closest matching sound
- What are the DataFrames, Evaluators, etc needed.



- Suppose we have a dataset in which each line has a recording of a noise, and its classification
- E.g., <bell.wav>, bell

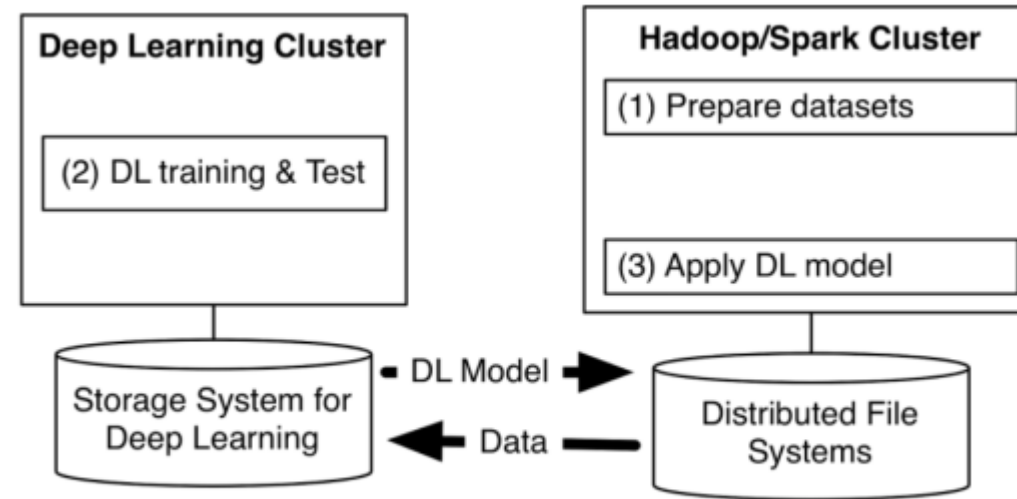
- Input DataFrame
  - <bell.wav>, bell
- Feature DataFrame
  - <bell.wav>, bell, frequencies
- Transformer (use same transformer for train/predict)
  - <bell.wav> Bell, average frequency
- Model
  - train(FeatureDataFrame)
    - Associate average frequency for “Bell”
  - predict(PredictDataFrame)
    - Output closest matching sound

- Classification
  - Logistic Regression
  - Decision Tree
  - Random Forest
  - Gradient boosted tree
  - Multilayer Perceptron
  - SVM
  - Naïve Bayes
- Clustering
  - K-Means
  - LDA
  - GMM
- Collaborative Filtering
  - ALS
- Frequent Pattern Mining

# Deep Learning with Big Data

- Heterogenous cluster
- Deep Learning (Tensorflow)
  - Iterative
  - Matrix vector multiplication – Linear algebra
- Initially evolved on a single machine – only scale up
- Then had its own cluster
  - Typically heterogenous with CPUs, GPUs, TPUs

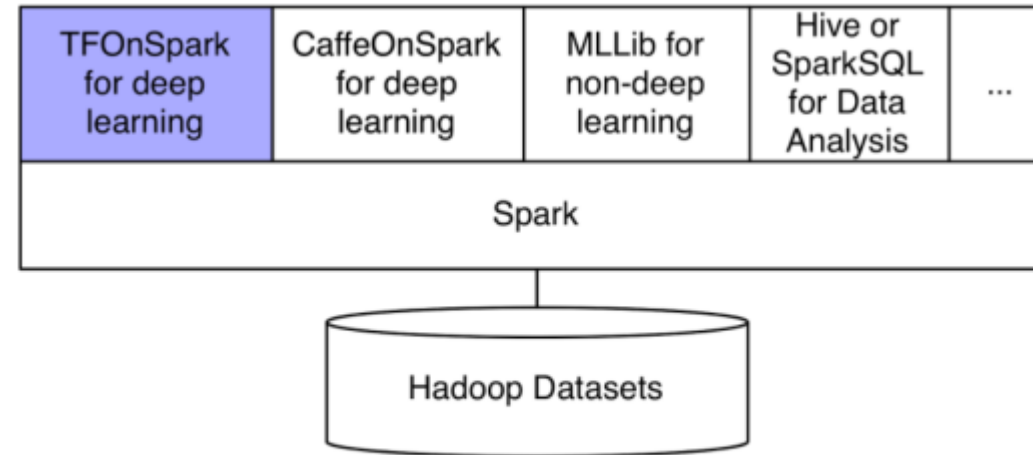
- But data resides on HDFS and big data platform uses Spark
- How should the two work together.
- Typically the two clusters are different



**Figure 1:** ML Pipeline with multiple programs on separated clusters

<https://developer.yahoo.com/blogs/157196317141/>

- Can we use the same cluster?
- Tensorflow on Spark
  - From Yahoo



**Figure 2:** TensorFlowOnSpark for deep learning on Spark clusters

<https://developer.yahoo.com/blogs/157196317141/>

- Supports both
  - Model parallelism
  - Data parallelism
- <10 lines of code change reqd
- Algorithm and parameter server run on Spark executors
  - Can read data directly from HDFS
  - Spark RDD data is fed to spark executor which passes it to Tensorflow
- RDMA → faster network transfers

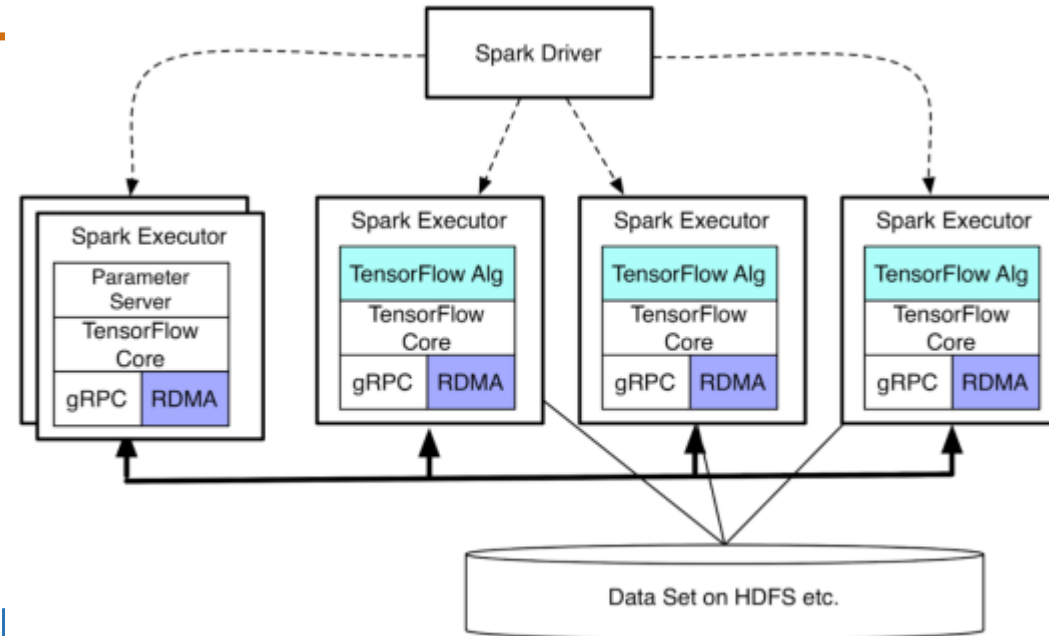


Figure 3: TensorFlowOnSpark system architecture

<https://developer.yahoo.com/blogs/157196317141/>

- SystemML – [systemml.apache.org](http://systemml.apache.org)
  - IBM
  - SystemML: Declarative Machine Learning on Spark  
<http://www.vldb.org/pvldb/vol9/p1425-boehm.pdf>
  - Uses a declarative ML language
  - Translated to MR/Spark
- Intel BigDL
  - Modeled on Torch





# THANK YOU

---

**K V Subramaniam, Usha Devi**

Dept. of Computer Science and Engineering

[subramaniamkv@pes.edu](mailto:subramaniamkv@pes.edu)

**ushadevibg@pes.edu**