



# Unix System Programming

## File I/O

---

**Chandravva Hebbi**

**Department of Computer Science and Engineering**

**`chandravvahebbi@pes.edu`**

# Unix System Programming

## File I/O

---



**Chandravva Hebbi**

Department of Computer Science and Engineering

# UNIX SYSTEM PROGRAMMING

## Topics to be Covered

---



- ❖ Read and write system calls.
- ❖ Programming Examples on Basic system calls.

```
ssize_t read(int fd, void *buf, size_t count);
```

- ❖ `read()` attempts to read up to `count` bytes from file descriptor `fd` into the buffer starting at `buf`.
- ❖ If `count` is zero, `read()` returns zero
- ❖ If `count` is greater than `SSIZE_MAX`, the result is unspecified.
- ❖ On error, `-1` is returned. `Errno` is set, check the error no.
- ❖ Error: `EBADF`

```
ssize_t write(int fd, const void *buf, size_t count);
```

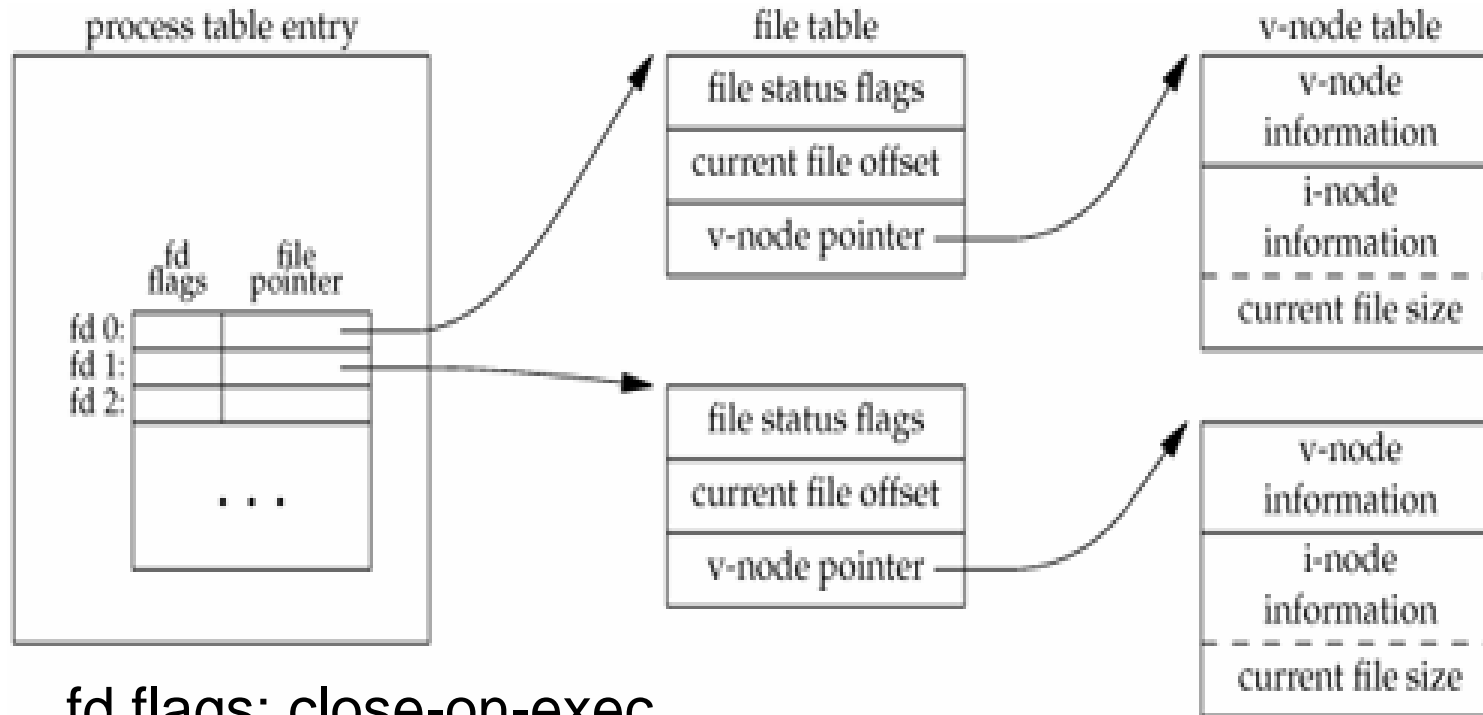
- ❖ writes up to count bytes to the file referenced by the file descriptor fd from the buffer starting at buf.
- ❖ On success, the number of bytes written are returned.
- ❖ Write: current position in a file or to the end if O\_APPEND
- ❖ On error, -1 is returned, and errno is set appropriately.

Most file systems support read-ahead to improve performance

Sequential reads – system tries to read in more data than an application requests

# UNIX SYSTEM PROGRAMMING

## Kernel Data Structures

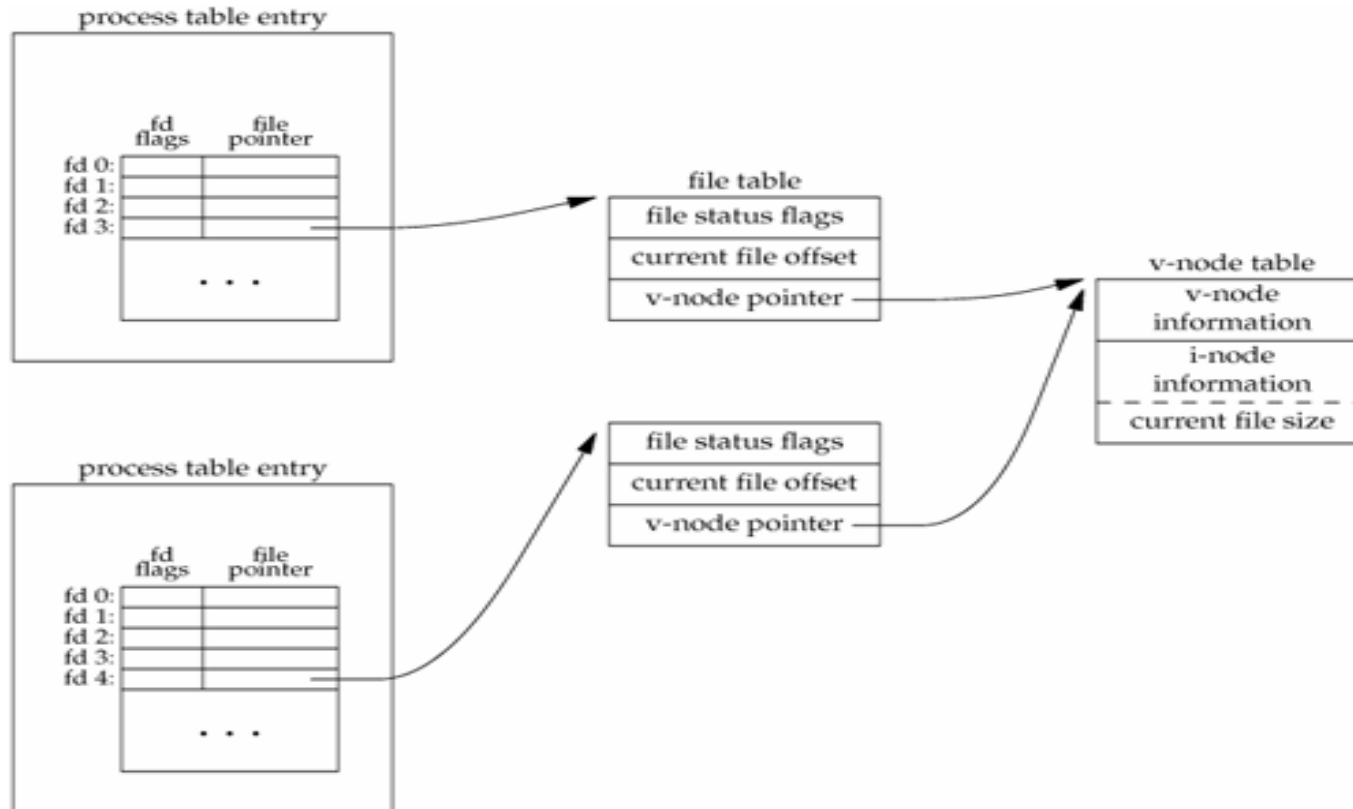


fd flags: close-on-exec  
Value: 0,1

v-node structure that contains information about the type of file and pointers to functions that operate on the file.

# UNIX SYSTEM PROGRAMMING

## File Sharing



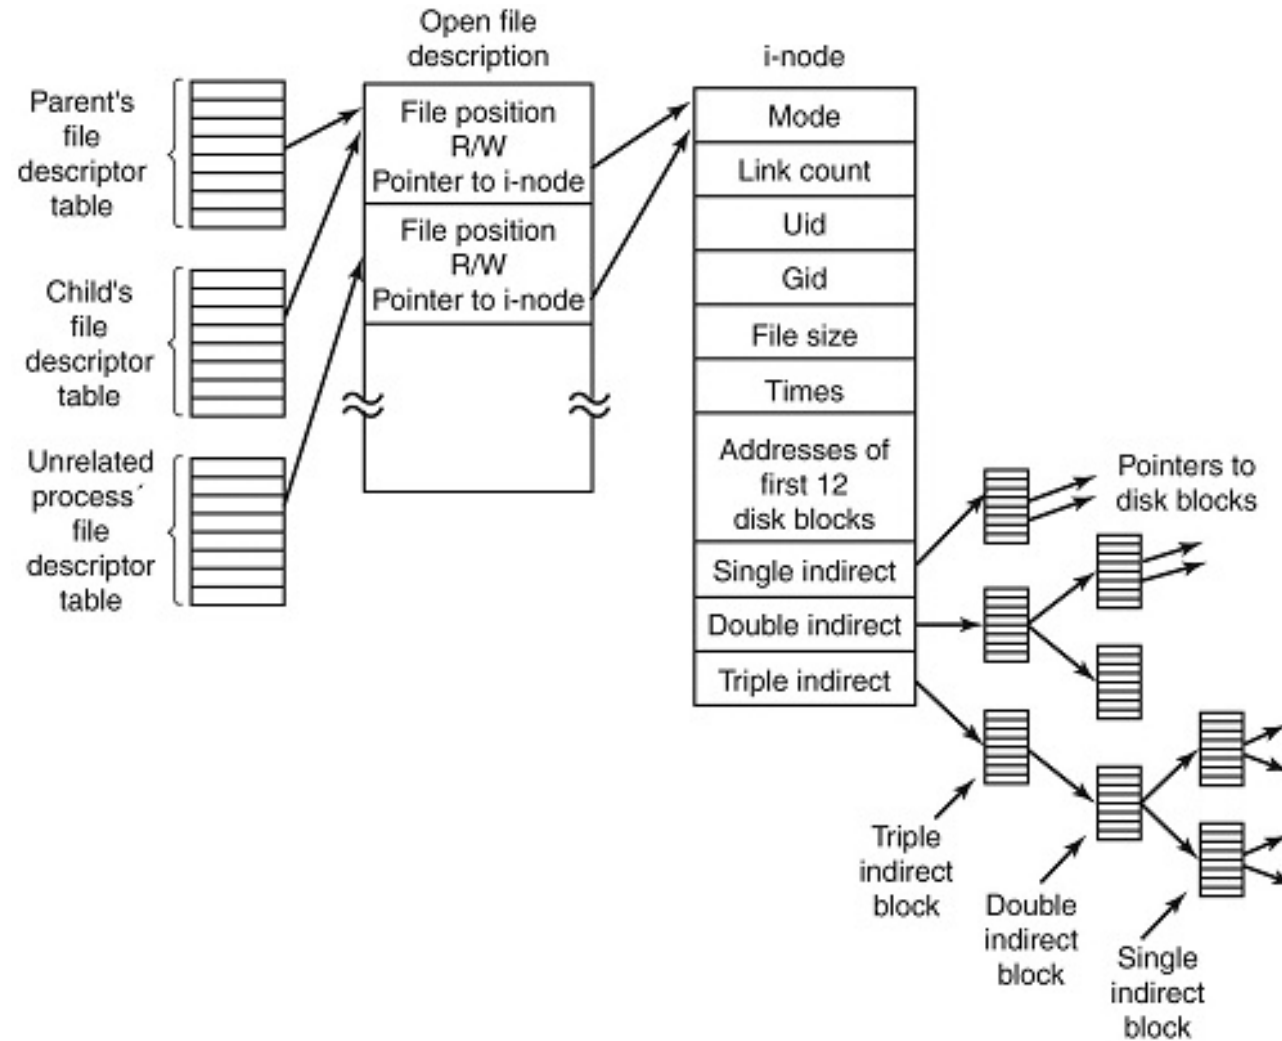


# UNIX SYSTEM PROGRAMMING

## i-node Structure



**PES**  
UNIVERSITY  
ONLINE



Source: <https://practice.geeksforgeeks.org/problems/explain-the-structure-of-inode-table>

```
if (lseek(fd, 0L, 2) < 0)          /* position to EOF */  
    err_sys("lseek error");  
if (write(fd, buf, 100) != 100)    /* and write */  
    err_sys("write error");
```

What happens if,

Case 1: single process access the file with descriptor,fd

Case 2: Multiple processes appending to the same file.

# UNIX SYSTEM PROGRAMMING

## Atomic Operations

---

The Single UNIX Specification includes XSI extensions that allow applications to seek and perform I/O atomically.

These extensions are `pread` and `pwrite`.

**`ssize_t pread(int fildes, void *buf, size_t nbytes, off_t offset);`** returns:

number of bytes read, 0 if end of file, -1 on error

**`ssize_t pwrite(int fildes, const void *buf, size_t nbytes, off_t offset);`**

Returns: number of bytes written if OK, -1 on error



Calling `pread` is equivalent to calling `lseek` followed by a call to `read`, with the following exceptions.

*There is no way to interrupt the two operations using `pread`. The file pointer is not updated.*

Calling `pwrite` is equivalent to calling `lseek` followed by a call to `write`, with similar exceptions.

# UNIX SYSTEM PROGRAMMING

## chown() - Unix, Linux System Call

---

```
int chown(const char *path, uid_t owner, gid_t group);  
int fchown(int fd, uid_t owner, gid_t group);  
int lchown(const char *path, uid_t owner, gid_t group);
```



The fcntl function can change the properties of a file that is already open.

***int fcntl(int filedes, int cmd, ... /\* int arg \*/ );***

Returns: depends on cmd if OK (see following), -1 on error system calls.

The fcntl function is used for five different purposes.

- ❖ Duplicate an existing descriptor (cmd = F\_DUPFD)
- ❖ Get/set file descriptor flags (cmd = F\_GETFD or F\_SETFD)
- ❖ Get/set file status flags (cmd = F\_GETFL or F\_SETFL)
- ❖ Get/set asynchronous I/O ownership (cmd = F\_GETOWN or F\_SETOWN)
- ❖ Get/set record locks (cmd = F\_GETLK, F\_SETLK, or F\_SETLKW)

F\_DUPFD: Duplicate the file descriptor `filedes`.

F\_GETFD: Return the file descriptor flags for `filedes` as the value of the function. Currently, only one file descriptor flag is defined: the `FD_CLOEXEC` flag.

F\_SETFD: Set the file descriptor flags for `filedes`. The new flag value is set from the third argument

F\_GETFL: Return the file status flags for `filedes` as the value of the function.

- O\_RDONLY: open for reading only
- O\_WRONLY: open for writing only
- O\_RDWR: open for reading and writing
- O\_APPEND: append on each write
- O\_NONBLOCK: nonblocking mode
- O\_SYNC: wait for writes to complete (data and attributes)
- O\_DSYNC: wait for writes to complete (data only)
- O\_RSYNC: synchronize reads and writes
- O\_FSYNC: wait for writes to complete (FreeBSD and Mac OS X only)
- O\_ASYNC: asynchronous I/O (FreeBSD and Mac OS X only)



**F\_SETFL:** Set the file status flags to the value of the third argument (taken as an integer).

The only flags that can be changed are `O_APPEND`, `O_NONBLOCK`, `O_SYNC`, `O_DSYNC`, `O_RSYNC`, `O_FSYNC`, and `O_ASYNC`.

**F\_GETOWN:** Get the process ID or process group ID currently receiving the `SIGIO` and `SIGURG` signals.

**F\_SETOWN:** Set the process ID or process group ID to receive the `SIGIO` and `SIGURG` signals.



**THANK YOU**

---

**Chandravva Hebbi**

Department of Computer Science and Engineering

**[chandravvahebbi@pes.edu](mailto:chandravvahebbi@pes.edu)**