



Unix System Programming

Process Control

Chandravva Hebbi
Department of Computer Science and Engineering
chandravvahebbi@pes.edu

Unix System Programming

Process Control



Chandravva Hebbi

Department of Computer Science and Engineering

UNIX SYSTEM PROGRAMMING

Topics to be Covered



- ❖ Process Identifiers

- ❖ fork()

- ❖ vfork()

- ❖ exec()

- ❖ Programming examples

- Every process has a unique process ID, a non-negative integer.
- Identifier of a process is always unique.
- process IDs are reused.
- Delay in reuse.
- special processes: Process ID 0 is a swapper process(scheduler Process).
- Process ID 1: init it is invoked by the kernel at the end of bootstrap.

- The program file for this process was `/etc/init` in older versions of the UNIX System and `/sbin/init` in newer versions.
- This process is responsible for bringing up a UNIX system after the kernel has been bootstrapped.
- `init` usually reads the system-dependent initialization files the `/etc/rc*` files or `/etc/inittab` and the files in `/etc/init.d` and brings the system to a certain state, such as multiuser.
- The `init` process never dies. It is a normal user process, not a system process within the kernel, like the swapper.
- It runs with superuser privileges.

- process ID 2 is the *pagedaemon*. It supports paging for Virtual memory system
- Other processes and their id's

```
#include <unistd.h>
```

```
pid_t getpid(void);
```

Returns: process ID of calling process

```
pid_t getppid(void);
```

Returns: parent process ID of calling process

```
uid_t getuid(void);
```

Returns: real user ID of calling process

```
uid_t geteuid(void);
```

Returns: effective user ID of calling process

```
gid_t getgid(void);
```

Returns: real group ID of calling process

```
gid_t getegid(void);
```

Returns: effective group ID of calling process

Note that none of these functions has an error return

- An existing process can create a new one by calling the **fork** function.

```
#include <unistd.h>
```

```
pid_t fork(void);
```

Returns: 0 in child, process ID of child in parent, -1 on error

- `fork()` creates a new process called child.
- Returns twice: 0 in child and child id in parent
- No user process with ID 0

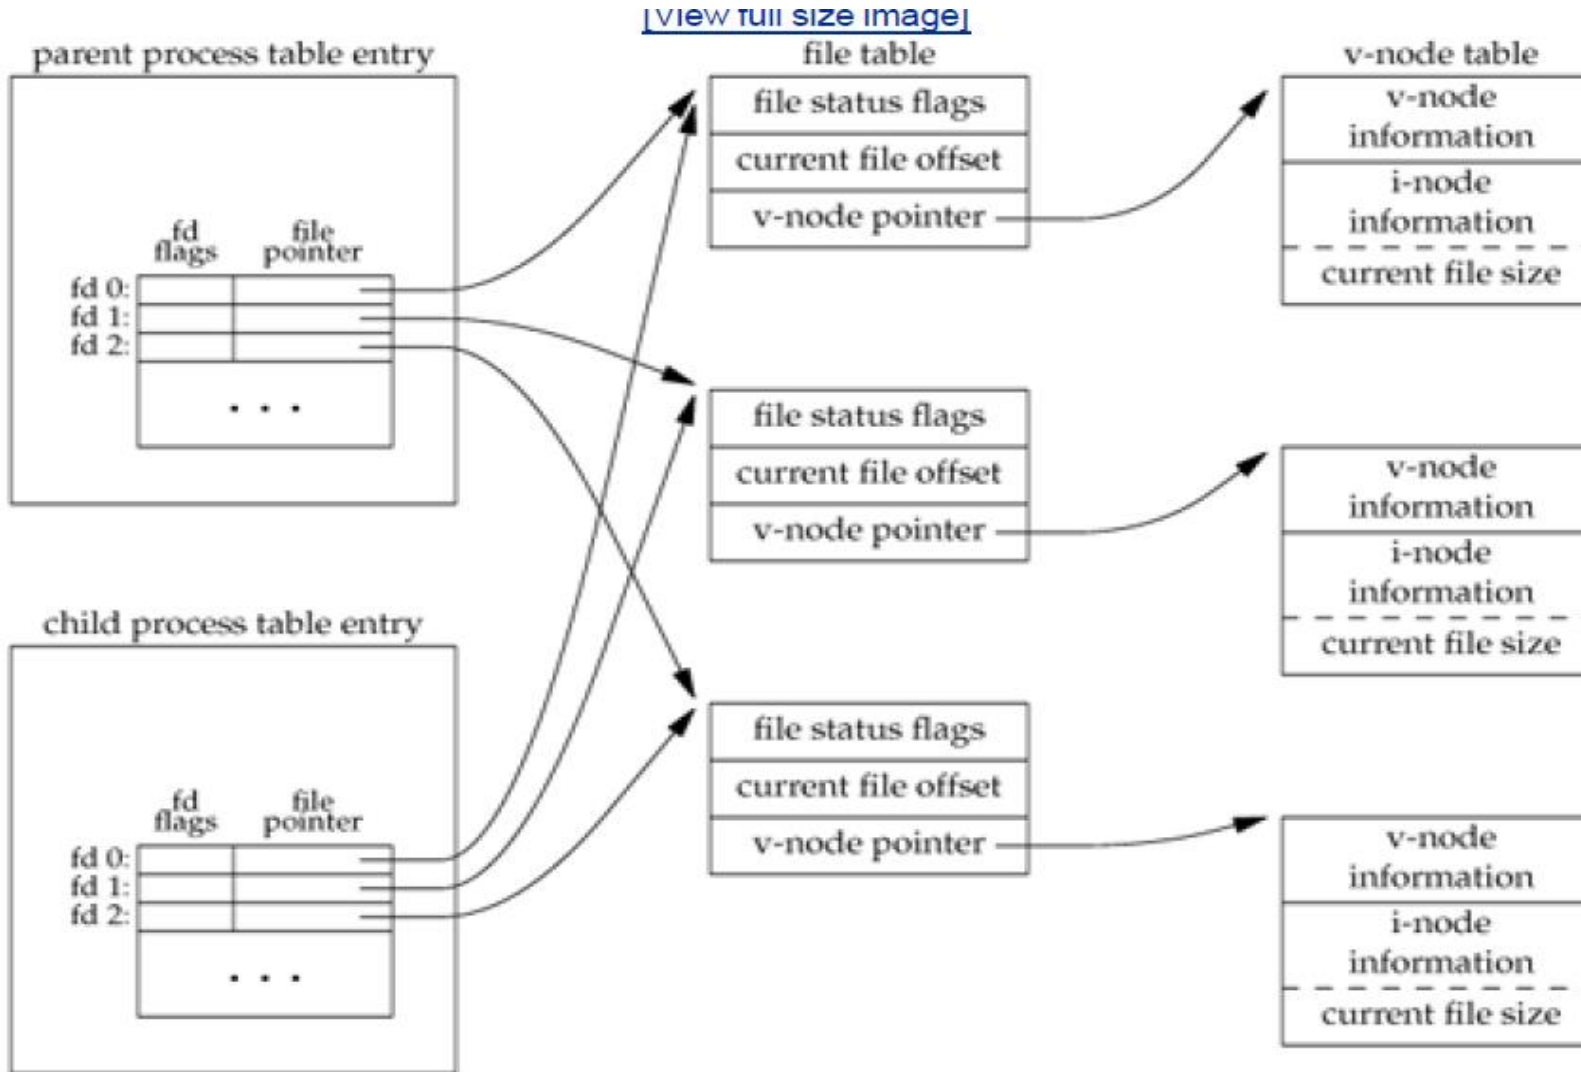
- The child is a copy of the parent. The child gets a copy of the parent's data space, heap, and stack.
- The parent and the child do not share these portions of memory.
- The parent and the child share the text segment.
- Current implementations don't perform a complete copy of the parent's data, stack, and heap, since a fork is often followed by an exec.
- Instead, a technique called copy-on-write (COW) is used.
- These regions are shared by the parent and the child and have their protection changed by the kernel to read-only.
- Copy is given when one of the process tries to modify.

Variations in the fork function

- Linux 2.4.22 also provides new process creation through the **clone**(2) system call
- FreeBSD 5.2.1 provides the **rfork**(2) system call.
- Solaris provides the **fork1** function: creates a process that duplicates only the calling thread.

UNIX SYSTEM PROGRAMMING

Fork function: File Sharing



Two normal cases for handling the descriptors after a **fork**.

- The parent waits for the child to complete
- Both the parent and the child go their own ways.

UNIX SYSTEM PROGRAMMING

fork function: File Sharing

Properties of the parent that are inherited by the child:

Real user ID, real group ID, effective user ID, effective group ID

Supplementary group IDs

Process group ID

Session ID

Controlling terminal

The set-user-ID and set-group-ID flags

Current working directory

Root directory

File mode creation mask

Signal mask and dispositions

The close-on-exec flag for any open file descriptors

Environment

Attached shared memory segments

Memory mappings

Resource limits



The differences between the parent and child are.

- The return value from `fork`
- The process IDs are different
- The two processes have different parent process IDs: the parent process ID of the child is the
parent; the parent process ID of the parent doesn't change
- The child's `tms_utime`, `tms_stime`, `tms_cutime`, and `tms_cstime` values are set to 0
- File locks set by the parent are not inherited by the child
- Pending alarms are cleared for the child
- The set of pending signals for the child is set to the empty set

Reasons for fork failure

- (a) if too many processes are already in the system
- (b) if the total number of processes for this real user ID exceeds the system's limit.

Uses for **fork**

Application 1: Networks

- When a process wants to duplicate itself so that the parent and child can each execute different sections of code at the same time.
- Application: In network servers the parent waits for a service request from a client.
- When the request arrives, the parent calls **fork** and lets the child handle the request.
- The parent goes back to waiting for the next service request to arrive.

Application 2: shells

- When a process wants to execute a different program.
- This is common for shells.
- In this case, the child does an **exec** right after it returns from the **fork**.



THANK YOU

Chandravva Hebbi

Department of Computer Science and Engineering

chandravvahebbi@pes.edu