



# **DATA ANALYTICS**

## **Unit 3: Concept of stationarity, DF and ADF test, transformations, ARIMA and SARIMA**

---

**Jyothi R.**

Department of Computer Science  
and Engineering

## Auto Regressive Integrated Moving Average (ARIMA) Process

ARIMA model was proposed by Box-Jenkins (1970)

and so known as [Box-Jenkins Methodology](#)

It has three components and is represented as  $ARIMA(p,d,q)$ :


1. Auto-regressive with lag  $p$
2. Integration component ( $d$ )
3. Moving average ( $q$ )

Integration component's objective: To convert a nonstationary signal to stationary

Nonstationarity could arise from deterministic or stochastic trend

## Identifying nonstationarity - ACF

- ACF will not cut off to zero quickly; may show a very slow decline



## Quantitative Test - Dickey Fuller (DF) Test

---

Consider AR(1) process defined below:

$$Y_{t+1} = \beta Y_t + \varepsilon_{t+1}$$

In Section 13.11, we proved that the AR(1) process can become very large when  $\beta > 1$  and is non-stationary when  $|\beta| = 1$ . Dickey–Fuller test (Dickey and Fuller, 1979) is a hypothesis test in which the null hypothesis and alternative hypothesis are given by

$H_0: \beta = 1$  (the time series is non-stationary)

$H_1: \beta < 1$  (the time series is stationary)

The AR(1) can be written as

$$Y_{t+1} - Y_t = \Delta Y_t = (\beta - 1)Y_t + \varepsilon_{t+1} = \psi Y_t + \varepsilon_{t+1} \quad (13.46)$$

## Dickey Fuller Test

---

- $\psi = 0$  is same as  $\beta = 1$ . So, the Dickey–Fuller test can be written in terms of  $\psi$  as

$H_0: \psi = 0$  (the time series is non-stationary)

$H_A: \psi < 0$  (the time series is stationary)

- The test statistic is given by

- DF Test Statistic = 
$$\frac{\psi}{S_e(\psi)}$$

- where  $S_e$  is the standard error. Note that DF test statistic is not  $t$ -statistic since the null hypothesis is on non-stationary process.
- Critical values are derived based on simulation

## Augmented Dickey–Fuller Test

- Dickey–Fuller test is valid only when the residual  $\varepsilon_{t+1}$  follows a white noise.
- When  $\varepsilon_{t+1}$  is not white noise, the actual series may not be AR(1); it may have more significant lags.
- To address this issue, we augment  $p$ -lags of the dependent variable  $Y$ .
- The model can be rewritten as:

$$\Delta Y_t = \psi Y_t + \sum_{i=0}^p \alpha_i \Delta Y_{t-i} + \varepsilon_{t+1}$$

- The above equation can be now tested for non-stationarity.
- Again the null and alternative hypotheses are
- $H_0: \psi = 0$  (the time series is non-stationary)
- $H_0: \psi < 0$  (the time series is stationary)

### Transforming Non-Stationary Process to Stationary Process Using Differencing

- The first step in ARIMA is to identify the order of differencing ( $d$ ) required to convert a non-stationary process into a stationary process.
- Many time-series data will be non-stationary due to factors such as trend and seasonality.
- If the non-stationary behaviour is due to trend, then it can be converted into a stationary process by de-trending the data.
- De-trending is usually achieved by fitting a trend line and subtracting it from the time series; this is known as **trend stationarity**.
- When the reason is not due to trend stationarity, then differencing the original time series may be useful for converting the non-stationary process into a stationary process (called **difference stationarity**).

### Transforming Non-Stationary Process to Stationary Process Using Differencing

- The first difference ( $d = 1$ ) is the difference between consecutive values of the time series ( $Y_t$  and  $Y_{t-1}$ )
- That is, the first difference  $\Delta Y_t$  is given by
$$\Delta y_t = Y_t - Y_{t-1}$$
- The second difference ( $d = 2$ ) is the difference of the first differences and is given by
$$\nabla^2 Y_t = \nabla(\nabla Y_t) = Y_t - 2Y_{t-1} + Y_{t-2}$$
- In most cases,  $d \leq 2$  will be sufficient to convert a non-stationary process to a stationary process.



## Stationarity and differencing

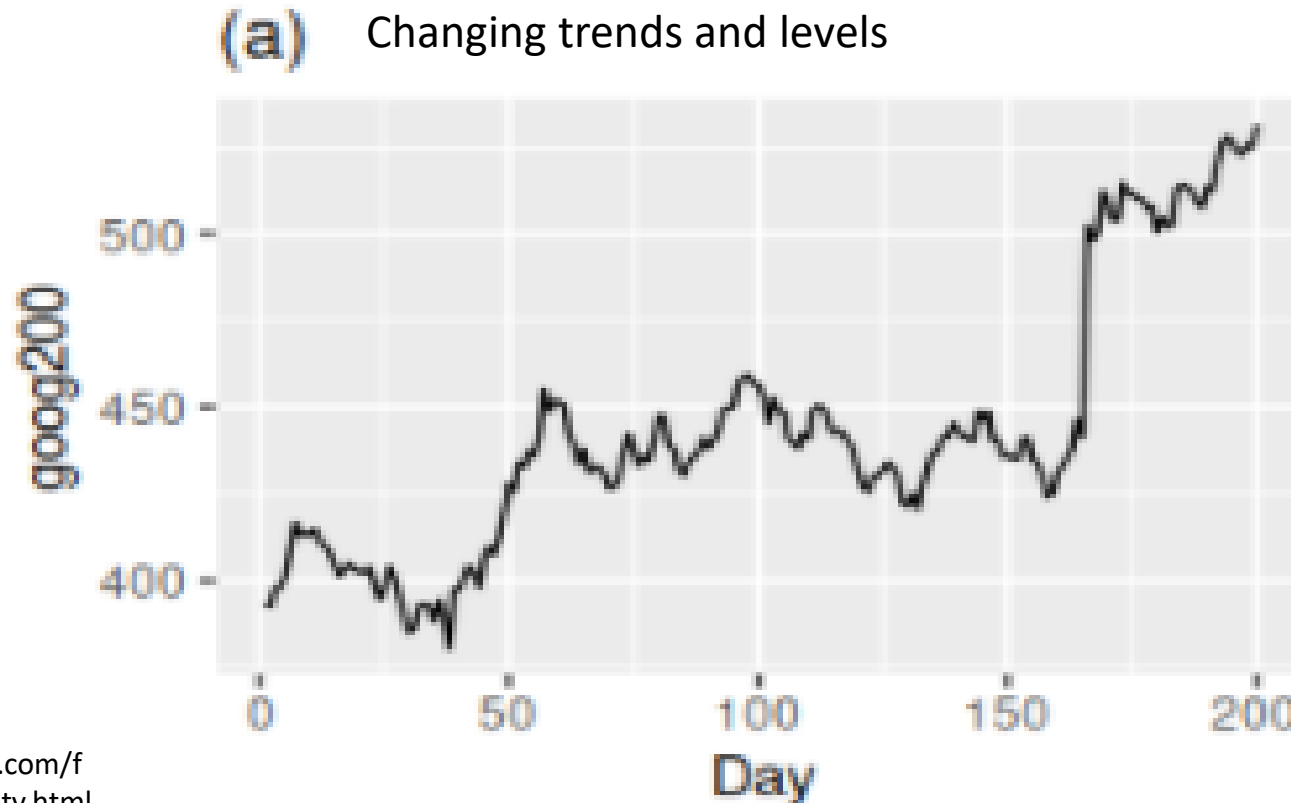
---

- A stationary time series is one whose properties do not depend on the time at which the series is observed
- Time series with trends, or with seasonality, are not stationary — the trend and seasonality will affect the value of the time series at different times.
- On the other hand, a white noise series is stationary — it does not matter when you observe it, it should look much the same at any point in time.
- Some cases can be confusing — a time series with cyclic behaviour but with no trend or seasonality is stationary.
- This is because the cycles are not of a fixed length, so before we observe the series we cannot be sure where the peaks and troughs of the cycles will be.
- In general, a stationary time series will have no predictable patterns in the long-term.

# DATA ANALYTICS

## Which of these series are stationary?

- Time plots will show the series to be roughly horizontal although some cyclic behaviour is possible, with constant variance.



<https://otexts.com/fpp2/stationarity.html>

Figure1. (a) Google stock price for 200 consecutive days

## DATA ANALYTICS

Which of these series are stationary?





Figure 2: (b) Daily change in the Google stock price for 200 consecutive days;

## Which of these series are stationary?

(c) Changing trends and levels




<https://otexts.com/fpp2/stationarity.html>

Figure 3 :(c) Annual number of strikes in the US

# DATA ANALYTICS

## Which of these series are stationary?




<https://otexts.com/fpp2/stationarity.html>

Figure 4: (d) Monthly sales of new one-family houses sold in the US

## Which of these series are stationary?

(e) Changing trends and levels




<https://otexts.com/fpp2/stationarity.html>

Figure 5: (e) Annual price of a dozen eggs in the US (constant dollars);

## DATA ANALYTICS

Which of these series are stationary?

(f) Changing trends and levels



<https://otexts.com/fpp2/stationarity.html>

Figure 6: (f) Monthly total of pigs slaughtered in Victoria, Australia;

## DATA ANALYTICS

Which of these series are stationary?




Figure 7: (g) Annual total of lynx trapped in the McKenzie River district of north-west Canada;



## Which of these series are stationary?

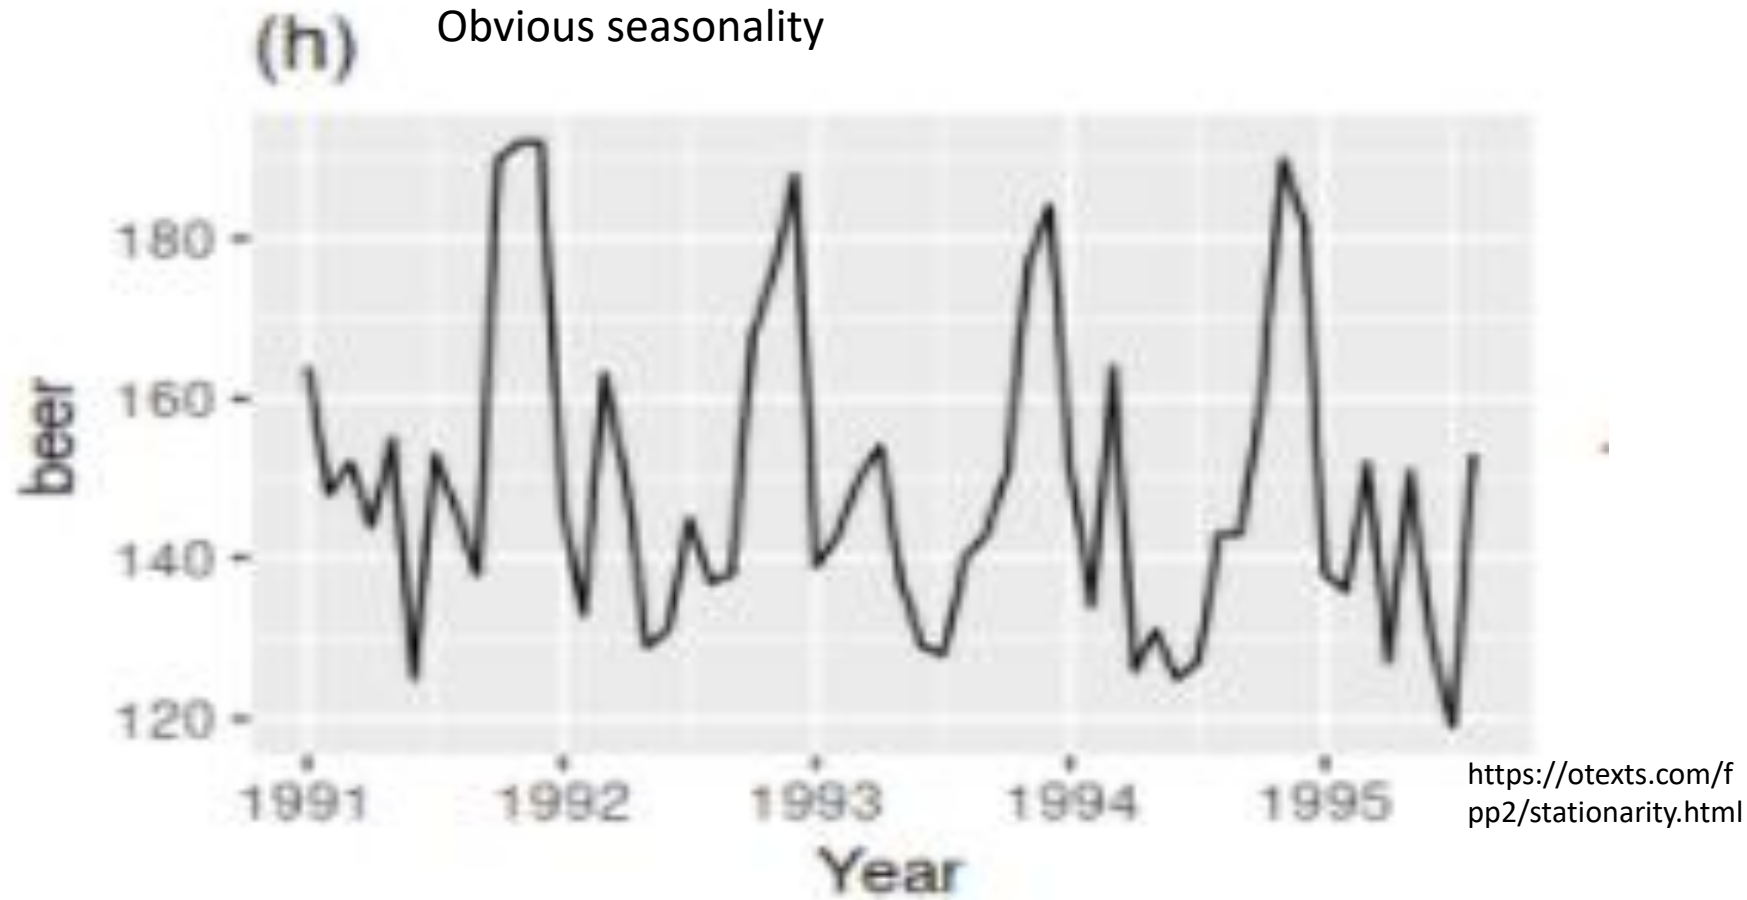


Figure 8: (h) Monthly Australian beer production;

# DATA ANALYTICS

Which of these series are stationary?

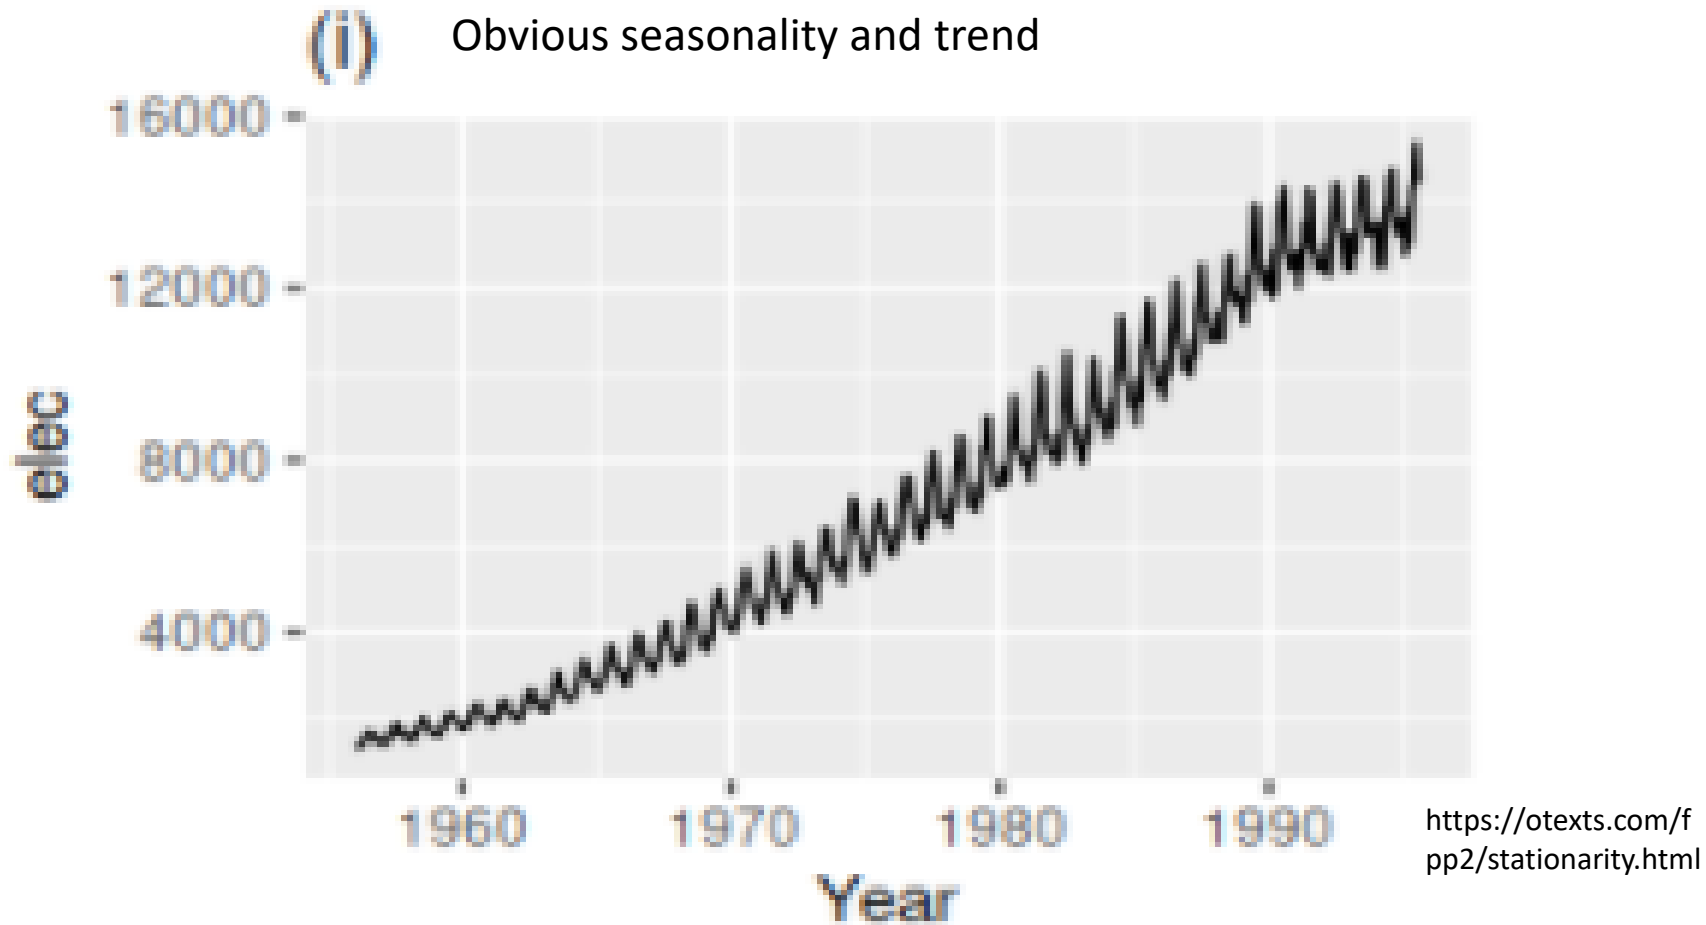


Figure 9:(i) Monthly Australian electricity production.

## Which of these series are stationary?


---

Consider the nine series plotted in Figure 1 to 9:

- Which of these do you think are stationary?
- Obvious seasonality rules out series (d), (h) and (i).
- Trends and changing levels rules out series (a), (c), (e), (f) and (i).
- Increasing variance also rules out (i).
- That leaves only (b) and (g) as stationary series.
- At first glance, the strong cycles in series (g) might appear to make it non-stationary.
- But these cycles are **aperiodic** — they are caused when the lynx population becomes too large for the available feed, so that they stop breeding and the population falls to low numbers, then the regeneration of their food sources allows the population to grow again, and so on.
- In the long-term, the timing of these cycles is not predictable. Hence the series is stationary.

- In Figure 1 to 9: The Google stock price was non-stationary in panel (a)
- But the daily changes were stationary in panel (b). This shows one way to make a non-stationary time series stationary — compute the differences between consecutive observations. This is known as **differencing**.
- Transformations such as [logarithms can help to stabilise the variance](#) of a time series.
- [Differencing can help stabilise the mean of a time series](#) by removing changes in the level of a time series, and therefore eliminating or reducing trend and seasonality.
- By looking at the time plot of the data, the ACF plot is also useful for identifying non-stationary time series.
- For a stationary time series, the ACF will drop to zero relatively quickly, while the ACF of non-stationary data decreases slowly.
- Also, for non-stationary data, the value of  $r_1$  is often large and positive.

Series: goog200



Series: diff(goog200)





Figure 10: The ACF of the Google stock price (left) and of the daily changes in Google stock price (right).


Figure 8.2: The ACF of the Google stock price (left) and of the daily changes in Google stock price (right).

- The ACF of the differenced Google stock price looks just like that of a white noise series.
- There are no autocorrelations lying outside the 95% limits
- This suggests that the *daily change* in the Google stock price is essentially a random amount which is uncorrelated with that of previous days.


Level



First-difference



Second-difference



- The differenced series is the *change* between consecutive observations in the original series, and can be written as
- $y'_t = y_t - y_{t-1}$ .
- The differenced series will have only  $T-1$  values, since it is not possible to calculate a difference  $y'_1$  for the first observation.
- When the differenced series is white noise, the model for the original series can be written as  $y_t - y_{t-1} = \varepsilon_t$ , where  $\varepsilon_t$  denotes white noise.
- Rearranging this leads to the “random walk” model  $y_t = y_{t-1} + \varepsilon_t$



- Random walk models are widely used for non-stationary data, particularly financial and economic data.

Random walks typically have:

- long periods of apparent trends up or down
- sudden and unpredictable changes in direction.
- The forecasts from a random walk model are equal to the last observation, as future movements are unpredictable, and are equally likely to be up or down.
- Thus, the random walk model underpins naïve forecasts,

- A closely related model allows the differences to have a non-zero mean.

Then

- $y_t - y_{t-1} = c + \varepsilon_t$  or  $y_t = c + y_{t-1} + \varepsilon_t$ .
- $y_t - y_{t-1} = c + \varepsilon_t$  or  $y_t = c + y_{t-1} + \varepsilon_t$ .
- The value of  $c$  is the average of the changes between consecutive observations.
- If  $c$  is positive, then the average change is an increase in the value of  $y_t$ .
- Thus,  $y_t$  will tend to drift upwards.
- However, if  $c$  is negative,  $y_t$  will tend to drift downwards.
- This is the model behind the drift method

## What ARIMA stands for

---

- A series which needs to be differenced to be made stationary is an “integrated” (**I**) series
- Lags of the stationarized series are called “auto- regressive” (**AR**) terms
- Lags of the forecast errors are called “moving average” (**MA**) terms
- We’ve already studied these time series tools separately: differencing, moving averages, lagged values of the dependent variable in regression

## ARIMA models put it all together

---

- Generalized random walk models: fine-tuned to eliminate all residual autocorrelation
- Generalized exponential smoothing models: that can incorporate long-term trends and seasonality
- Stationarized regression models: that use lags of the dependent variables and/or lags of the forecast errors as regressors.
- A general class of forecasting models for time series that can be stationarized by transformations such as differencing, logging, and or deflating.

## ARIMA( $p, d, q$ ) Model Building

---

- The first step in ARIMA( $p, d, q$ ) is the model identification, that is, identifying the values of  $p$ ,  $d$ , and  $q$ .
- Box and Jenkins (1970) proposed the following procedure to build the ARIMA( $p, d, q$ ) model.
- The main objective of model identification stage is to identify the right values of
  - $p$  (auto-regressive lags),
  - $d$  (order of differencing), and
  - $q$  (moving average lags).

## ARIMA( $p, d, q$ ) Model Building

---

- The following flow chart can be used during the model identification stage
- The first step is to plot the ACF and PACF to identify whether the time series is stationary or not.
- If the time series is stationary then  $d = 0$  and
- the model is ARIMA( $p, 0, q$ ) or ARMA( $p, q$ ) model.
- If the time series is non-stationary then it has to be converted into a stationary
- process by identifying the order of differencing.
- Once the value of  $d$  is known that will make the
- process stationary, then  $p$  and  $q$  are identified for the stationary process.

## ARIMA( $p, d, q$ ) Model Building

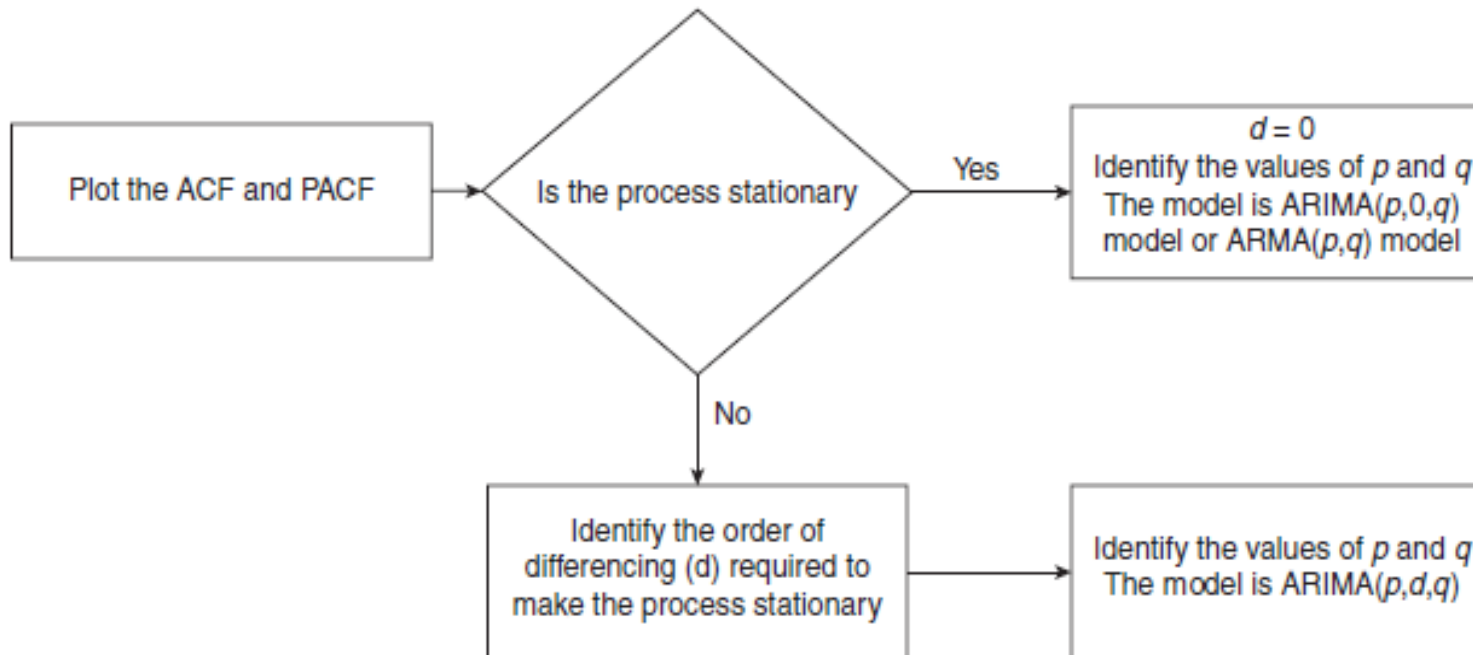


FIGURE 13.14 Model identification in ARIMA model.

## Parameter Estimation and Model Selection

---

- Once the model is identified (values of  $p$ ,  $d$ , and  $q$ ),
- the next step in ARIMA model building is the parameter estimation.
- That is, the estimation of coefficients in AR and MA components which are
- achieved using ordinary least squares.
- The model selection may be carried using several criteria such as RMSE, MAPE, Akaike Information Criteria (AIC), or Bayesian Information Criteria (BIC).
- AIC and BIC are measures of distance from the actual values to the forecasted values.



## Parameter Estimation and Model Selection

---



- AIC is given by  $AIC = -2LL + 2K$

where  $LL$  is the log likelihood function and  $K$  is the number of parameters estimated (in this case  $p + q$ ).

- BIC is given by  $BIC = -2LL + K \ln(n)$


In BIC equation,  $n$  is the number of observations in the sample. BIC assigns higher penalty compared to AIC for every additional variable added to the model.

Lower values of AIC and BIC are preferred.

- ARIMA model is a regression model and thus has to satisfy all the assumptions of regression.
- The residual should be white noise. We can also perform a [goodness of fit test using Ljung–Box test](#) (coming up tomorrow!) before accepting the model.

# DATA ANALYTICS

## The ARIMA “filtering box”



Objective: adjust the knobs until the residuals are “white noise” (uncorrelated)

# DATA ANALYTICS

## ARIMA models we have already met

---



- $\text{ARIMA}(0,0,0)+c$  = mean (constant) model
- $\text{ARIMA}(0,1,0)$  = RW model
- $\text{ARIMA}(0,1,0)+c$  = RW with drift model
- $\text{ARIMA}(1,0,0)+c$  = regress  $Y$  on  $Y_{\text{LAG}1}$
- $\text{ARIMA}(1,1,0)+c$  = regr.  $Y_{\text{DIFF}1}$  on  $Y_{\text{DIFF}1\_LAG1}$
- $\text{ARIMA}(2,1,0)+c$  = " " plus  $Y_{\text{DIFF\_LAG}2}$  as well
- $\text{ARIMA}(0,1,1)$  = SES model
- $\text{ARIMA}(0,1,1)+c$  = SES + constant linear trend
- $\text{ARIMA}(1,1,2)$  = LES w/ damped trend (leveling off)
- $\text{ARIMA}(0,2,2)$  = generalized LES (including Holt's)

## ARIMA forecasting equation

---

- Let  $Y$  denote the original series
- Let  $y$  denote the differenced (stationarized) series

No difference ( $d=0$ ):  $y_t = Y_t$

First difference ( $d=1$ ):  $y_t = Y_t - Y_{t-1}$

Second difference ( $d=2$ ):  $y_t = (Y_t - Y_{t-1}) - (Y_{t-1} - Y_{t-2})$   
 $= Y_t - 2Y_{t-1} + Y_{t-2}$

## Forecasting equation for $y$

Not as bad as it looks! Usually  $p+q \leq 2$  and either  $p=0$  or  $q=0$  (pure AR or pure MA model)

$$\hat{y}_t = \underbrace{\mu}_{\text{constant}} + \underbrace{\phi_1 y_{t-1} + \dots + \phi_p y_{t-p}}_{\text{AR terms (lagged values of } y)} - \underbrace{\theta_1 e_{t-1} \dots - \theta_q e_{t-q}}_{\text{MA terms (lagged errors)}}$$

## Undifferencing the forecast

---

- The differencing (if any) must be reversed to obtain a forecast for the original series:

$$\text{If } d = 0: \quad \hat{Y}_t = \hat{y}_t$$

$$\text{If } d = 1: \quad \hat{Y}_t = \hat{y}_t + Y_{t-1}$$

$$\text{If } d = 2: \quad \hat{Y}_t = \hat{y}_t + 2Y_{t-1} - Y_{t-2}$$

- Fortunately, your software will do all of this automatically!

## Do you need both AR and MA terms?

---

- In general, you don't: usually it suffices to use only one type or the other.
- Some series are better fitted by AR terms, others are better fitted by MA terms (at a given level of differencing).
- Rough rules of thumb:
  - If the stationarized series has positive autocorrelation at lag 1, AR terms often work best.
  - If it has negative autocorrelation at lag 1, MA terms often work best.
  - An MA(1) term often works well to fine-tune the effect of a nonseasonal difference, while an AR(1) term often works well to compensate for the lack of a nonseasonal difference, so the choice between them may depend on whether a difference has been used.



## Interpretation of AR terms

---


- A series displays autoregressive (AR) behavior if it apparently feels a “restoring force” that tends to pull it back toward its mean.
- In an AR(1) model, the AR(1) coefficient determines how fast the series tends to return to its mean. If the coefficient is near zero, the series returns to its mean quickly; if the coefficient is near 1, the series returns to its mean slowly.
- In a model with 2 or more AR coefficients, the sum of the coefficients determines the speed of mean reversion, and the series may also show an oscillatory pattern.

## AR and MA “signatures”


---

- ACF that dies out gradually and PACF that cuts off sharply after a few lags => **AR signature**
  - An AR series is usually positively autocorrelated at lag 1 (or even borderline nonstationary)
- ACF that cuts off sharply after a few lags and PACF that dies out more gradually => **MA signature**
  - An MA series is usually negatively autocorrelated at lag 1 (or even mildly overdifferenced)


Time Series Plot for MA



Estimated Autocorrelations for MA



Estimated Partial Autocorrelations for MA



AR signature: mean-reverting behavior, slow decay in ACF (usually positive at lag 1), sharp cutoff after a few lags in PACF.

Here the signature is AR(2) because of 2 spikes in PACF.

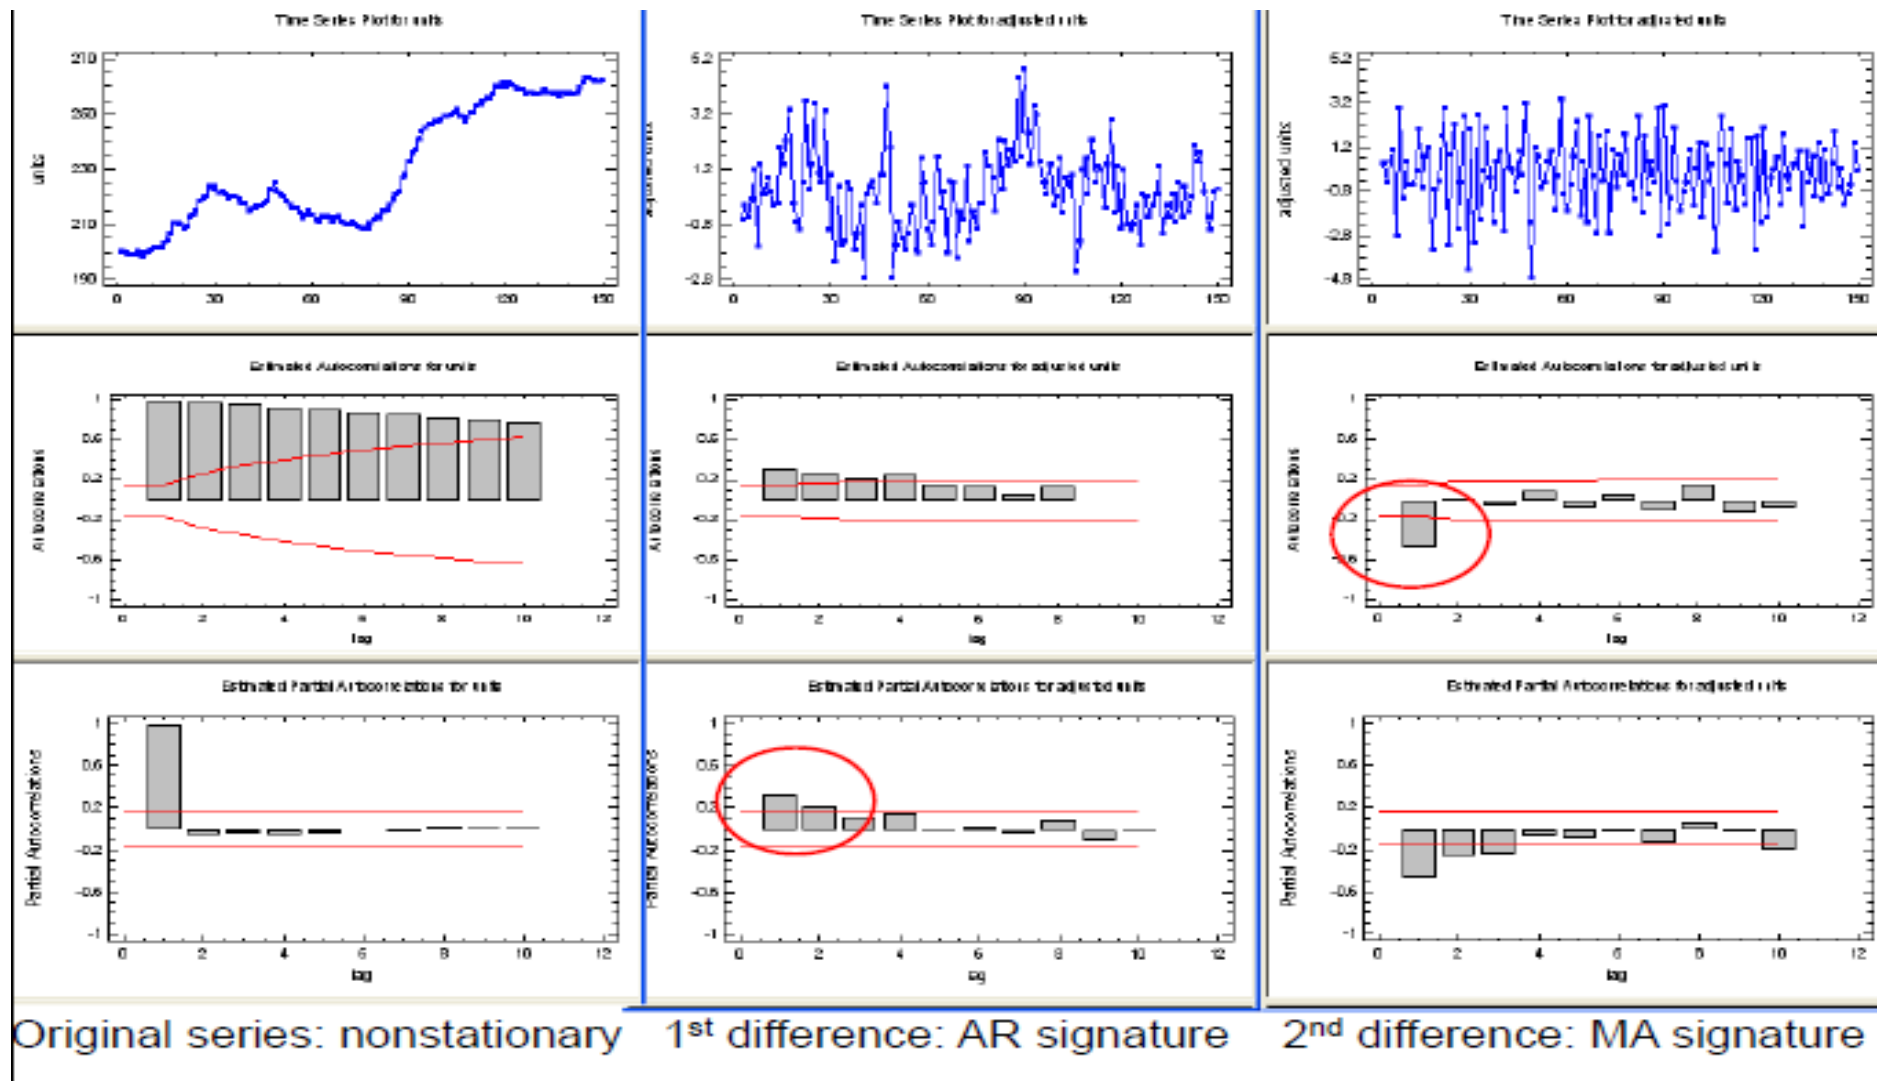
## AR or MA? It depends!

---

- Whether a series displays AR or MA behavior often depends on the extent to which it has been differenced.
- An “underdifferenced” series has an AR signature (positive autocorrelation)
- After one or more orders of differencing, the autocorrelation will become more negative and an MA signature will emerge
- Don’t go too far: if series already has zero or negative autocorrelation at lag 1, don’t difference again

# DATA ANALYTICS

## An example



1. Determine the order of differencing
2. Determine the numbers of AR & MA terms
3. Fit the model—check to see if residuals are “white noise,” highest-order coefficients are significant (w/ no “unit “roots”), and forecasts look reasonable.  
If not, return to step 1 or 2.

## Tuning Parameters - Technical issues

---

- Backforecasting
  - Estimation algorithm begins by forecasting backward into the past to get start-up values
- Unit roots
  - Look at sum of AR coefficients and sum of MA coefficients—if they are too close to 1 you may want to consider higher or lower order of differencing
- Overdifferencing
  - A series that has been differenced one too many times will show very strong negative autocorrelation and a strong MA signature, probably with a unit root in MA coefficients

## Seasonal ARIMA terminology

---

- The seasonal part of an ARIMA model is summarized by three *additional* numbers:

***P = # of seasonal autoregressive terms***

***D = # of seasonal differences***


***Q = # of seasonal moving-average terms***

- The complete model is called an “SARIMA( $p, d, q$ )( $P, D, Q$ )” model



# DATA ANALYTICS

The “filtering box” now has 6 knobs:



- How non-seasonal & seasonal differences are combined to stationarize the series:

If  $d=0, D=1$ :  $y_t = Y_t - Y_{t-s}$   $s$  is the seasonal period,  
e.g.,  $s=12$  for monthly data

$$\begin{aligned}\text{If } d=1, D=1: \quad y_t &= (Y_t - Y_{t-1}) - (Y_{t-s} - Y_{t-s-1}) \\ &= Y_t - Y_{t-1} - Y_{t-s} + Y_{t-s-1}\end{aligned}$$

$D$  should never be more than 1, and  $d+D$  should never be more than 2. Also, if  $d+D=2$ , the constant term should be suppressed.

- How SAR and SMA terms add coefficients to the model:
- Setting  $P=1$  (i.e., SAR=1) adds a multiple of
  - $y_{t-s}$  to the forecast for  $y_t$
- Setting  $Q=1$  (i.e., SMA=1) adds a multiple of
  - $e_{t-s}$  to the forecast for  $y_t$
- Total number of SAR and SMA factors usually should not be more than 1 (i.e., either SAR=1 or SMA=1, not both)

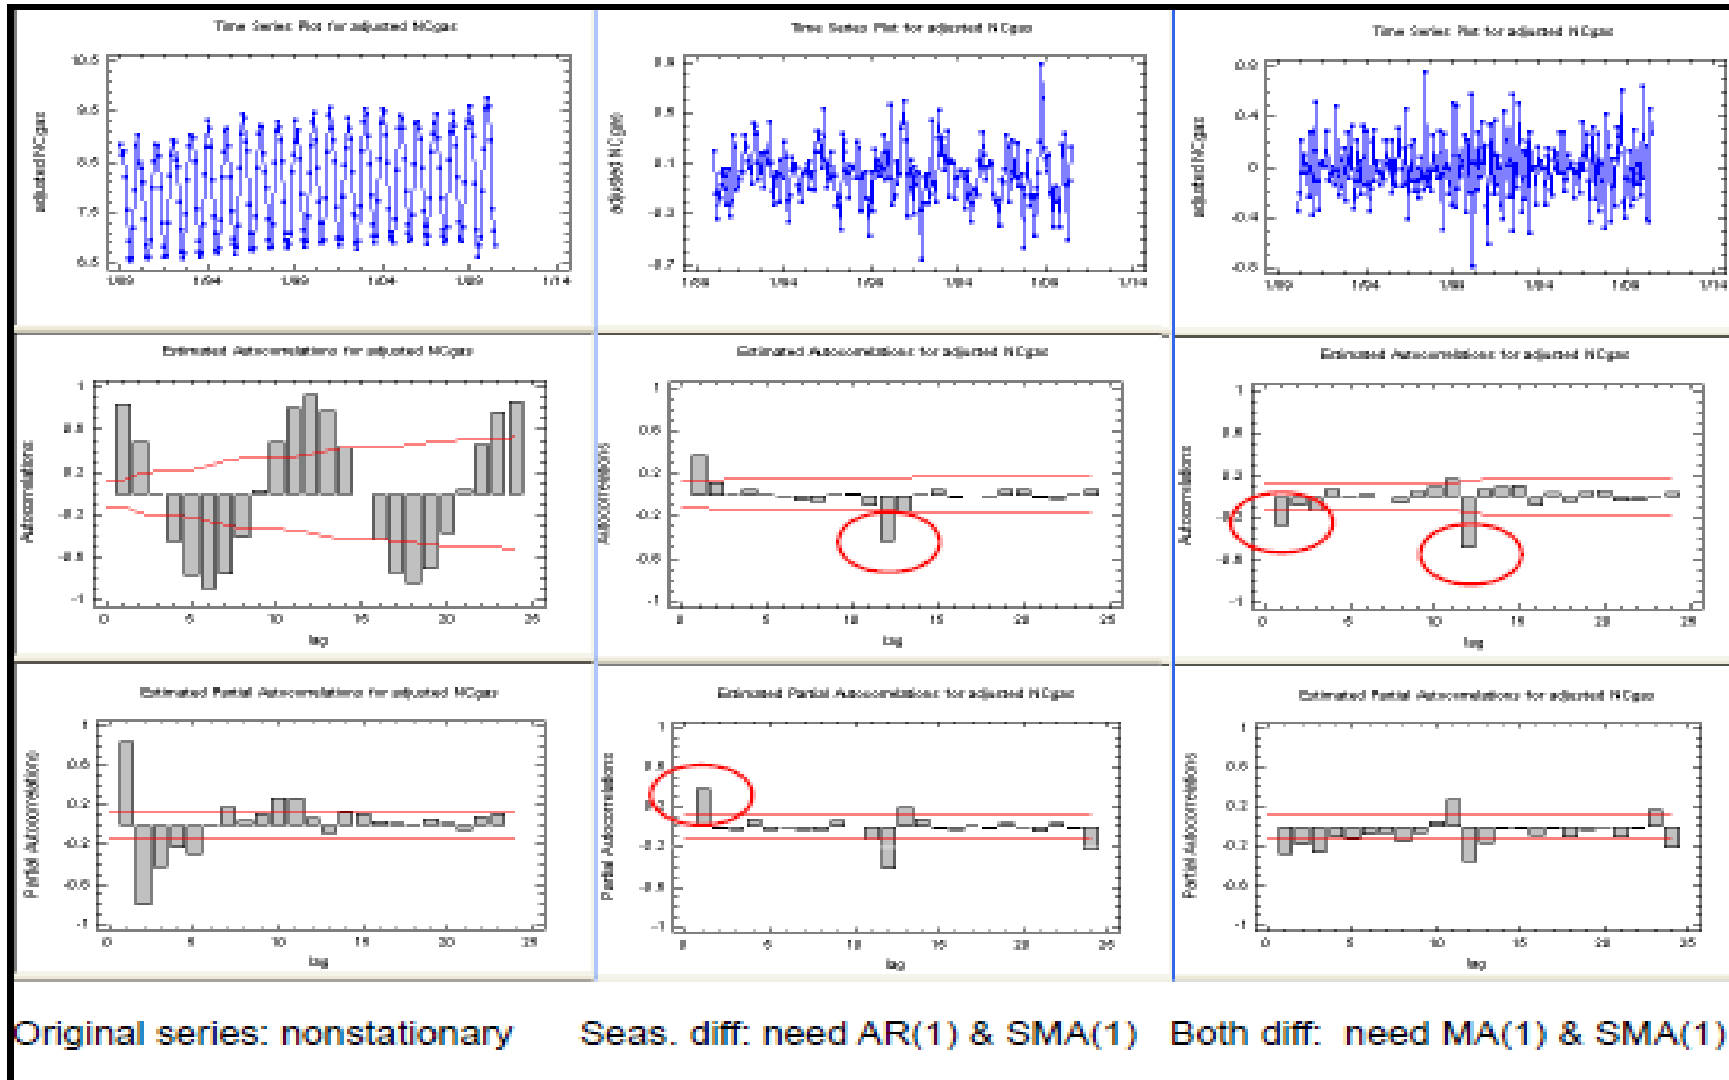
- Start by trying various combinations of one seasonal difference and/or one non-seasonal difference to stationarize the series and remove gross features of seasonal pattern.
- If the seasonal pattern is strong and stable, you **MUST** use a seasonal difference (otherwise it will “die out” in long-term forecasts)

## Model-fitting steps, continued

---

- After differencing, inspect the ACF and PACF at multiples of the seasonal period ( $s$ ):
  - Positive spikes in ACF at lag  $s$ ,  $2s$ ,  $3s$ ..., single positive spike in PACF at lag  $s$
  - $s \Rightarrow \text{SAR}=1$
  - Negative spike in ACF at lag  $s$ , negative spikes in PACF at lags  $s$ ,  $2s$ ,  $3s$ ,...  $\Rightarrow \text{SMA}=1$
  - $\text{SMA}=1$  often works well in conjunction with a seasonal difference.
  - Same principles as for non-seasonal models, except focused on what happens at multiples of lag  $s$  in ACF and PACF.

## Model-fitting steps, continued



## A common seasonal ARIMA model

---

- Often you find that the “correct” order of differencing is  $d=1$  and  $D=1$ .
- With one difference of each type, the autocorr. often negative at both lag 1 and lag  $s$ .
- This suggests an  $SARIMA(0,1,1)(0,1,1)$  model, a common seasonal ARIMA model.
- Similar to Winters’ model in estimating time-varying trend and time-varying seasonal pattern

## Another common seasonal ARIMA model

---

- Often with  $D=1$  (only) you see a borderline nonstationary pattern with  $AR(p)$  signature, where  $p=1$  or 2, sometimes 3
- After adding  $AR=1, 2$ , or 3, you may find negative autocorrelation at
- lag  $s$  ( $\Rightarrow$   $SMA=1$ )
- This suggests  $ARIMA(p,0,0) \times (0,1,1) + c$ , another common seasonal ARIMA model.
- Key difference from previous model: assumes a constant annual trend



## Bottom-line suggestion

---

- When fitting a time series with a strong seasonal pattern, you generally should try

**ARIMA(0,1,q)(0,1,1) model (q=1 or 2)**

**ARIMA(p,0,0)(0,1,1)+c model (p=1, 2 or 3)**

- ... in addition to other models (e.g., Random Walk, Single Exponential Smoothing, etc., with seasonal adjustment; or Winters Method)
- If there is a significant trend and/or the seasonal pattern is multiplicative, you should also try a natural log transformation.

# DATA ANALYTICS

## Seasonal differencing

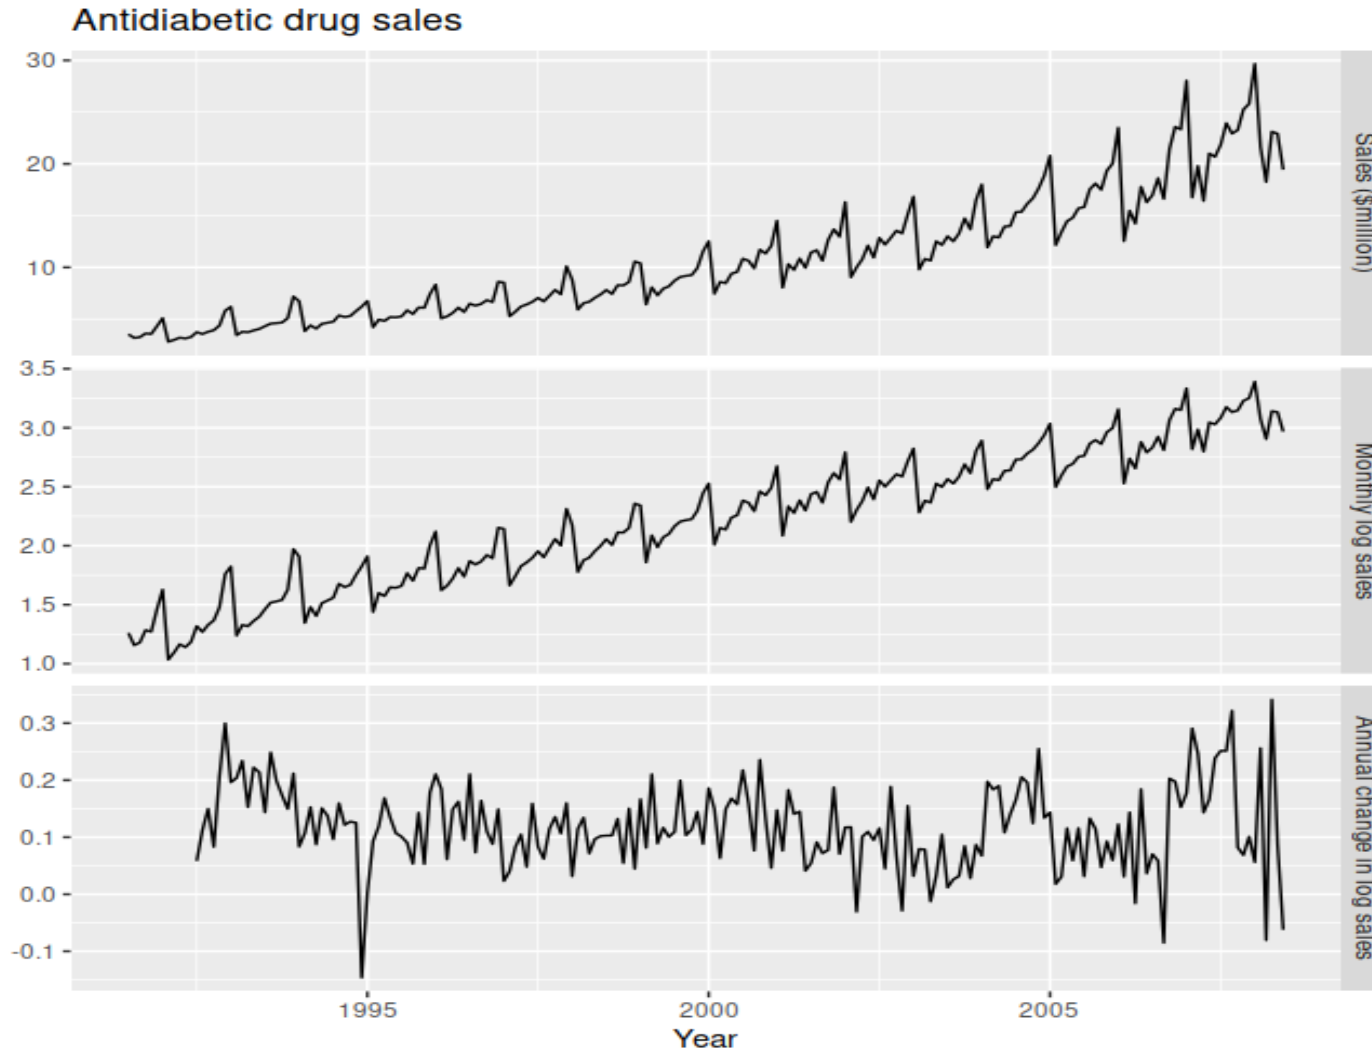


Figure 11:  
Logs and  
seasonal  
differences of  
the A10  
(antidiabetic)  
sales data.

- The bottom panel in Figure 11, shows the seasonal differences of the logarithm of the monthly scripts for A10 (antidiabetic) drugs sold in Australia.
- The transformation and differencing have made the series look relatively stationary.
- The logarithms stabilise the variance, while the seasonal differences remove the seasonality and trend.
- To distinguish seasonal differences from ordinary differences, we refer to ordinary differences as “first differences”, meaning differences at lag 1.

# DATA ANALYTICS

## Seasonal differencing




Figure 12: Monthly Net Electricity Generation in the US

- Sometimes it is necessary to take both a seasonal difference and a first difference to obtain stationary data, as is shown in Figure 12.
- Here, the data are first transformed using logarithms (second panel), then seasonal differences are calculated (third panel).
- The data still seem somewhat non-stationary, and so a further lot of first differences are computed (bottom panel).
- Figure 12: Top panel: US net electricity generation (billion kWh).
- Other panels show the same data after transforming and differencing.




Figure 12: Monthly Net Electricity Generation in the US

## Take-aways

---

- Seasonal ARIMA models (especially the  $(0,1,q) \times (0,1,1)$  and  $(p,0,0) \times (0,1,1) + c$  models) compare favorably with other seasonal models and often yield better short-term forecasts.
- Advantages: solid underlying theory, stable estimation of time-varying trends and seasonal patterns, relatively few parameters.
- Drawbacks: no explicit seasonal indices, hard to interpret coefficients or explain “how the model works”, danger of overfitting or misidentification if not used with care.

### **Text Book:**

“Business Analytics, The Science of Data-Driven Making”, U. Dinesh Kumar, Wiley 2017 ([Ch. 13.14 -13.14.4](#))



# DATA ANALYTICS

## Image Courtesy

---



<https://www.abs.gov.au/websitedbs/D3310114.nsf/home/Time+Series+Analysis:+The+Basics>

<https://otexts.com/fpp2/stationarity.html>



# THANK YOU

---

**Jyothi R**

Assistant Professor, Department of  
Computer Science

[jyothir@pes.edu](mailto:jyothir@pes.edu)