



OPERATING SYSTEMS

Memory Management -2

Nitin V Pujari
Faculty, Computer Science
Dean - IQAC, PES University

OPERATING SYSTEMS

Course Syllabus - Unit 3



Unit-3: Unit 3: Memory Management: Main Memory

Hardware and control structures, OS support, Address translation, Swapping, Memory Allocation (Partitioning, relocation), Fragmentation, Segmentation, Paging, TLBs context switches

Virtual Memory - Demand Paging, Copy-on-Write, Page replacement policy - LRU (in comparison with FIFO & Optimal), Thrashing, design alternatives - inverted page tables, bigger pages.

Case Study: Linux/Windows Memory

OPERATING SYSTEMS


Course Outline



25	Main Memory: Hardware and control structures, OS support, Address translation	8.1	64.2
26	Dynamic linking, Swapping	8.2	
27	Memory Allocation (Partitioning, relocation), Fragmentation	8.3	
28	Segmentation	8.4	
29	Paging: OS Support, TLBs, Address Translation	8.5	
30	Structure of the Page Table	8.6	
31	Design Alternatives - Inverted Page Tables, Bigger Pages	8.7-8.8	
32	Virtual Memory: Demand Paging, Copy-OnWrite	9.1-9.3	
33	Page replacement policy - LRU	9.4	
34	FIFO & Optimal	9.5	
35	Thrashing	9.6	
36	Case Study: Linux/ Windows Memory Management	9.10	

- **Dynamic Allocation using Relocation Register**
- **Dynamic Linking**
- **Process Swapping**
- **Schematic view of Swapping**
- **Context Switching time including swapping**
- **Swapping on Mobile Systems**

- Routine is not loaded until it is called
- Better memory-space utilization; unused routine is never loaded
- All routines kept on disk in relocatable load format
- Useful when large amounts of code are needed to handle infrequently occurring cases
- No special support from the operating system is required
- Implemented through program design
- OS can help by providing libraries to implement dynamic loading



- When a process is running, what does its memory look like ? A collection of regions called sections. Basic memory layout for Linux and other Unix systems:
 - Code (or "text" in Unix terminology): starts at location 0
 - Data: starts immediately above code, grows upward
 - Stack: starts at highest address, grows downward
- System components that take part in managing a process's memory:
- **Compiler and assembler:**
 - Generate one object file for each source code file containing information for that source file.
 - Information is incomplete, since each source file generally references some things defined in other source files.

- **Linker:**

- Combines all of the object files for one program into a single object file.
- Linker output is complete and self-sufficient.

- **Operating system:**

- Loads object files into memory.
- Allows several different processes to share memory at once.
- Provides facilities for processes to get more memory after they've started running.

- **Run-time library:**

- Works together with OS to provide dynamic allocation routines, such as malloc and free in C.

- **Static linking** => system libraries and program code combined by the loader into the binary program image
- **Dynamic linking** => linking postponed until execution time
- Small piece of code, stub, used to locate the appropriate memory-resident library routine
- Stub replaces itself with the address of the routine, and executes the routine
- Operating system checks if routine is in process's memory address

- Since late 1980's most systems have supported shared libraries and dynamic linking
- For common library packages, only keep a single copy in memory, shared by all processes.
- Don't know where library is loaded until runtime; must resolve references dynamically, when program runs.


- A process can be swapped temporarily out of memory to a backing store, and then brought back into memory for continued execution
 - Total physical memory space of processes can exceed physical memory
- Backing store is fast disk large enough to accommodate copies of all memory images for all users; providing direct access to these memory images

- Roll out, roll in => swapping variant used for priority based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed
- Major part of swap time is transfer time; total transfer time is directly proportional to the amount of memory swapped
- System maintains a ready queue of ready-to-run processes which have memory images on disk

- Does the swapped out process need to swap back in to same physical addresses ?
 - Depends on address binding method
 - Plus consider pending I/O to / from process memory space
- Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows)
 - Swapping normally disabled
 - Started if more than threshold amount of memory allocated
 - Disabled again once memory demand reduced below threshold

OPERATING SYSTEMS

Schematic View of Process Swapping



- If next processes to be put on CPU is not in memory, need to swap out a process and swap in target process
- Context switch time can then be very high
- 100MB process swapping to hard disk with transfer rate of 50MB / sec \Rightarrow 50MB / 1000 ms \Rightarrow .05 ms per Mb
 - Swap out time of 2000 ms
 - Plus swap in case of same sized process
 - Total context switch swapping component time of 4000 ms (4 seconds)

- Can reduce => if reduce size of memory swapped => by knowing how much memory really being used ?
- System calls to inform OS of memory use via `request_memory()` and `release_memory()`

Context Switch Time including Process Swapping



- Other constraints as well on swapping
 - Pending I/O – can't swap out as I/O would occur to wrong process
 - or always transfer I/O to kernel space, then to I/O device
 - Known as double buffering, adds overhead
- Standard swapping not used in modern operating systems
 - But modified version common.
 - Swap only when free memory extremely low

Process Swapping on Mobile System

- Not typically supported
 - Flash memory based
 - Small amount of space
 - Limited number of write cycles
 - Poor throughput between flash memory and CPU on mobile platform

Process Swapping on Mobile System

- Instead use other methods to free memory if low
 - iOS asks apps to voluntarily relinquish allocated memory
 - Read-only data thrown out and reloaded from flash if needed
 - Failure to free can result in termination
 - Android terminates apps if low free memory, but first writes application state to flash for fast restart



THANK YOU

Nitin V Pujari
Faculty, Computer Science
Dean - IQAC, PES University

nitin.pujari@pes.edu

For Course Deliverables by the Anchor Faculty click on www.pesuacademy.com and complete reading assignments provided by the Anchor on Edmodo