



OPERATING SYSTEMS

Memory Management - 5

Nitin V Pujari
Faculty, Computer Science
Dean - IQAC, PES University

OPERATING SYSTEMS

Course Syllabus - Unit 3



Unit-3: Unit 3: Memory Management: Main Memory

Hardware and control structures, OS support, Address translation, Swapping, Memory Allocation (Partitioning, relocation), Fragmentation, Segmentation, Paging, TLBs context switches

Virtual Memory - Demand Paging, Copy-on-Write, Page replacement policy - LRU (in comparison with FIFO & Optimal), Thrashing, design alternatives - inverted page tables, bigger pages.

Case Study: Linux/Windows Memory

OPERATING SYSTEMS

Course Outline




25	Main Memory: Hardware and control structures, OS support, Address translation	8.1	64.2
26	Dynamic linking, Swapping	8.2	
27	Memory Allocation (Partitioning, relocation), Fragmentation	8.3	
28	Segmentation	8.4	
29	Paging: OS Support, TLBs, Address Translation	8.5	
30	Structure of the Page Table	8.6	
31	Design Alternatives - Inverted Page Tables, Bigger Pages	8.7-8.8	
32	Virtual Memory: Demand Paging, Copy-OnWrite	9.1-9.3	
33	Page replacement policy - LRU	9.4	
34	FIFO & Optimal	9.5	
35	Thrashing	9.6	
36	Case Study: Linux/ Windows Memory Management	9.10	

- **Non Contiguous memory allocation**
- **Paging**
- **Address Translation Scheme**
- **Paging Hardware**
- **Paging model of Logical and Physical Memory**
- **Paging Example**
- **Free Frames**

- **Implementation of Page Table**
- **Associative Memory**
- **Table Lookaside Buffer**
- **Effective Access Time**
- **Memory Protection**
- **Valid - Invalid Bit**
- **Shared Pages**
- **Shared Pages Example**

Non - Contiguous Memory Allocation

- Non-contiguous memory allocation is a memory allocation technique.
- It allows to store parts of a single process in a non-contiguous fashion.
- Thus, different parts of the same process can be stored at different places in the main memory.



Non - Contiguous Memory Allocation

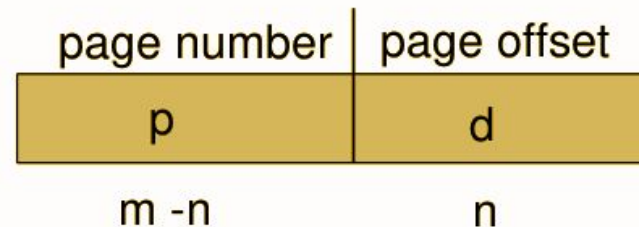
- Physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
 - Avoids external fragmentation
 - Avoids problem of varying sized memory chunks

Paging

- Divide physical memory into fixed-sized blocks called **Frames**
 - Size is power of 2, between 512 bytes and 16 Mbytes
- Divide logical memory into blocks of same size called **Pages**
- Keep track of all free frames
- To run a program of size N pages, need to find N free frames and load program
- Set up a page table to translate logical to physical addresses
- Backing store likewise split into pages
- Still have Internal fragmentation

Address Translation Scheme


- Address generated by CPU is divided into:
 - **Page number (p)** => used as an index into a page table which contains base address of each page in physical memory, which is nothing but the frame address
 - **Page offset (d)** => combined with base address to define the physical memory address that is sent to the memory unit



- For given logical address space 2^m and page size 2^n

OPERATING SYSTEMS

Paging Hardware




logical
memory

0	1
1	4
2	3
3	7

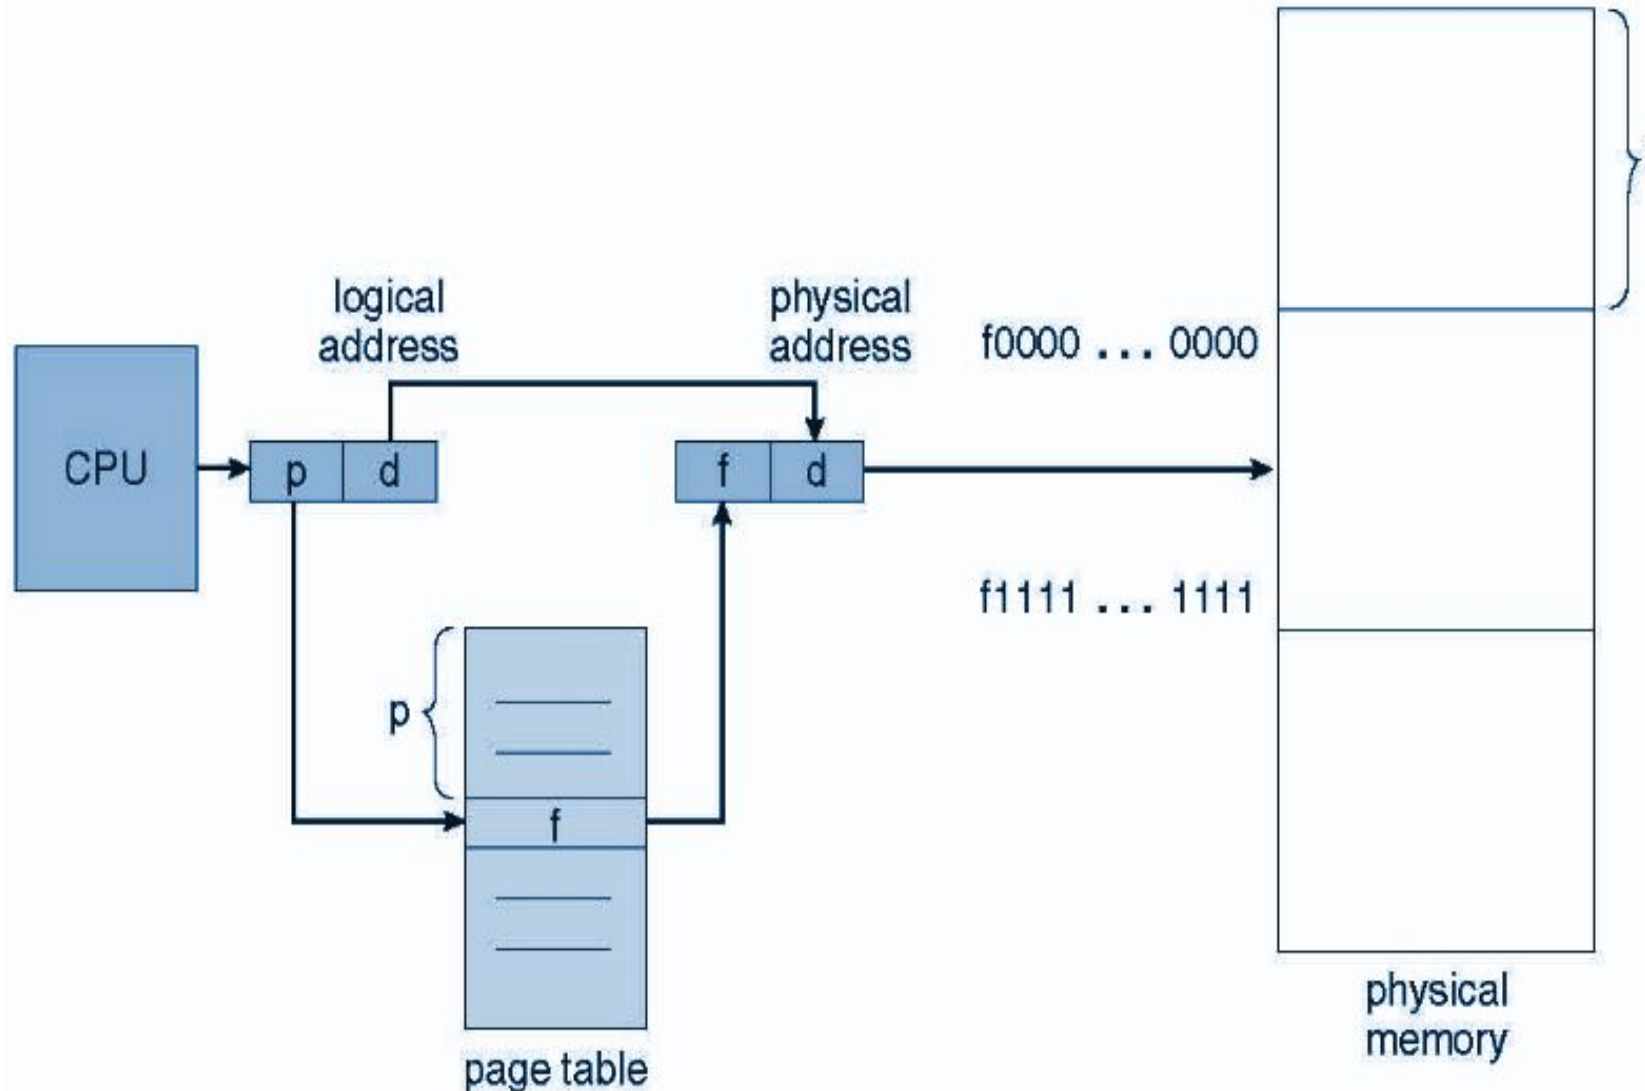
page table

frame
number



physical
memory

Paging Model of Logical and Physical Memory



OPERATING SYSTEMS

Paging Example

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5
1	6
2	1
3	2

page table

page number	page offset
p	d
m - n	n

- For given logical address space 2^m and page size 2^n

for $n=2$ and $m=4$

Logical address range $\Rightarrow 2^4 = (0..15)$

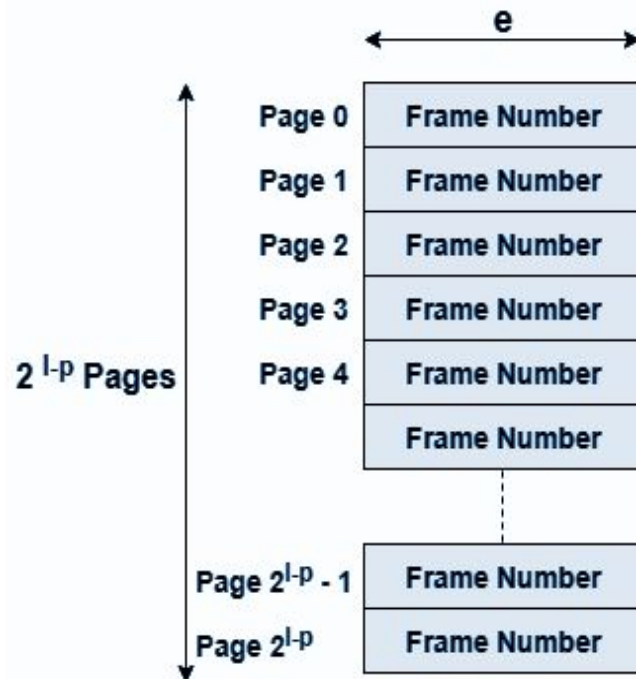
Page Size $\Rightarrow 2^2 = 4$ Pages

0	
4	i
	j
	k
8	m
	n
	o
	p
12	
16	
20	a
	b
	c
	d
24	e
	f
	g
	h
28	



physical memory

OPERATING SYSTEMS

Paging Example



Page Table



No. of entries in Page Table = No. of the pages in the process

Page Table Size = $2^{l-p} \times e$ bytes

$e = m - p$ (Frame Size) bits

Physical Address Space = M words
Logical Address Space = L words
Page Size = P words

Physical Address = $\log_2 M = m$ bits
Logical Address = $\log_2 L = l$ bits
page offset = $\log_2 P = p$ bits

<https://www.kernel.org/doc/gorman/html/understand/understand006.html>

Paging Example

Page Number	Logical Address	Content
0 => 00	0000	a
	0001	b
	0010	c
	0011	d
1 => 01	0100	e
	0101	f
	0110	g
	0111	h
2 => 10	1000	i
	1001	j
	1010	k
	1011	l
3 => 11	1100	m
	1101	n
	1110	o
	1111	p

Page Table	
Page Number	Frame Number
0	3
1	1
2	2
3	0

For Logical address		
Logical Address	Page #	
1001	2	
1100	3	
0001	0	
0101	1	

Frame#	Physical Address	Content
0	0000	m
	0001	n
	0010	o
	0011	p
1	0100	e
	0101	f
	0110	g
	0111	h
2	1000	i
	1001	j
	1010	k
	1011	l
3	1100	a
	1101	b
	1110	c
	1111	d

OPERATING SYSTEMS

Paging Example

Logical Address generated by CPU	Physical Address generated in runtime
1010 => [1010] => (10,10) => (Page 2, Offset 2) => k	(Frame 2, Offset 2) => (1000,10) => 1010 => [1010] => k
0011 => [0011] => (00,11) => (Page 0, Offset 3) => d	(Frame 3, Offset 3) => (1100,11) => 1111 => [1111] => d
1101 => [1101] => (11,01) => (Page3, offset 1) => n	(Frame 0, Offset 1) => (0000, offset 1) => 0001 => [0001] => n
1111 => ? => ? => ? => ?	(Frame ?, Offset ?) => ? => ? => ? => ?

For Logical address	
Logical Address	Page #
1001	2
1100	3
0001	0
0101	1

Page Table	
Page Number	Frame Number
0	3
1	1
2	2
3	0

Frame#	Physical Address	Content
0	0000	m
	0001	n
	0010	o
	0011	p
1	0100	e
	0101	f
	0110	g
	0111	h
2	1000	i
	1001	j
	1010	k
	1011	l
3	1100	a
	1101	b
	1110	c
	1111	d



Paging : Calculating Internal Fragmentation

- Page size = 2,048 bytes
- Process size = 72,766 bytes
- 35 pages + 1,086 bytes
- Internal fragmentation of $2,048 - 1,086 = 962$ bytes
- Worst case fragmentation = 1 frame – 1 byte
- On average fragmentation = $1 / 2$ frame size
- So small frame sizes desirable ?
- But each page table entry takes memory to track
- Page sizes growing over time

- Process view and physical memory now very different
- By implementation process can only access its own memory

OPERATING SYSTEMS

Free Frame List



Implementation of Page table

- Page table is kept in main memory
- Page-table base register (PTBR) points to the page table
- Page-table length register (PTLR) indicates size of the page table
- In this scheme every data/instruction access requires two memory accesses
- One for the page table and one for the data / instruction
- The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **Associative Memory or Translation Lookaside Buffers (TLBs)**

Implementation of Page table

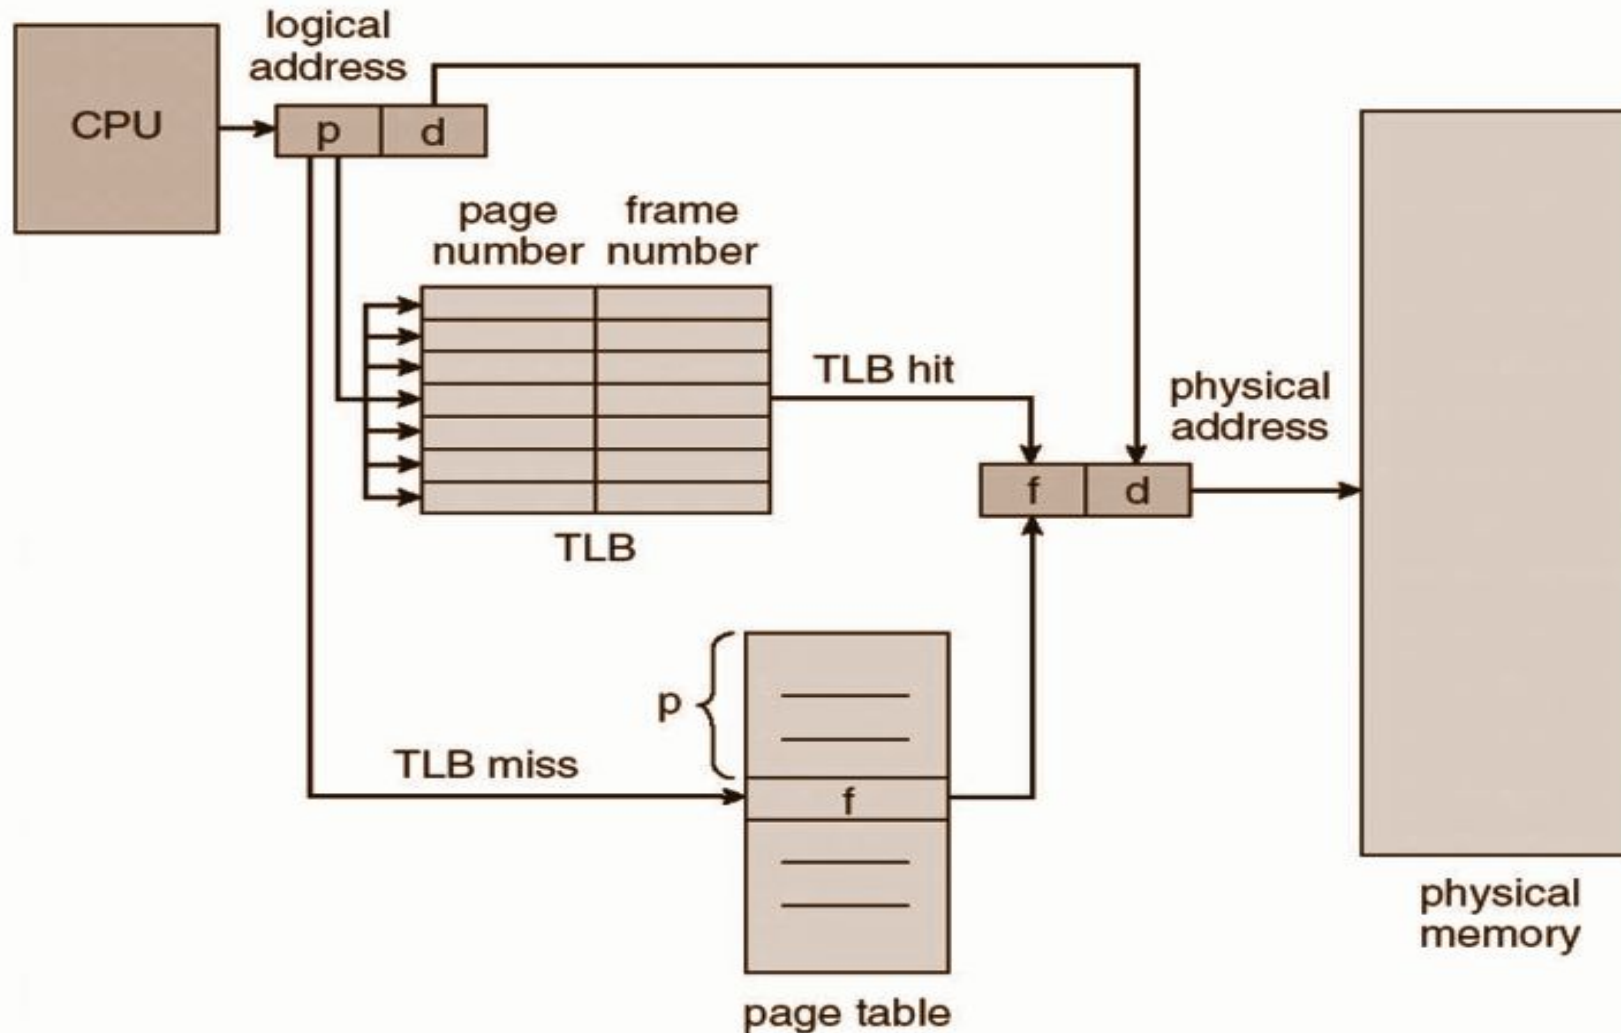
- Some TLBs store **Address-Space Identifiers (ASIDs)** in each TLB entry that uniquely identifies each process to provide address-space protection for that process
- Otherwise need to flush at every context switch
- TLBs typically small (64 to 1,024 entries)
- On a TLB miss, value is loaded into the TLB for faster access next time
- Replacement policies must be considered
- Some entries can be wired down for permanent fast access

- Associative memory – parallel search

Page #	Frame #

- Address translation (p, d)
 - If p is in associative register, get frame # out
 - Otherwise get frame # from page table in memory

Paging Hardware with TLB



- Associative Lookup = time unit
 - Can be $< 10\%$ of memory access time
- Hit ratio = α
- TLB search and access time $\Rightarrow \epsilon$ (Epsilon)
- Hit ratio – percentage of times that a page number is found in the associative registers; ratio related to number of associative registers

Effective Access Time

- Effective Access Time (EAT)
 - $EAT = (1 + \epsilon) \alpha + (2 + \epsilon)(1 - \alpha)$
 $= 2 + \epsilon - \alpha$
 - **EAT = Found + NotFound (Considering Epsilon)**
 - **Found**=> $\alpha \times 120 \text{ ns} \Rightarrow 0.8 \times 120 \Rightarrow 96 \text{ ns}$
 - **Not Found**=> $(1 - \alpha) \times 200 \text{ ns} \Rightarrow 0.2 \times 200 \text{ ns} \Rightarrow 40 \text{ ns}$
- Consider $\alpha = 80\%$, $\epsilon = 20 \text{ ns}$ for TLB search, 100 ns for memory access
 - **EAT = $0.80 \times 120 + 0.20 \times 200 = 136 \text{ ns}$ (Considering Epsilon)**
- Consider more realistic hit ratio => 99% , $\epsilon = 20 \text{ ns}$ for TLB search, 100 ns for memory access
 - **EAT = $0.99 \times 120 + 0.01 \times 200 = 119 \text{ ns}$ (Considering Epsilon)**



Effective Access Time

- Effective Access Time (EAT)
 - $EAT = (1 + \epsilon) \alpha + (2 + \epsilon)(1 - \alpha)$
 $= 2 + \epsilon - \alpha$
 - $EAT = \text{Found} + \text{NotFound}$ (ignoring Epsilon)
 - **Found** $\Rightarrow \alpha \times 100\text{ns} \Rightarrow 0.8 \times 100 \Rightarrow 80\text{ns}$
 - **Not Found** $\Rightarrow (1 - \alpha) \times 200\text{ns} \Rightarrow 0.2 \times 200 \text{ ns} \Rightarrow 40\text{ns}$
- Consider $\alpha = 80\%$, $\epsilon = 20\text{ns}$ for TLB search, 100ns for memory access
 - $EAT = 0.80 \times 100 + 0.20 \times 200 = 120\text{ns}$ (ignoring Epsilon)
- Consider more realistic hit ratio $\Rightarrow 99\%$, $\epsilon = 20\text{ns}$ for TLB search, 100ns for memory access
 - $EAT = 0.99 \times 100 + 0.01 \times 200 = 101\text{ns}$ (ignoring Epsilon)

Memory Protection

- Memory protection implemented by associating protection bit with each frame to indicate if read-only or read-write access is allowed
 - Can also add more bits to indicate page execute-only, and so on
- Valid-invalid bit attached to each entry in the page table:
 - “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page
 - “invalid” indicates that the page is not in the process’ logical address space
 - Or use page-table length register (PTLR)
- Any violations result in a trap to the kernel

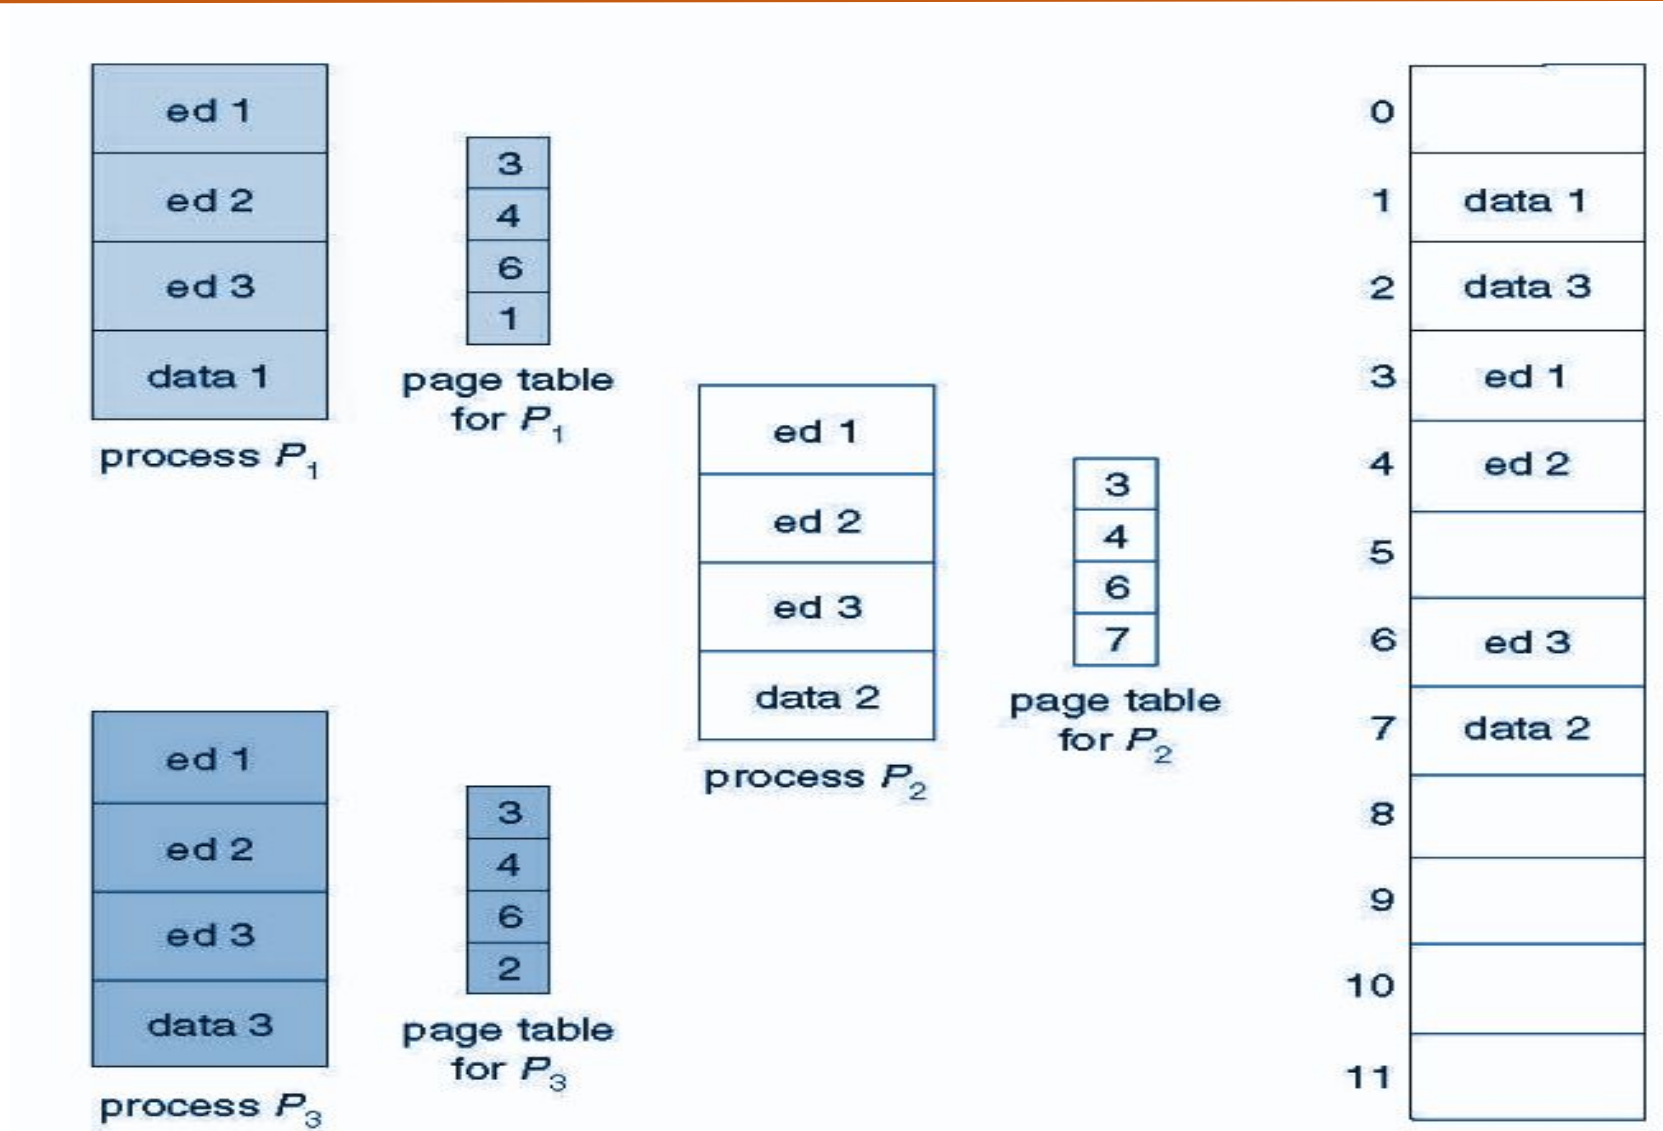
Valid Invalid Bit in the Page Table



Shared Pages

- **Shared code**
 - One copy of read-only (reentrant) code shared among processes (i.e., text editors, compilers, window systems)
 - Similar to multiple threads sharing the same process space
 - Also useful for interprocess communication if sharing of read-write pages is allowed
- **Private Code and Data**
 - Each process keeps a separate copy of the code and data
 - The pages for the private code and data can appear anywhere in the logical address space

Shared Pages - Example





THANK YOU

Nitin V Pujari
Faculty, Computer Science
Dean - IQAC, PES University

nitin.pujari@pes.edu

For Course Deliverables by the Anchor Faculty click on www.pesuacademy.com and complete reading assignments provided by the Anchor on Edmodo