



OPERATING SYSTEMS

Memory Management - 11

Nitin V Pujari
Faculty, Computer Science
Dean - IQAC, PES University

OPERATING SYSTEMS

Course Syllabus - Unit 3



Unit-3: Unit 3: Memory Management: Main Memory

Hardware and control structures, OS support, Address translation, Swapping, Memory Allocation (Partitioning, relocation), Fragmentation, Segmentation, Paging, TLBs context switches
Virtual Memory - Demand Paging, Copy-on-Write, Page replacement policy - LRU (in comparison with FIFO & Optimal), Thrashing, design alternatives - inverted page tables, bigger pages.
Case Study: Linux/Windows Memory

OPERATING SYSTEMS

Course Outline



25	Main Memory: Hardware and control structures, OS support, Address translation	8.1	64.2
26	Dynamic linking, Swapping	8.2	
27	Memory Allocation (Partitioning, relocation), Fragmentation	8.3	
28	Segmentation	8.4	
29	Paging: OS Support, TLBs, Address Translation	8.5	
30	Structure of the Page Table	8.6	
31	Design Alternatives - Inverted Page Tables, Bigger Pages	8.7-8.8	
32	Virtual Memory: Demand Paging, Copy-OnWrite	9.1-9.3	
33	Page replacement policy - LRU	9.4	
34	FIFO & Optimal	9.5	
35	Thrashing	9.6	
36	Case Study: Linux/ Windows Memory Management	9.10	

- **Virtual Memory - Page replacement**
- **What happens if there is no free Frame ?**
- **Basic Page Replacement**
- **Page and Frame Replacement Algorithms**
- **Graph of Page Faults versus the number of Frames**

- **First-In-First-Out (FIFO) Algorithm**
- **FIFO illustrating Belady's Anomaly**
- **Optimal Page Replacement Algorithm**
- **Least Recently Used (LRU) Algorithm**
- **Use of a Stack to Record Most Recent Page References**

- LRU Algorithm Implementation
- LRU Approximation Algorithm
- Second-Chance (clock) Page-Replacement Algorithm
- Enhanced Second-Chance Algorithm
- Counting Algorithms
- Page-Buffering Algorithms

- Applications and Page Replacement
- Allocation of Frames
- Fixed Allocation
- Priority Allocation
- Global vs. Local Allocation
- Non-Uniform Memory Access

LRU Approximation Algorithms

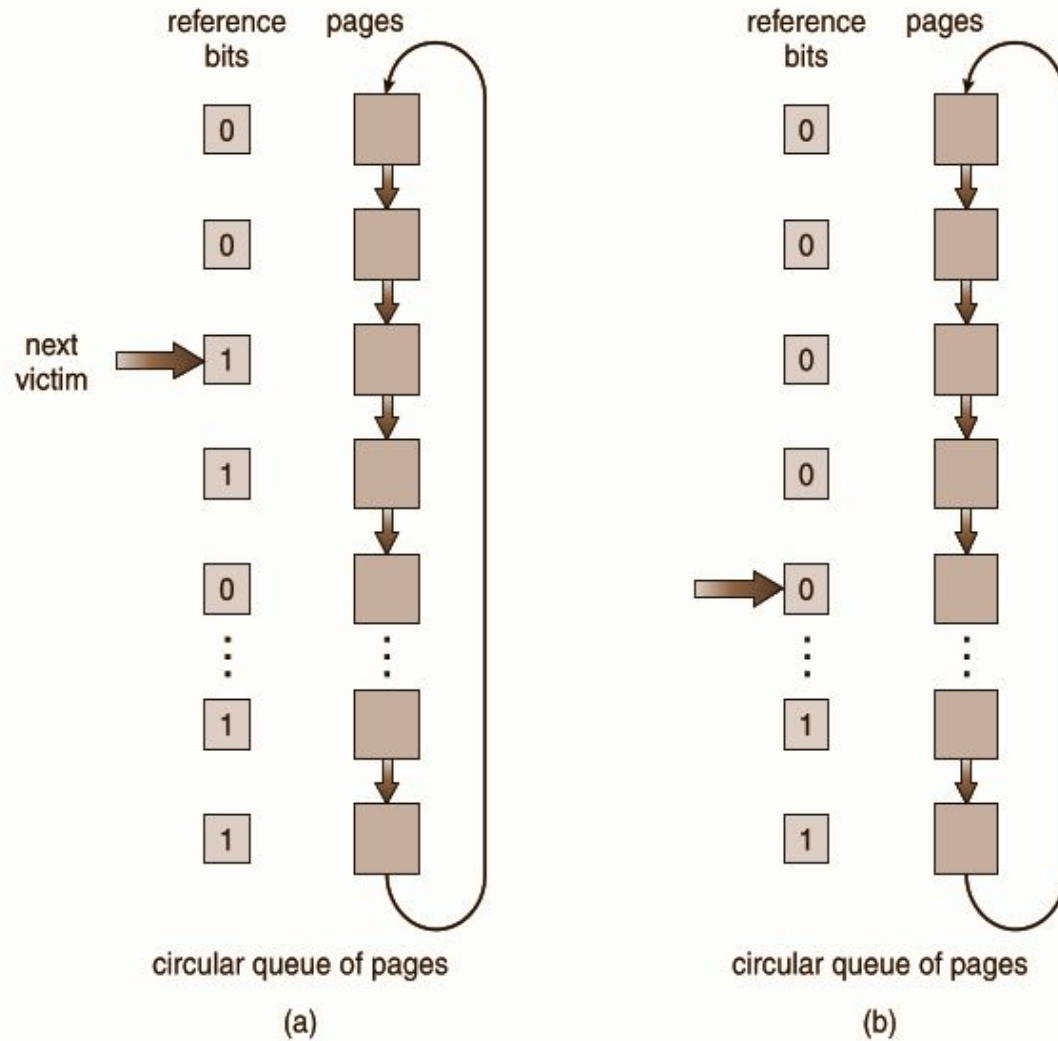


- LRU needs special hardware and still slow
- Reference bit
 - With each page associate a bit, initially = 0
 - When page is referenced bit set to 1
 - Replace any with reference bit = 0 (if one exists)
 - We do not know the order, however



Second-Chance (clock) Page-Replacement Algorithm

- Second-chance algorithm
 - Generally FIFO, plus hardware-provided reference bit
 - Clock replacement
 - If page to be replaced has reference bit = 0 -> replace it
 - if reference bit = 1 then:
 - set reference bit 0, leave page in memory
 - replace next page, subject to same rules


Second-Chance (clock) Page-Replacement Algorithm



Second-Chance (clock) Page-Replacement Algorithm



Second-Chance (clock) Page-Replacement Algorithm



- We can see notably that the bad replacement decisions made by FIFO is not present in Second chance !!!
- There are a total of 9 page read operations to satisfy the total of 18 page requests - just as good as the more computationally expensive LRU method !!!

Enhanced Second-Chance Algorithm

- Improve algorithm by using reference bit and modify bit (if available) in concert
- Take ordered pair (reference, modify)
 1. (0, 0) neither recently used nor modified – best page to replace
 2. (0, 1) not recently used but modified – not quite as good, must write out before replacement
 3. (1, 0) recently used but clean – probably will be used again soon
 4. (1, 1) recently used and modified – probably will be used again soon and need to write out before replacement
- When page replacement called for, use the clock scheme but use the four classes replace page in lowest non-empty class
 - Might need to search circular queue several times

- Keep a counter of the number of references that have been made to each page
 - Not common
- **Least Frequently Used (LFU) Algorithm:** replaces page with smallest count
- **Most Frequently Used (MFU) Algorithm:** based on the argument that the page with the smallest count was probably just brought in and has yet to be used

Req #	R1	R2	R3	R4	R5	R6	R7	R9	R10
Page #	7	0	1	2	0	3	0	3	7

Fr #	R1	FQ	Fr #	R2	FQ	Fr #	R3	FQ	Fr #	R4	FQ	Fr #	R5	FQ	Fr #	R6	FQ	Fr #	R7	FQ
1	7	1	1	7	1	1	7	1	1	2	1	1	2	1	1	2	1	1	2	1
2			2	0	1	2	0	1	2	0	1	2	0	2	2	0	2	2	0	3
3			3			3	1	1	3	1	1	3	1	1	3	3	1	3	3	2
Flag	PF	NA	Flag	PF	NA	Flag	PF	NA	Flag	PF	NA	Flag	PH	NA	Flag	PF	NA	Flag	PH	NA

Working Set => { }

Total Number of Page Requests = > 7

Total Number of Frames => 3

Total Number of Page faults =>

% of Page Faults = > => %

% of Page Hits => %

Legend

Page Fault => PF

Page Hit => PH

FQ => Frequency Count

Fr #	R10	FQ
1	7	1
2	0	3
3	3	2
Flag	PF	NA

Req #	R1	R2	R3	R4	R5	R6	R7	R8
Page #	7	0	1	0	7	7	2	7

Fr #	R1	FQ	Fr #	R2	FQ	Fr #	R3	FQ	Fr #	R4	FQ	Fr #	R5	FQ	Fr #	R6	FQ	Fr #	R7	FQ
1	7	1	1	7	1	1	7	1	1	7	1	1	7	2	1	7	3	1	2	1
2			2	0	1	2	0	1	2	0	2	2	0	2	2	0	2	2	0	2
3			3			3	1	1	3	1	1	3	1	1	3	1	1	3	1	1
Flag	PF	NA	Flag	PF	NA	Flag	PF	NA	Flag	PH	NA	Flag	PH	NA	Flag	PH	NA	Flag	PF	NA

Working Set => { }

Total Number of Page Requests = > 22

Total Number of Frames => 3

Total Number of Page faults =>

% of Page Faults = > => %

% of Page Hits => %

Legend

Page Fault => PF

Page Hit => PH

FQ => Frequency Count

Fr #	R7	FQ
1	2	1
2	7	1
3	1	1
Flag	PF	NA

- Keep a pool of free frames, always
 - Then frame available when needed, not found at fault time
 - Read page into free frame
 - Select victim to evict and add to freed pool
 - When convenient, evict Victim
- Possibly, keep list of modified pages
 - When backing store otherwise idle, write pages there and set to non-dirty
 - Possibly, keep free frame contents intact and note what is in them
 - If referenced again before reused, no need to load contents again from disk
 - Generally useful to reduce penalty if wrong victim frame selected

- All of these algorithms have OS guessing about future page access
- Some applications have better knowledge – i.e. databases
- Memory intensive applications can cause double buffering
 - OS keeps copy of page in memory as I/O buffer
 - Application keeps page in memory for its own work
- Operating system can given direct access to the disk, getting out of the way of the applications
 - Raw disk mode
- Bypasses buffering, locking, etc

- Each process needs minimum number of frames
- Example: IBM 370 – 6 pages to handle SS MOVE instruction:
 - instruction is 6 bytes, might span 2 pages
 - 2 pages to handle from
 - 2 pages to handle to
- Maximum of course is total frames in the system
- Two major allocation schemes
 - fixed allocation
 - priority allocation
- Many variations

- Equal allocation – For example, if there are 100 frames (after allocating frames for the OS) and 5 processes, give each process 20 frames
 - Keep some as free frame buffer pool
- Proportional allocation – Allocate according to the size of process
- Dynamic as degree of multiprogramming, process sizes $m = 64$ change

s_i = size of process p_i

$$S = \sum s_i$$

m = total number of frames

$$a_i = \text{allocation for } p_i = \frac{s_i}{S} \times m$$

$$m = 64$$

$$s_1 = 10$$

$$s_2 = 127$$

$$a_1 = \frac{10}{137} \times 62 \approx 4$$

$$a_2 = \frac{127}{137} \times 62 \approx 57$$

- Use a proportional allocation scheme using priorities rather than size
- If process P_i generates a page fault,
 - Select for replacement one of its frames
 - Select for replacement a frame from a process with lower priority number

- Global replacement – process selects a replacement frame from the set of all frames; one process can take a frame from another
 - But then process execution time can vary greatly
 - But greater throughput so more common

Non-Uniform Memory Access

- So far all memory accessed equally
- Many systems are NUMA – speed of access to memory varies
 - Consider system boards containing CPUs and memory, interconnected over a system bus
- Optimal performance comes from allocating memory “close to” the CPU on which the thread is scheduled
- And modifying the scheduler to schedule the thread on the same system board when possible

- Solved by Solaris by creating **lggroups**
 - Structure to track CPU / Memory low latency groups
 - Used my schedule and pager
 - When possible schedule all threads of a process and allocate all memory for that process within the lgroup

- **Second-Chance (Clock)
Page-Replacement Algorithm**
- **Enhanced Second-Chance Algorithm**
- **Counting Algorithms**
- **Page-Buffering Algorithms**

- Applications and Page Replacement
- Allocation of Frames
- Fixed Allocation
- Priority Allocation
- Global vs. Local Allocation
- Non-Uniform Memory Access



THANK YOU

Nitin V Pujari
Faculty, Computer Science
Dean - IQAC, PES University

nitin.pujari@pes.edu

For Course Deliverables by the Anchor Faculty click on www.pesuacademy.com and complete reading assignments provided by the Anchor on Edmodo