



OPERATING SYSTEMS

Memory Management - 12

Nitin V Pujari
Faculty, Computer Science
Dean - IQAC, PES University

OPERATING SYSTEMS

Course Syllabus - Unit 3



Unit-3: Unit 3: Memory Management: Main Memory

Hardware and control structures, OS support, Address translation, Swapping, Memory Allocation (Partitioning, relocation), Fragmentation, Segmentation, Paging, TLBs context switches

Virtual Memory - Demand Paging, Copy-on-Write, Page replacement policy - LRU (in comparison with FIFO & Optimal), Thrashing, design alternatives - inverted page tables, bigger pages.

Case Study: Linux/Windows Memory

OPERATING SYSTEMS

Course Outline



25	Main Memory: Hardware and control structures, OS support, Address translation	8.1	64.2
26	Dynamic linking, Swapping	8.2	
27	Memory Allocation (Partitioning, relocation), Fragmentation	8.3	
28	Segmentation	8.4	
29	Paging: OS Support, TLBs, Address Translation	8.5	
30	Structure of the Page Table	8.6	
31	Design Alternatives - Inverted Page Tables, Bigger Pages	8.7-8.8	
32	Virtual Memory: Demand Paging, Copy-OnWrite	9.1-9.3	
33	Page replacement policy - LRU	9.4	
34	FIFO & Optimal	9.5	
35	Thrashing	9.6	
36	Case Study: Linux/ Windows Memory Management	9.10	

- **Virtual Memory - Page replacement**
- **What happens if there is no free Frame ?**
- **Basic Page Replacement**
- **Page and Frame Replacement Algorithms**
- **Graph of Page Faults versus the number of Frames**

- **First-In-First-Out (FIFO) Algorithm**
- **FIFO illustrating Belady's Anomaly**
- **Optimal Page Replacement Algorithm**
- **Least Recently Used (LRU) Algorithm**
- **Use of a Stack to Record Most Recent Page References**

- **LRU Algorithm Implementation**
- **LRU Approximation Algorithm**
- **Second-Chance (clock) Page-Replacement Algorithm**
- **Enhanced Second-Chance Algorithm**
- **Counting Algorithms**
- **Page-Buffering Algorithms**

- Applications and Page Replacement
- Allocation of Frames
- Fixed Allocation
- Priority Allocation
- Global vs. Local Allocation
- Non-Uniform Memory Access

- Virtual Memory - Thrashing
- Demand Paging and Thrashing
- Working-Set Model
- Keeping Track of the Working Set
- Page-Fault Frequency

- If a process does not have “enough” pages, the page-fault rate is very high
 - Page fault to get page
 - Replace existing frame
 - But quickly need replaced frame back

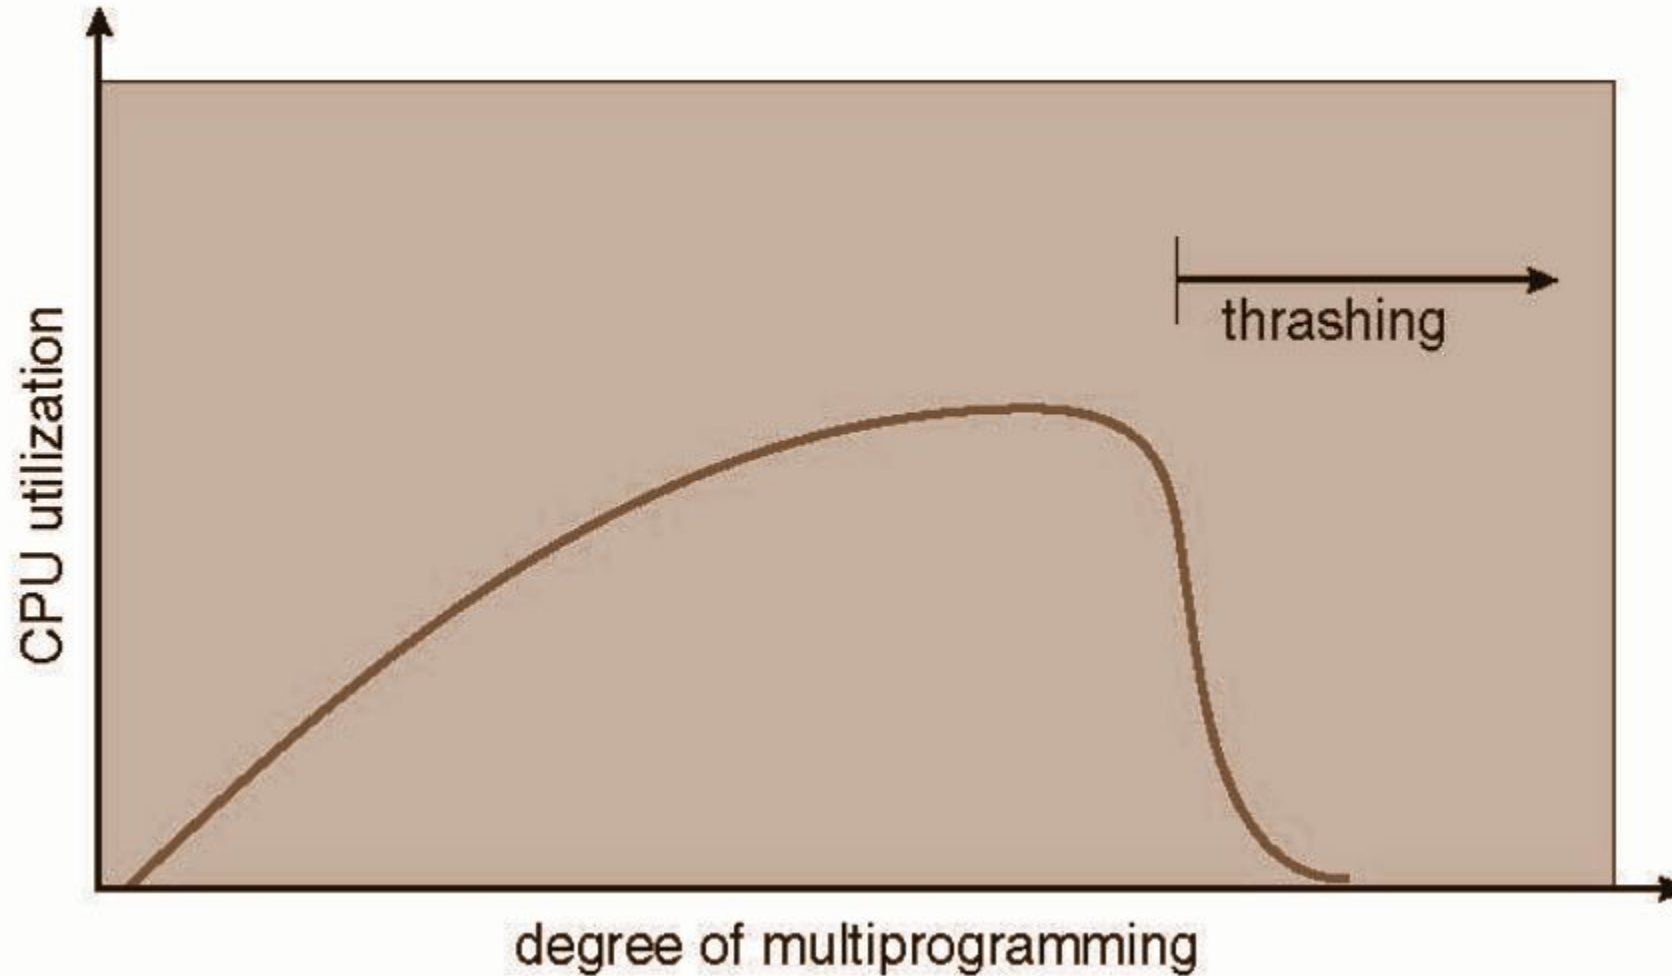
Virtual Memory - Thrashing



- If a process does not have “enough” pages, the page-fault rate is very high
 - This leads to:
 - Low CPU utilization
 - Operating system thinking that it needs to increase the degree of multiprogramming
 - Another process added to the system

Thrashing is a
process of keeping
system busy with
swapping pages **in**
and **out**

Virtual Memory - Thrashing



- Why does demand paging work?

- Locality model

- Process migrates from one locality to another

- Why does thrashing occur ?
 - **Size of Locality** > Total Memory Size
 - Limit effects by using local or priority page replacement

Working Set Model


- $\Delta \equiv$ working-set window \Rightarrow a fixed number of page references
- Example: 10,000 instructions
- WSSi (working set of Process P_i) = total number of pages referenced in the most recent (varies in time)
 - if Δ is too **small** \Rightarrow will **not** encompass entire locality
 - if Δ is too **large** \Rightarrow **will** encompass several localities
 - if $\Delta = \infty$ will encompass **entire program**
- $D = \sum WSS_i$ **total** demand frames
 - Approximation of locality

Working Set Model


- If $D > m \Rightarrow$ Thrashing
- Policy if $D > m$, then suspend or swap out one of the processes

page reference table

... 2 6 1 5 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 1 3 2 3 4 4 4 3 4 4 4 ...



$$WS(t_1) = \{1, 2, 5, 6, 7\}$$



$$WS(t_2) = \{3, 4\}$$

Working Set Model - Additional Input

- Working Sets: Conceptual model proposed by Peter Denning to prevent thrashing.
- **Informal definition:** the collection of pages a process is using actively, and which must thus be memory-resident to prevent this process from thrashing.
- If the sum of all working sets of all runnable threads or active process exceeds the size of memory, then stop running some of the threads or process for a while.

Working Set Model - Additional Input

- Divide processes into two groups: **active** and **inactive**:
 - When a process is **active** its entire working set must always be in memory: never execute a thread or a process whose working set is not resident.
 - When a process becomes **inactive**, its working set can migrate to disk.
 - Threads or child processes from inactive processes are never scheduled for execution.
 - The collection of **active** processes is called the **balance set**.
 - The system must have a mechanism for gradually moving processes into and out of **the balance set**.
 - As working sets change, **the balance set** must be adjusted.

Working Set Model - Additional Input

- How to compute working sets ?
 - Denning proposed a working set parameter T : all pages referenced in the last T seconds comprise the working set.
 - Can extend the clock algorithm to keep an idle time for each page.
 - Pages with idle times less than T are in the working set.
- Difficult questions for the working set approach:
 - How long should T be (typically minutes) ?
 - How to handle changes in working sets ?
 - How to manage the balance set ?
 - How to account for memory shared between processes ?

Keeping track of the Working Set

- Approximate with interval timer + a reference bit
- Example: $\Delta = 10,000$
 - Timer interrupts after every 5000 time units
 - Keep in memory 2 bits for each page
 - Whenever a timer interrupts copy and sets the values of all reference bits to 0
 - If one of the bits in memory = 1 page in working set
- Why is this not completely accurate ?
- Improvement = 10 bits and interrupt every 1000 time units

Current Situation - Additional Input

In practice, today's operating systems don't worry much about thrashing:

- With personal computers, users can notice thrashing and handle it themselves:
 - Typically, just buy more memory
 - manage balance set by hand
- Thrashing was a bigger issue for timesharing machines with dozens or hundreds of users:
 - **For a typical question like why should I stop my processes just so you can make progress ?**
 - **Answered by making the System to handle thrashing automatically.**
- Technology changes make it unreasonable to operate machines in a range where memory is even slightly overcommitted; better to just buy more memory.

Page-Fault Frequency




- More direct approach than Working set of process S
- Establish “acceptable” page-fault frequency (PFF) rate for a process and use local replacement policy
 - If actual rate too **low**, process **loses** frame
 - A low page fault rate indicates that the process has **too many frames**.
 - If actual rate too **high**, process **gains** frame
 - If the page fault rate is too high, it indicates that the process **has too few frames** allocated to it.

Page-Fault Frequency - Additional Input

Page Fault Frequency: another approach to preventing thrashing.

- Per-process replacement; at any given time, each process is allocated a fixed number of physical page frames.
- Monitor the rate at which page faults are occurring for each process.
- If the rate gets too high for a process, assume that its memory is overcommitted; increase the size of its memory pool.
- If the rate gets too low for a process, assume that its memory pool can be reduced in size.
- If the sum of all memory pools don't fit in memory, deactivate some processes.

Page-Fault Frequency - Additional Input



- Virtual Memory - Thrashing
- Demand Paging and Thrashing
- Working-Set Model
- Keeping Track of the Working Set
- Page-Fault Frequency



THANK YOU

Nitin V Pujari
Faculty, Computer Science
Dean - IQAC, PES University

nitin.pujari@pes.edu

For Course Deliverables by the Anchor Faculty click on www.pesuacademy.com