

Introduction to Apache ZooKeeper™



<http://zookeeper.apache.org/>

Who am I?

- Saurav Haloi
- Engineer at Symantec
- Work in Hadoop & Distributed System
- FOSS enthusiast

What is a Distributed System?

“A distributed system consists of multiple computers that communicate through a computer network and interact with each other to achieve a common goal.”

- Wikipedia

Fallacies of Distributed Computing

- The network is reliable.
- Latency is zero.
- Bandwidth is infinite.
- The network is secure.
- Topology doesn't change.
- There is one administrator.
- Transport cost is zero.
- The network is homogeneous.

Reference: http://en.wikipedia.org/wiki/Fallacies_of_Distributed_Computing

- *Coordination*: An act that multiple nodes must perform together.
- Examples:
 - Group membership
 - Locking
 - Publisher/Subscriber
 - Leader Election
 - Synchronization
- Getting node coordination correct is very hard!



Copyright by Shamus O'Reilly via Flickr

“ZooKeeper allows distributed processes to coordinate with each other through a shared hierarchical name space of data registers.”

- ZooKeeper Wiki

ZooKeeper is much more than a
distributed lock server!


What is ZooKeeper?

- An open source, high-performance coordination service for distributed applications.
- Exposes common services in simple interface:
 - naming
 - configuration management
 - locks & synchronization
 - group services

... developers don't have to write them from scratch
- Build your own on it for specific needs.

- Configuration Management
 - Cluster member nodes bootstrapping configuration from a centralized source in unattended way
 - Easier, simpler deployment/provisioning
- Distributed Cluster Management
 - Node join / leave
 - Node statuses in real time
- Naming service – e.g. DNS
- Distributed synchronization - locks, barriers, queues
- Leader election in a distributed system.
- Centralized and highly reliable (simple) data registry

The ZooKeeper Service



- ZooKeeper Service is replicated over a set of machines
- All machines store a copy of the data (in memory)
- A leader is elected on service startup
- Clients only connect to a single ZooKeeper server & maintains a TCP connection.
- Client can read from any Zookeeper server, writes go through the leader & needs majority consensus.

Image: <https://cwiki.apache.org/confluence/display/ZOOKEEPER/ProjectDescription>

The ZooKeeper Data Model

- ZooKeeper has a hierarchal name space.
- Each node in the namespace is called as a *ZNode*.
- Every ZNode has data (given as byte[]) and can optionally have children.

```
parent : "foo"
```

```
|-- child1 : "bar"
```

```
|-- child2 : "spam"
```

```
`-- child3 : "eggs"
```

```
    |-- grandchild1 : "42"
```

- ZNode paths:
 - canonical, absolute, slash-separated
 - no relative references.
 - names can have Unicode characters

- Maintain a stat structure with version numbers for data changes, ACL changes and timestamps.
- Version numbers increases with changes
- Data is read and written in its entirety

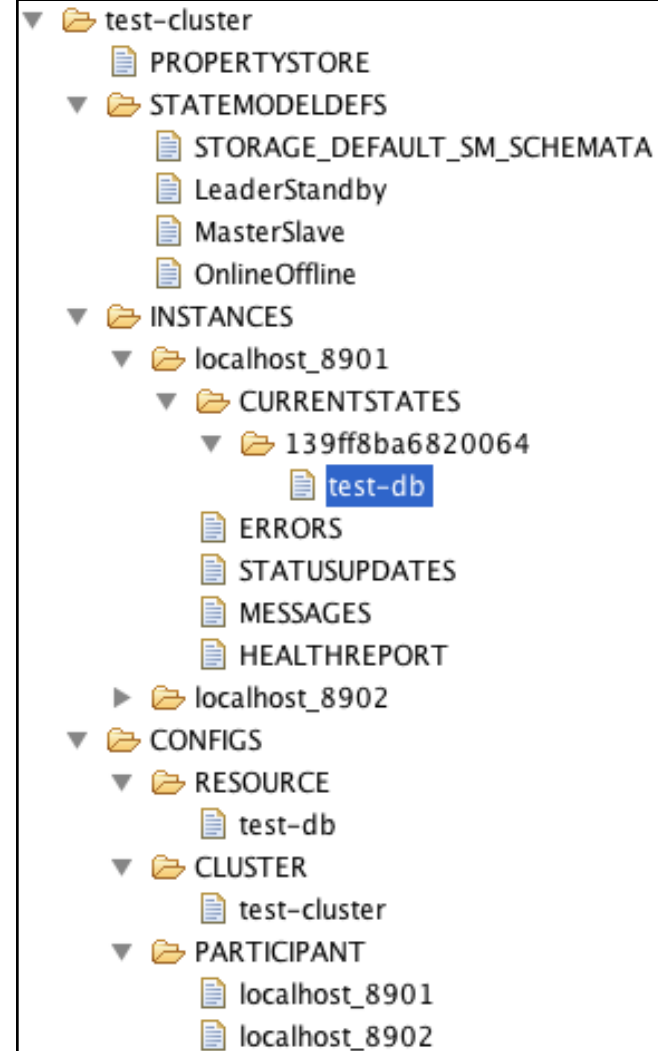



Image: <http://helix.incubator.apache.org/Architecture.html>

- Persistent Nodes
 - exists till explicitly deleted
- Ephemeral Nodes
 - exists as long as the session is active
 - can't have children
- Sequence Nodes (Unique Naming)
 - append a monotonically increasing counter to the end of path
 - applies to both persistent & ephemeral nodes

ZNode Operations

Operation	Type
create	Write
delete	Write
exists	Read
getChildren	Read
getData	Read
setData	Write
getACL	Read
setACL	Write
sync	Read

ZNodes are the main entity that a programmer access.

[zk: localhost:2181(CONNECTED) 0] help

ZooKeeper -server host:port cmd args

connect host:port

get path [watch]

ls path [watch]

set path data [version]

rmr path

delquota [-n|-b] path

quit

printwatches on|off

create [-s] [-e] path data acl

stat path [watch]

close

ls2 path [watch]

history

listquota path

setAcl path acl

getAcl path

sync path

redo cmdno

addauth scheme auth

delete path [version]

setquota -n|-b val path

[zk: localhost:2181(CONNECTED) 1] ls /

[hbase, zookeeper]

[zk: localhost:2181(CONNECTED) 2] ls2 /zookeeper

[quota]

cZxid = 0x0

ctime = Tue Jan 01 05:30:00 IST 2013

mZxid = 0x0

mtime = Tue Jan 01 05:30:00 IST 2013

pZxid = 0x0

cversion = -1

dataVersion = 0

aclVersion = 0

ephemeralOwner = 0x0

dataLength = 0

numChildren = 1

[zk: localhost:2181(CONNECTED) 3] create /test-znode HelloWorld

Created /test-znode

[zk: localhost:2181(CONNECTED) 4] ls /

[test-znode, hbase, zookeeper]

[zk: localhost:2181(CONNECTED) 5] get /test-znode

HelloWorld

- Clients can set watches on znodes:
 - NodeChildrenChanged
 - NodeCreated
 - NodeDataChanged
 - NodeDeleted
- Changes to a znode trigger the watch and ZooKeeper sends the client a notification.
- Watches are one time triggers.
- Watches are always ordered.
- Client sees watched event before new znode data.
- Client should handle cases of latency between getting the event and sending a new request to get a watch.

- API methods are sync as well as async

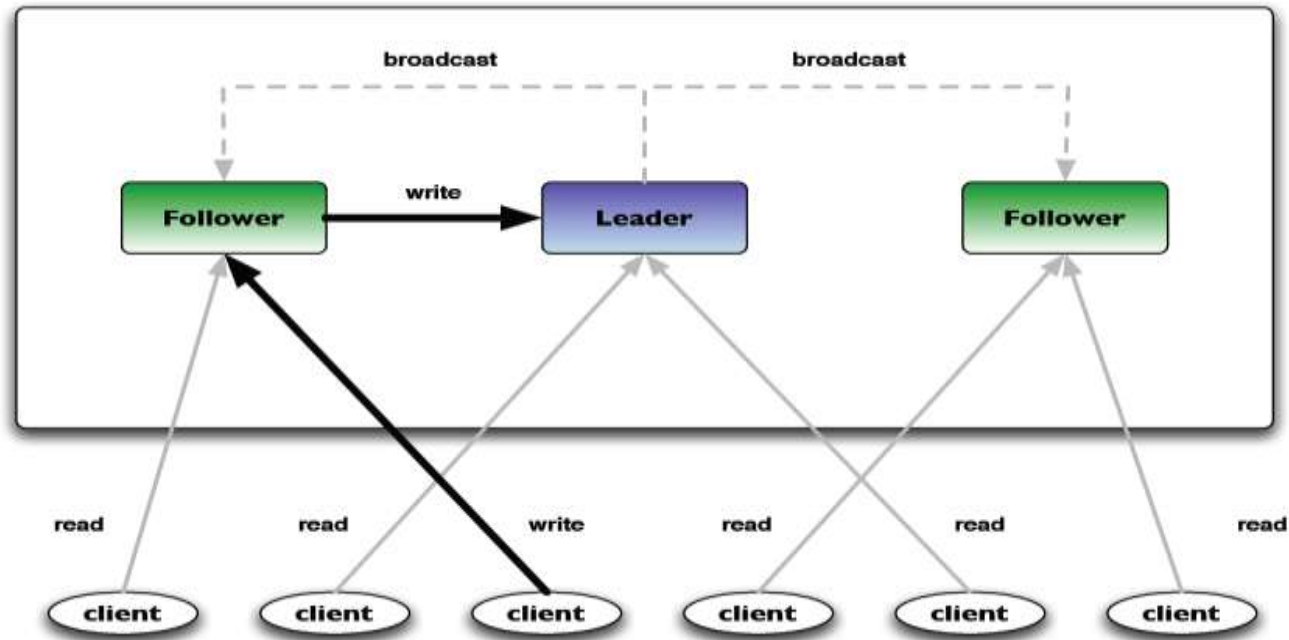
- Sync:

```
exists("/test-cluster/CONFIGS", null);
```

- Async:

```
exists("/test-cluster/CONFIGS", null, new StatCallback() {  
    @Override  
    public processResult(int rc, String path, Object ctx, Stat stat)  
    {  
        //process result when called back later  
    }  
}, null);
```

ZNode Reads & Writes



- Read requests are processed locally at the ZooKeeper server to which the client is currently connected
- Write requests are forwarded to the leader and go through majority consensus before a response is generated.


Image: <http://www.slideshare.net/scottleber/apache-zookeeper>

- **Sequential Consistency:** Updates are applied in order
- **Atomicity:** Updates either succeed or fail
- **Single System Image:** A client sees the same view of the service regardless of the ZK server it connects to.
- **Reliability:** Updates persists once applied, till overwritten by some clients.
- **Timeliness:** The clients' view of the system is guaranteed to be up-to-date within a certain time bound. (Eventual Consistency)

Recipe #1: Cluster Management

Each Client Host i , $i:=1 \dots N$

1. Watch on `/members`
2. Create `/members/host- $\{i\}$` as ephemeral nodes
3. Node Join/Leave generates alert
4. Keep updating `/members/host- $\{i\}$` periodically for node status changes
(load, memory, CPU etc.)



Recipe #2: Leader Election

1. A znode, say “/svc/election-path”
2. All participants of the election process create an ephemeral-sequential node on the same election path.
3. ***The node with the smallest sequence number is the leader.***
4. Each “follower” node listens to the node with the next lower seq. number
5. Upon leader removal go to election-path and find a new leader, or become the leader if it has the lowest sequence number.
6. Upon session expiration check the election state and go to election if needed

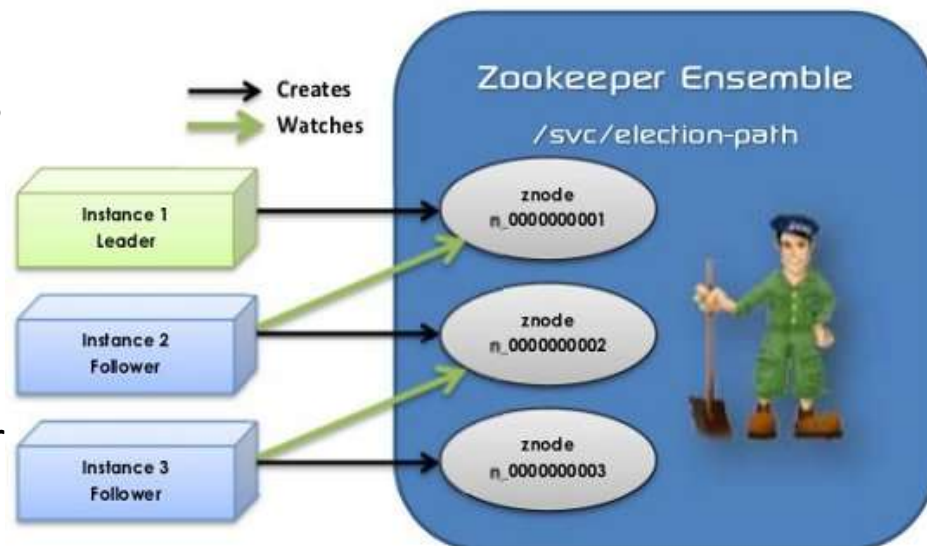


Image: <http://techblog.outbrain.com/2011/07/leader-election-with-zookeeper/>

Recipe #3: Distributed Exclusive Lock

Assuming there are N clients trying to acquire a lock

- Clients creates an ephemeral, sequential znode under the path /Cluster/_locknode_
- Clients requests a list of children for the lock znode (i.e. _locknode_)
- ***The client with the least ID according to natural ordering will hold the lock.***
- Other clients sets watches on the znode with id immediately preceding its own id
- Periodically checks for the lock in case of notification.
- The client wishing to release a lock deletes the node, which triggering the next client in line to acquire the lock.

ZK

```
|---Cluster
+---config
+---memberships
+---_locknode_
+---host1-3278451
+---host2-3278452
+---host3-3278453
+--- ...
\---hostN-3278XXX
```

- ZooKeeper ships client libraries in:
 - Java
 - C
 - Perl
 - Python
- Community contributed client bindings available for Scala, C#, Node.js, Ruby, Erlang, Go, Haskell
<https://cwiki.apache.org/ZOOKEEPER/zkclientbindings.html>

A few points to remember

- Watches are one time triggers
 - Continuous watching on znodes requires reset of watches after every events / triggers
- Too many watches on a single znode creates the “*herd effect*” - causing bursts of traffic and limiting scalability
- If a znode changes multiple times between getting the event and setting the watch again, carefully handle it!
- Keep session time-outs long enough to handle long garbage-collection pauses in applications.
- Set Java max heap size correctly to *avoid swapping*.
- Dedicated disk for ZooKeeper transaction log

Who uses ZooKeeper?

Companies:

- Yahoo!
- Zynga
- Rackspace
- LinkedIn
- Netflix
- *and many more...*

Projects in FOSS:

- Apache Map/Reduce (Yarn)
- Apache HBase
- Apache Solr
- Neo4j
- Katta
- *and many more...*

Reference: <https://cwiki.apache.org/confluence/display/ZOOKEEPER/PoweredBy>

ZooKeeper In Action @Twitter

- Used within Twitter for service discovery
- How?
 - Services register themselves in ZooKeeper
 - Clients query the production cluster for service “A” in data center “XYZ”
 - An up-to-date host list for each service is maintained
 - Whenever new capacity is added the client will automatically be aware
 - Also, enables load balancing across all servers.



Reference: <http://engineering.twitter.com/>

- The Chubby lock service for loosely-coupled distributed systems
Google Research (7th USENIX Symposium on Operating Systems Design and Implementation (OSDI), {USENIX} (2006))
- ZooKeeper: Wait-free coordination for Internet-scale systems
Yahoo Research (USENIX Annual Technology Conference 2010)
- Apache ZooKeeper Home: <http://zookeeper.apache.org/>
- Presentations:
 - <http://www.slideshare.net/mumrah/introduction-to-zookeeper-trihug-may-22-2012>
 - <http://www.slideshare.net/scottleber/apache-zookeeper>
 - <https://cwiki.apache.org/confluence/display/ZOOKEEPER/ZooKeeperPresentations>

- The Google File System
- The Hadoop Distributed File System
- MapReduce: Simplified Data Processing on Large Clusters
- Bigtable: A Distributed Storage System for Structured Data
- PNUTS: Yahoo!'s Hosted Data Serving Platform
- Dynamo: Amazon's Highly Available Key-value Store
- Spanner: Google's Globally Distributed Database
- Centrifuge: Integrated Lease Management and Partitioning Cloud Services (Microsoft)
- ZAB: A simple totally ordered broadcast protocol (Yahoo!)
- Paxos Made Simple by Leslie Lamport.
- Eventually Consistent by Werner Vogel (CTO, Amazon)
- <http://www.highscalability.com/>

Questions?

Thank You!

Saurav Haloi

saurav.haloi@yahoo.com

Twitter: sauravhaloi